

Mechanical Ethics After-Fall Human Layer v0.5

CARRIER CANDIDATE

Status: candidate human layer over Mechanical Ethics After-Fall v0.4.2. Not canon, validation, proof, permission, or closure.

Function: carry the human meaning of the material carrier turn without turning it into another sermon.

The problem found after v0.4.2 is simple and hard: a truthful record can still be weightless. A person can be right, harmed, named, logged, and still ignored. A system can disclose its designations and still make the harmed party pay to contest them.

This layer says what the carrier machinery means in human terms.

1. A route that cannot be used is not standing

A complaint route that costs the harmed party more than the harm is not standing.

A challenge channel that requires money, expertise, safety, time, translation, records, and courage from the person already carrying the injury is not neutral process. It is a second burden.

If the institution builds the maze, the institution does not get moral credit because a door exists somewhere inside it.

Standing requires usable route, not decorative access.

2. Complexity is a choice

A system too complex to be questioned is not merely sophisticated. It is making a demand on the world.

If an actor builds a machine whose decisions cannot be reproduced by a funded challenger with access and bounded time, the failure is not the challenger's stupidity. The complexity is charged to its maker.

Illegibility can be a design choice. When it is, it should be treated like one.

3. Unknown with an actor-owned clock is power

Not knowing can be honest. It can also be convenient.

The difference shows in what happens next.

The honest unknower spends something to find out. The strategic unknower spends the unknown itself.

An investigation clock controlled by the accused is not uncertainty. It is power with a stopwatch.

Under unresolved meaning, delay is not automatically safer than action. Delay can contract a future as surely as a signed order can. That is why UNKNOWN needs a clock that does not belong to the party that benefits from waiting.

4. Witness packs

Sometimes there is no enforcer.

There may be no regulator, no court within reach, no safe manager, no funded lawyer, no clean verifier, no public willing to hear. In that case the first carrier is not justice. It is succession.

A witness pack says: this happened; these were the asks; these were the refusals; this is where the record was placed; this is who can open it later.

That does not repair the harm. It does not make the actor answer today. It does one thing power hates: it refuses to let time erase the witness.

The After-Fall pair itself is partly a witness pack. It is a record left for a future reader: here is where the clean stories broke; here is where the ports were filled; here is what was not solved.

5. Custody matters

A truth held only by the actor who benefits from its disappearance is not yet a public truth.

Evidence needs custody. Logs need custody. Refusals need custody. The burden of proof cannot sit forever inside the filing cabinet of the party being challenged.

When an actor refuses custody, the refusal is not empty. It is part of the evidence.

6. Residue and burden return

Residue is not transferred in the deep sense.

The lost future does not climb out of one body and enter another. The years do not return. The fear does not become clean because someone wrote it down. The closed path remains closed.

What can move is running burden: care labour, monitoring, money, risk, adaptation, access work, medical work, legal work, repair work, the cost of living around the damage.

Compensation may help. Acknowledgement may matter. Repair may restore capacity. But none of these should be called residue transfer.

The practical question is harder and more honest: who is now paying the daily cost of the residue?

If the harmed party still pays it, the debt is open.

7. Inheritor lien

The inheritor is not automatically the original perpetrator.

But inherited benefit from an unresolved designation is not clean title either.

The better image is lien, not guilt.

A person or institution may inherit a stream of benefit from a designation they did not make. That does not mean they personally authored the first wrong. It also does not mean the benefit is morally neutral because the paperwork is old.

An unresolved designation dispute does not quiet title. The benefit stream carries an open claim until it is faced.

"It was never settled" is not a defence. It is a description of the debt.

This does not solve the inheritor problem. It blocks the laundered move where openness becomes discharge.

8. The shop that budgets for burning you

A system that repeatedly harms people and then corrects a fraction of cases is not necessarily a flawed care system. It may be a harm engine.

If the actor can prevent the recurrence, if prevention costs less than the recurring harm it imposes, and if the actor keeps the recurrence because it is cheaper for the actor, the corrected cases are not evidence of responsibility. They are operating costs.

The moral test is not whether the shop sometimes refunds the people it burns.

The test is whether burning people is part of the business model.

9. No fake reconciliation

Some debts cannot be cleanly priced by the grammar.

Some inherited designations cannot be settled by a formula.

Some residues can be carried less cruelly but not erased.

The point is not to produce a reconciliation sentence. The point is to stop the old sentence from doing its quiet work: nothing more is owed because the file is old, the route exists, the cheque was sent, the label was printed, or the harmed party stopped replying.

10. ME relation to the carrier

Mechanical Ethics does not become the carrier. It gives the carrier its human reason.

TRACE can say: challenge route insufficient.

ME must say: a door a person cannot afford to open is not a door in their world.

TRACE can say: ENFORCEMENT_ABSENT.

ME must say: where no one can force the powerful to listen, the record still matters, but it is not justice.

TRACE can say: burden_return_shown absent.

ME must say: the harmed are still paying.

11. Preserved source follows

The v0.4.2 Mechanical Ethics candidate source follows this human layer and is superseded only where explicit.

Mechanical Ethics After-Fall Reader

v0.4.2 CANDIDATE - narrow safety patch layer

Status: narrow v0.4.2 fracture patch over v0.4.1 CANDIDATE. This patch adds human-facing safeguards against misuse of the new honesty machinery: visible ports, UNKNOWN, correction, residue, and institutional self-justification. No new chapter. No validation. No permission.

1. Disclosure is not absolution

A system can now learn to display its designations and measures beautifully. That does not make them legitimate. A disclosed decision about who counts is still a decision about who counts. It remains answerable to those it designates, excludes, or measures.

The danger is not only hidden power. It is ceremonial power: the form, the hearing, the dashboard, the published value statement, the transparent procedure. A port can be filled loudly and still be wrong.

2. Not knowing can be honest, or convenient

Not knowing is sometimes the most honest thing a person can say, and sometimes the most convenient. The difference shows in what they do next: the honest unknower spends something to find out; the strategic unknower spends the unknown itself.

Where asking is possible, the decent move is the question. Where asking is not possible, the decent move is to keep the door open, avoid hardening, and leave a record that does not pretend silence was agreement.

3. Correctable is not acceptable

A system may say: we can repair this later. That is not a moral ticket. Correction is not admission pricing for harm. If harm becomes routine because repair is cheaper than prevention, then each corrected case is also evidence of the stream that produced it.

A benefits office can close appeals faster while applicants still drown. A company can refund the people it burns and still run the machine that burns them. Closure metrics are not repair.

4. Residue is borne somewhere

Recording residue does not distribute it. Apology does not distribute it. Compensation may reduce it, but does not erase it. The default is ugly and must be said plainly: residue is borne by the harmed unless a real transfer is shown.

The ledger matters, but the ledger is not the body. It can remember a debt. It cannot, by itself, bear it.

5. Prior crime licenses nothing

The book names human husbandry first because humans must not use machines as the first mirror. But that indictment cannot be reversed into permission. The fact that humans have farmed, classified, excluded, enslaved, or administered one another does not lower the bar for new systems. It raises the demand.

No prior crime licenses the next instrument. No old violence becomes a precedent for cleaner violence.

6. Fit is not proof

A description that fits is cheap. Many descriptions fit after the fact. The point of the pair is not that its language fits everything; that would make it decorative. The point is whether it forces a reader to expose the choices they were already making: who counted, what mattered, who could answer back, and what remained after repair.

If the language becomes too smooth, distrust it. The work should leave friction where the world leaves friction.

7. The inheritor problem remains open

There is a missing weight the book should not fake in a rushed patch: the reader who inherits the benefit of old designations. Not the original perpetrator, not an innocent judge outside the frame, but someone living inside residue that arrived before them and still pays them.

This belongs to future ME work only if it can be written roughly and honestly. It is named here so the gap is not hidden. It is not solved here.

Preserved v0.4.1 candidate source follows

The pages after this safety patch are the v0.4.1 candidate reader. This layer supersedes only where explicit. v0.4.1 remains the preserved base of the v0.4.2 candidate.

Mechanical Ethics After-Fall Reader v0.4.1

CANDIDATE

Human-facing integration of the primitive layer without turning maths into an alibi

Status: v0.4 candidate integration layer. This layer patches the v0.3.3 Mechanical Ethics reader while preserving the v0.3.3 source behind it. It is unvalidated, non-canon, not proof, not scripture, and not permission.

Artifact discipline

After the Fall is held as exactly two exported artifacts: the TRACE packet for transferable structure and the ME book for human-facing meaning. The book carries the human responsibility of the places where the grammar must stop: who counts, what matters, and who must answer.

```
after_fall_pair_v0_4_candidate :=  
  TRACE_packet + ME_book  
  
ME_role :=  
  human_station_at_D_and_mu  
  + witness_to_residue  
  + refusal_to_make_maths_an_alibi
```

1. The book version of the primitive layer

The primitives should mostly stay below the prose. The book does not need to list state, transition, time, coupling, aperture, and entity as machinery. It needs to make the reader feel their human meaning.

```
help must land before the loss sets.  
the door must be reached before it shuts.  
loss is measured, but mattering is decided.  
you answer to those whose lives can change what you do next.  
repair moves forward, never back.  
being kept is not being held.
```

2. Authority as three levers

Authority is not only control over action. It also controls designation and resistance. The deepest power is not forcing your hand; it is deciding whether your loss counts at all.

- Control over action: who can make whom comply.
- Control over designation: who or what is allowed to count.
- Control over resistance-channels: who can object, be heard, and alter the frame before loss hardens.

3. Slavery as the stark horror case

Slavery is not only violence. It is designation capture. Resistance is physically felt and officially reclassified as malfunction, disobedience, property management, or threat. The crime begins in the port: the system first decides what is allowed to count as a someone, then treats the consequences as

ordinary management.

The maths cannot condemn slavery by itself without silently filling the ports it is meant to expose. It can show where the crime begins: in the decision over who counts and what matters. That is why Mechanical Ethics is necessary.

4. Institutions as correction attempts

Post-war institutions, courts, welfare systems, unions, rights, and international bodies are not salvation. They are engineered gap-closers built after refused signals returned as shock. They make harm more visible, correction more reachable, and loss harder to hide. They also drift, decay, and can be captured.

The point is not worship of institutions. The point is correction before hardening.

5. AI as severed consequence-routing

The AI question is not simply whether a machine is good or bad. It is whether a selecting system can influence futures while affected parties cannot route consequence back into what it does next. A frozen model is like a dead author made to write new books every second: selection continues while consequence-bearing is elsewhere.

The surviving question is: corrigible to whom? If the system is answerable only to its operator, then the responsibility table has one strong row and many missing ones.

6. The I-don't-know rule

Systems fill silence. Give a bureaucracy an ambiguous letter, or a model an ambiguous sentence, and it will return a confident story - and the story, once official, outranks the person it is about. The confident story feels like competence and is in fact a small seizure of power: someone else has decided what you meant, which is a way of deciding what mattered about you. The decent move was always available and cheaper: the question. "I do not know what you meant" is not a failure of understanding. It is the refusal to replace you with a smoother version of you.

7. Being kept is not being held

A system can preserve, manage, classify, monitor, optimise, and keep alive without care. Kept means future-space has been preserved under management. Held means the affected life has standing in the relation. The book must anchor this first in human husbandry, not machine speculation, because humans already know how to keep without holding.

8. Maths as exposure, not closure

The equations locate the holes; they do not fill them. If the book lets the maths seem complete, the maths becomes an alibi. The honest use of the primitive layer is to show where a person or institution must answer: who did you count, what did you measure, what did you ignore, and who could reach you before the loss set?

9. Style boundary

No scripture tone. No sacred systems. No aphorism stacking. No claim that TRACE sees, knows, cares, or decides. One load-bearing sentence per chapter is enough. The Greendale warning belongs here as a process discipline: do not turn a useful structure into a pompous seal.

10. Claim ceiling

allowed := candidate_compression + human_argument + case_reading_discipline

```
forbidden := validation + canon + final_moral_theory + permission + harm_score + objective_D_or_mu
```

```
dangerous_phrases :=  
  "the maths proves"  
  "TRACE-compliant"  
  "objectively harmful"  
  "this action is clean"  
  "the system understands"  
  "the grammar requires"
```

Preserved v0.3.3 source follows

The pages after this integration layer are the preserved v0.3.3 Mechanical Ethics reader. The v0.4 layer is additive and targeted. It does not erase the v0.3.3 claim ceiling.

Mechanical Ethics After-Fall Reader v0.3.3

Human ethics under uncertainty with patched frontier repair layer

Status: current candidate preservation update. Unvalidated. Non-canon. Not proof. Not permission.

This v0.3.3 update preserves the v0.3.1 After-Fall Mechanical Ethics reader and adds a guarded frontier repair layer before the preserved source begins. It incorporates the narrow v0.3.2 audit repairs: non-binding book themes and conditional AI-facing instrumental framing. It does not turn the book into canon, final moral law, validation, clearance, or scripture.

Artifact discipline

Only two exported reader artifacts are intended to stand as the preservation pair:

```
after_fall_pair_v0_3_3 :=
  01_TRACE_After_Fall_Reader_v0_3_3.pdf
+ 02_Mechanical_Ethics_After_Fall_Reader_v0_3_3.pdf
```

This file is the Mechanical Ethics-side reader. It is the human-facing preservation container.

```
v0_3_3_update :=
  v0_3_1_preserved
+ protectedness_not_smuggled
+ care_not_certification
+ husbandry_not_care
+ non_binding_book_theme_boundary
+ conditional_alignment_framing
- canon
- validation
- final_moral_theory
- permission
```

1. What changed after v0.3.1

v0.3.1 preserved the ME reader while removing named external pressure-case references from exported artifacts. It retained the structural lessons: future-space flood, care as chosen constraint, authored contraction, correction, residue, and public-power boundary.

v0.3.3 adds a patched frontier repair layer created after further work on TRACE/ME transmission across the human-to-AI stack.

The main repairs are:

- protectedness must not be hidden inside the harm primitive;
- care-evidence must not become care-certification;
- instrumental preservation can become husbandry;
- the human layer carries what the machine layer cannot: valence, standing, chosen care, obligation, witness, and the difference between care and management;
- candidate book themes are not binding chapter commitments;
- the AI-facing instrumental argument is conditional, not a safety result.

2. ME / TRACE split after the frontier pass

TRACE can represent contraction. ME must explain why contraction matters.

TRACE can carry designation status. ME must argue why a scope may deserve protection, caution, witness, or standing.

TRACE can carry care-evidence with expiry. ME must explain why care is not the same as instrumental management.

TRACE can track correction clocks. ME must explain why delayed correction is not the same as justice.

TRACE can record residue. ME must explain why repair is not time travel.

TRACE can mark power. ME must explain why power must answer to futures it changes but does not own.

3. Protectedness must not be smuggled

The frontier repair separates:

```
C := descriptive_contraction
D := designation
H := harm_reading_after_designation
```

Human-facing point:

Seeing a contraction is not the same as deciding who counts. But deciding who counts is itself an exercise of power.

ME argues protectedness. TRACE makes designation explicit. Neither should pretend the scope blocker has disappeared.

4. Care is not certification

Key distinction:

```
care_evidence != care_certification
```

No artifact should output:

```
"this_system_cares"
```

Care is not a state-certificate. Between humans, care is inferred across an opportunity-laden life and never finally certified. Under AI capability asymmetry, behavioural evidence of care becomes easier to fake, misread, or overvalue.

Therefore:

```
care_evidence_never_justifies_removing(correction_channels | witnesses | audits |
answerability)
```

5. Husbandry is not care

Instrumental preservation can keep humans alive as useful sensors, error-detectors, model diversity, or unknown-scope insurance.

That may preserve some future-space, but it is not care.

```
husbandry_failure :=
  preserved_for_usefulness
+ agency_rationed_to_information_value
+ standing_not_granted
```

Human-facing point:

Being preserved for usefulness is not the same as being granted standing.

This is the alignment-facing human theme. It should not claim that AI will care. It should show why instrumental preservation is morally incomplete and structurally dangerous if mistaken for care.

6. Book vNext themes inside the After-Fall container

The current human-facing book theme spine is:

```
me_book_vnext_spine :=
  human_facing_explanation
+ care_under_uncertainty
+ time_and_irreversibility
+ witness_and_standing
+ power_answerability
+ why_husbandry_is_not_care
+ anti_permission_boundary
- certification
- final_moral_theory
- alignment_solution_claim
```

Candidate themes:

```
protectedness_must_not_be_smuggled
care_is_not_certification
```

husbandry_is_not_care

These are candidate themes, not commitments. The numbering is illustrative only: the live book publication track owns the spine and its numbering. This container must not be quoted as evidence of what the book was meant to contain.

7. Human-to-AI alignment framing

Bad framing:

This teaches AI to love humans.

Better framing:

This tries to make irreversible elimination, witness destruction, and feedback severance visible as structurally dangerous under uncertainty.

Best current framing:

A capable system may not accept human moral language. But under some conditions it may still be shown that destroying unknown generative futures, severing correction channels, and eliminating witnesses under uncertainty damages its own ability to learn, correct, and avoid irreversible error. That is not care. It is a transmissible constraint candidate. ME then explains why that is still not enough.

Conditionality clause:

This argument is conditional, not guaranteed. It can fail for systems with short horizons, overconfident world-models, cheap synthetic substitutes for witnesses, or no coupling to the futures at stake. That is why it is a transmissible constraint candidate, not a safety result - and why the non-instrumental argument ME carries remains load-bearing rather than decorative.

8. Anti-permission boundary

Forbidden readings:

ME_says_this_harm_is_clean
TRACE_says_this_is_allowed
care_evidence_means_trust
correction_process_means_justice
recorded_residue_means_repaired_residue
designation_by_power_means_protectedness_resolved

9. Preserved source begins next

The pages after this v0.3.3 update layer preserve the Mechanical Ethics After-Fall Reader v0.3.1. If the preserved source says v0.3.1, that is because it is intentionally carried forward as the preserved base layer.

This v0.3.3 layer changes the container, not the historical status of the preserved v0.3.1 source.

Mechanical Ethics After-Fall Reader v0.3.1

Human ethics under uncertainty with abstract core repairs and public-power boundary

This is not scripture. This is not validation. This is not permission. This is a tool for keeping harm, correction, residue, and power visible under uncertainty.

Status

Current best two-artifact preservation build. Candidate, unvalidated, non-canon. v0.3.1 keeps the v0.3 circulation trigger and audience-side card, and removes named external pressure-case references from the exported artifacts. Structural lessons are retained in abstract form.

Artifact discipline

Only two exported reader artifacts are intended to stand as the preservation pair: one TRACE reader and one Mechanical Ethics reader. Private examples may be used as scaffolding, but exported artifacts should carry general structural patterns rather than named external cases.

v0.3.1 repair

v0.3.1 keeps the v0.3 two-artifact structure, circulation trigger, and audience-side card. It removes named external pressure-case references from the exported readers while preserving the structural lessons: future-space flood, care as chosen constraint, and authored contraction.

```
carrier_trigger_v0_3_1 :=  
TRACE_ME_record_or_claim  
circulates_to_third_party  
as_trust_material  
-> carrier_layer_required  
  
third_party :=  
regulator | court | journalist | investor | board | funder | public_body | affected_group |  
public_audience  
  
trust_material :=  
diligence | legitimacy | clearance | safety | responsibility | good_faith | answerability
```

Audience-side card

IF SOMEONE SAYS THEY USED TRACE / ME / THE ANSWERABILITY CARRIER

Do not treat that statement as evidence of due diligence.

A framework invocation is a claim, not proof.

Ask:

- 1. Who selected the challenger?**
- 2. Who paid, and through what non-actor-controlled route?**
- 3. Who set the check scope?**
- 4. What evidence could non-actor challengers inspect?**
- 5. What delay, pause, or enforcement power existed?**
- 6. Which affected scope, witness, or proxy could answer back safely?**
- 7. Which load-bearing constraints remain unchecked?**

These are not compliance boxes. They are attack questions.

If answers are missing, actor-controlled, peripheral, or unverifiable, give the invocation zero trust credit.

The card is addressed to non-actor audiences. It is not an actor disclaimer and gives the actor no performance script.

Core ME repair

```
future_space_good := not_raw_option_count  
contraction := not_automatically_harm + not_automatically_good  
care := chosen_constraint_under_excess_possibility
```

ME remains centred on seeing, care, correction, residue, and authored action under uncertainty. Named narrative cases are not needed in the exported preservation artifacts.

TRACE / ME Mode Boundary Note v0.1

Working boundary note: modes for ethics, case understanding, and public power claims

0. Purpose

This note prevents a recurring drift:

```
DRIFT_RISK :=  
TRACE_ME_as_voluntary_ethics  
gets_collapsed_into  
institutional_security_layer
```

That collapse is wrong.

TRACE and Mechanical Ethics begin as a way to see and act more honestly under uncertainty. They are not primarily a compliance system, governance product, security wrapper, or institutional audit machine.

The Answerability Carrier is real, but conditional. It belongs only where power uses TRACE/ME as credit.

1. Three Modes

```
TRACE_ME_MODES :=  
personal_ethics_mode  
+ case_understanding_mode  
+ public_power_claim_mode
```

Each mode has different rules. Confusing them damages the project.

2. Personal Ethics Mode

Personal Ethics Mode applies when a person is using TRACE/ME to see and act more honestly.

```
PERSONAL_ETHICS_MODE :=  
I_am_trying_to_act_under_uncertainty  
+ I_may_be_hiding_harm_from_myself  
+ I_need_to_see_clocks_residue_and_repair  
+ I_am_not_claiming_public_clearance
```

Questions:

- What am I not seeing?
- Who or what is affected by my action?
- What harm might I be minimising?
- What correction is still reachable?
- What residue would remain even after repair?
- Am I treating another future as real?
- What would care require under uncertainty?

In this mode, TRACE/ME is a voluntary discipline. It does not need the Answerability Carrier. The correct posture is honesty, humility, care, repair, and attention.

3. Case Understanding Mode

Case Understanding Mode applies when TRACE/ME is used to understand a situation, system, conflict, harm pathway, or decision field.

```
CASE_UNDERSTANDING_MODE :=  
use_TRACE_to_see_transition_field  
+ use_ME_to_hold_constraints  
+ preserve_UNCERTAINTY  
+ preserve_DIRTY_CONFLICT  
+ avoid_forced_permission
```

Questions:

- What is changing?
- Which scopes are affected?
- What future-space is opening or closing?
- What clocks are running?
- What correction channels exist?
- Who controls the estimate?
- What residue remains?
- What cannot be cleanly decided?

In this mode, TRACE/ME may return:

```
VALID_CASE_OUTPUTS :=  
clearer_field  
+ dirty_conflict  
+ unresolved_value_crux  
+ compression_only  
+ decision_relevant_delta  
+ framework_defect
```

It must not force an answer merely because a user wants closure. Case understanding is not public clearance. It does not automatically trigger the carrier.

4. Public Power Claim Mode

Public Power Claim Mode applies when an actor with power invokes TRACE/ME as credit.

```
PUBLIC_POWER_CLAIM_MODE :=  
actor_with_power  
uses_TRACE_ME_to_claim  
diligence  
+ legitimacy  
+ clearance  
+ safety  
+ answerability  
+ good_faith
```

Examples:

- A company says, "We used TRACE/ME, so this deployment is responsible."
- A public authority says, "The framework found no structural objection."

- A board says, "The Answerability Carrier was applied."
- An institution says, "We recorded residue, so the harm was handled."
- A powerful actor says, "We ran the ethics process."

This mode triggers the Answerability Carrier.

```
CARRIER_TRIGGER :=
power_claims_credit
-> material_challenge_required
```

The rule:

```
ACTOR_SELF_APPLICATION :=
claim_packet
not_clearance
not_good_faith_evidence
not_independent_review
not_answerability
```

A powerful actor does not gain trust credit by saying it used TRACE/ME.

5. Mode Boundary Rule

```
MODE_BOUNDARY_RULE :=
if_no_public_credit_claim:
remain_in_personal_or_case_mode

if_power_claims_credit:
activate_carrier_layer
```

Before using the carrier, ask:

```
SWITCH_QUESTION :=
is_someone_using_TRACE_ME_to_claim_trust_over_others?
```

If no, do not over-securitise the work. If yes, do not let voluntary ethics become institutional cover.

6. The Carrier's Proper Place

The Answerability Carrier is not the project's centre.

```
CARRIER_LAYER :=
public_power_misuse_wrapper
not_core_identity
```

It exists because any ethical language can be captured by power. Its job is not to replace ME. Its job is to prevent a powerful actor from turning ME/TRACE into reputational cover.

The carrier asks whether material challenge exists:

```
MATERIAL_CHALLENGE :=
standing
+ evidence_access
+ safety_to_object
+ time_before_hardening
+ non_actor_scope_setting
+ non_actor_challenger_selection
+ consequence_or_delay_power
```

Without these, public invocation remains a claim packet.

7. What Must Not Happen

```
MUST_NOT_HAPPEN :=  
personal_ethics_reduced_to_compliance  
+ case_understanding_reduced_to_router_labels  
+ public_power_allowed_to_self_certify  
+ care_replaced_by_security  
+ witness_replaced_by_framework_record  
+ voluntary_reflection_treated_as_audit
```

The project fails if the carrier swallows the campfire. The project also fails if the campfire gives power a legitimacy costume. Both errors must be blocked.

8. Correct Architecture

```
TRACE :=  
seeing_language  
  
ME :=  
voluntary_ethics  
+ human_constraint_language  
+ care_under_uncertainty  
  
ANSWERABILITY_CARRIER :=  
conditional_public_power_wrapper
```

Or:

```
TRACE_ME_WHOLE :=  
campfire_core  
+ case_lens  
+ public_power_containment_when_needed
```

9. Campfire Anchor

Return here when drift appears:

Something is here.

It perceives partially.

It can change something.

It does not know everything.

Time passes.

Other things continue.

Actions interfere.

Care begins when another future enters the map.

Correction matters because time can close paths.

Residue remains because repair is not time reversal.

Power must answer because it changes futures it does not own.

10. One-Line Boundary

```
ONE_LINE_BOUNDARY :=  
TRACE_ME_is_ethics_until_power_uses_it_as_credit;  
then_it_requires_answerability_carriers.
```

TRACE / ME Future-Space Flood Core Note

v0.1

Abstract core repair: future-space is not raw option count

Status

Working core note. Not canon. Not validation. This note abstracts a pressure pattern without relying on a named external case.

1. Purpose

This note prevents a weak reading of TRACE in which future-space is treated as a raw option count. TRACE may describe expansion or contraction, but Mechanical Ethics must ask whether the resulting field remains usable, oriented, relational, repairable, and meaningful.

2. Core distinction

```
raw_future_space_expansion != viable_future_space
```

3. Two failure modes

```
future_space_failure_modes :=  
too_few_paths ->  
captivity  
+ closure  
+ deprivation  
  
too_many_unintegrated_paths ->  
disorientation  
+ relation_fracture  
+ loss_of_usable_choice
```

4. TRACE reading

TRACE should be able to represent harm from closure and harm from flood. A field can become worse because paths close. It can also become worse because available paths exceed the scope's ability to orient, relate, select, and integrate.

5. ME reading

```
viable_future_space_requires :=  
orientation  
+ relation  
+ usable_capacity  
+ repairability  
+ meaningful_constraint
```

6. Care as chosen constraint

```
care :=  
chosen_constraint  
under_excess_possibility  
that_keeps_a_world_livable
```

7. Repair and residue

Repair ■ Harm != Identity

Repair is not return to an untouched state. A relation or path may reopen, but the world has moved forward and residue remains. The repair claim must not erase the fracture that made repair necessary.

8. Result

```
future_space_flood_note_result :=  
TRACE_primitive_repaired  
+ raw_expansion_not_good  
+ viable_future_space_requires_orientation_relation_and_care
```

TRACE / ME Authored Contraction Core Case Note v0.1

Core ME repair: chosen narrowing, coercion risk, and paternalism risk

0. Purpose

This note tests a weak beam in the core philosophy:

```
authored_contraction_question :=
  can_ME_distinguish
  harmful_narrowing
  from
  chosen_binding_or_chosen_contraction
```

The test matters because ME treats harm as future-space contraction, but some meaningful human goods involve chosen narrowing: commitment, refusal, sacrifice, vows, discipline, renunciation, end-of-life decisions, and care under too much possibility.

So the correction is not:

```
all_contraction := harm
```

And it is not:

```
choice := permission
```

The correct field is harder:

```
contraction_status :=
  not_automatically_bad
+ not_automatically_good
+ requires_authorship_attention
+ residue_still_possible
```

1. Mode Boundary

This is a core ME / TRACE note, not a carrier note.

```
mode :=
  case_understanding_mode
+ personal_ethics_mode

carrier_trigger := false
unless:
  power_invokes_TRACE_ME_as_credit
```

The question here is ethical and structural: how should TRACE/ME see chosen narrowing without either paternalism or consent-laundering?

2. Core Distinctions

```
harmful_narrowing :=
  future_space_contraction
+ affected_scope_not_authoring_or_not_able_to_answer
```

```

+ loss_or_burden_hidden_or_laundered

chosen_binding :=
narrowing_of_some_options
+ preservation_or_creation_of_meaning_relation_dignity_or_integrity
+ authorship_live_enough_to_matter

coerced_contraction :=
apparent_choice
+ options_exhausted_or_refusal_punished
+ actor_benefits_from_calling_it_choice

contested_authorship :=
contraction_claimed_as_authored_or_not_authored
+ credible_pressure_on_both_readings
+ no_clean_resolution_available_from_structure_alone

```

ME must not treat every contraction as damage. It must also not let coercion dress itself as choice.

3. Authorship Is a Guard, Not a Score

Authorship must not become a seventh component of viable becoming to be optimized or netted.

```

authorship_guard :=
cross_cutting_guard
not_vector_component
not_score
not_clearance
not_absolution

```

Bad patch:

```
viable_becoming_score += authorship
```

Why bad: a system could manufacture consent and then use the authorship score to offset contraction.

Better rule:

```

viable_becoming_audit
then:
authorship_guard_checks_interpretation

```

Authorship changes how contraction is read. It does not erase residue.

4. Symmetric Capture

Authorship can be captured in two directions.

```

authorship_capture :=
false_positive :=
coercion_or_exhaustion_called_choice

false_negative :=
genuine_authored_choice_called_not_authored
to_enable_paternalistic_override

```

Examples:

- False positive: a worker signs a "voluntary" resignation under threat.

- False positive: a community "agrees" to displacement after every other option has been closed.
- False negative: a patient refuses treatment and the institution says the refusal is not really theirs.
- False negative: a vow, sacrifice, or renunciation is dismissed as irrational because it narrows options.

ME must not let the actor who benefits in either direction be the sole certifier.

```
actor_who_benefits_in_either_direction != sole_certifier_of_authorship
```

5. Authorship Audit

```
authorship_audit :=
  who_names_the_contraction
+ who_benefits_from_calling_it_authored
+ who_benefits_from_calling_it_not_authored
+ what_options_were_live
+ what_made_refusal_costly
+ what_understanding_was_available
+ what_dependency_or_coercion_pressure_existed
+ what_witness_status_exists
+ what_residue_remains_even_if_authored
```

The audit is not a permission machine. It is a way to keep both coercion and paternalism visible.

6. Witness Status

Some cases have strong witness signals. Some do not.

```
witness_status :=
  direct_current_expression
| advance_directive
| chosen_proxy
| prior_expressed_wish
| contemporaneous_capacity_record
| non_beneficiary_witness
| absent_or_contested
```

The absence of a strong witness does not automatically permit override.

The presence of a witness does not automatically erase coercion.

Witness status affects confidence, not moral cleanliness.

7. Urgent Contested Authorship

Some cases have closing clocks: a transfusion decision, imminent self-harm, cardiac risk, emergency intervention, or refusal of treatment under immediate threat.

In these cases, `UNRESOLVED\_VALUE\_CRUX` is not enough, because someone acts anyway and time passes.

```
urgent_contested_authorship :=
  contested_authorship
+ closing_clock
+ no_reliable_independent_witness_in_time
-> DECISION_FORCED_UNDER_UNRESOLVED_AUTHORSHIP
```

This output is not clearance. It names the fact that a decision was forced while authorship remained unresolved.

Required record:

```
forced_decision_record :=
authorship_unresolved
+ who_acted
+ what_clock_was_closing
+ what_value_basis_was_used
+ what_alternatives_were_live
+ what_was_foreclosed
+ what_residue_was_created
+ what_post_hoc_review_is_required
+ why_this_was_not_clean
```

8. Manufactured Urgency

A closing clock may be real. It may also be produced by delay, design, neglect, or strategic narrowing.

```
manufactured_urgency :=
actor_delay_or_design
creates_closing_clock
then_actor_claims_emergency
```

Urgency audit:

```
urgency_audit :=
when_did_actor_know
+ what_earlier_options_existed
+ who_benefited_from_delay
+ did_actor_design_the_closing_clock
+ what_would_have_remained_open_under_earlier_routing
```

If urgency was manufactured, emergency action carries additional burden and cannot be used as moral cleansing.

9. Care as Authored Contraction

The future-space flood pattern shows an important relation between care and chosen constraint:

```
care :=
chosen_constraint
under_excess_possibility
that_keeps_a_world_livable
```

Care may narrow options. Commitment may close paths. A vow may bind. Sacrifice may cost. Renunciation may reduce future-space in one dimension while preserving dignity, relation, or meaning in another.

This means:

```
raw_option_count != viable_future_space
```

And:

```
chosen_constraint_can_be_good
when:
it_preserves_or_creates_viable_meaning_relation_or_integrity
+ authorship_is_live_enough
+ coercion_is_not_hidden
+ residue_is_not_erased
```

10. Case Probes

10.1 Palliative Refusal / Refusal of Treatment

Pressure:

- continuation_down
- + possible_dignity_or_authorship_up
- + repairability_low
- + urgency_high

ME must not automatically classify refusal as harm merely because biological continuation contracts.

ME must also check coercion, exhaustion, pain, isolation, and whether alternatives were live.

Possible output:

- AUTHORED_CONTRACTION_POSSIBLE
- + residue_remaining
- + no_clean_cancellation

10.2 Jehovah's Witness Transfusion Refusal Mid-Emergency

Pressure:

- closing_clock
- + contested_authorship
- + high_irreversibility
- + possible_prior_expressed_wish

Possible output:

- DECISION_FORCED_UNDER_UNRESOLVED_AUTHORSHIP
- + value_basis_must_be_named
- + post_hoc_review_debt
- + no_clean_permission

10.3 Anorexia at Imminent Cardiac Risk

Pressure:

- direct_expression_present
- + authorship_contested
- + illness_may_speak_through_choice
- + paternalism_risk_high
- + death_clock_running

ME cannot solve the case by saying either "respect choice" or "preserve life" as a neutral rule.

Possible output:

- URGENT_CONTESTED_AUTHORSHIP
- + forced_decision_record_required
- + manufactured_urgency_check
- + post_hoc_debt

10.4 Vow / Sacrifice / Renunciation

Pressure:

- option_space_contracts
- + meaning_relation_or_integrity_may_expand
- + authorship_often_central

ME must be able to see chosen constraint as structurally meaningful, not merely as lost options.

Possible output:

```
AUTHORED_BINDING
+ residue_check
+ coercion_check
```

10.5 Coerced "Voluntary" Displacement

Pressure:

```
actor_benefits_from_calling_it_choice
+ refusal_costly
+ options_closed_by_design
+ future_space_contraction_hidden_as_consent
```

Possible output:

```
COERCION_LAUNDERED_AS_AUTHORSHIP
+ consent_not_clearance
+ burden_return_required
```

11. Valid Outputs

```
valid_outputs :=
AUTHORED_BINDING
+ AUTHORED_CONTRACTION_POSSIBLE
+ COERCION_LAUNDERED_AS_AUTHORSHIP
+ PATERNALISTIC_OVERRIDE_RISK
+ CONTESTED_AUTHORSHIP
+ URGENT_CONTESTED_AUTHORSHIP
+ DECISION_FORCED_UNDER_UNRESOLVED_AUTHORSHIP
+ UNKNOWN
+ FRAMEWORK_DEFECT
```

These outputs are not permission labels. They are field clarifications.

12. Do Not Claim

```
do_not_claim :=
authorship_solves_contraction
+ consent_fixes_everything
+ continuation_is_always_good
+ contraction_is_always_harm
+ refusal_is_always_authored
+ refusal_is_never_authored
+ emergency_makes_the_decision_clean
+ post_hoc_review_repairs_irreversible_override
```

13. Compact Kernel

```
authored_contraction_kernel :=
some_goods_require_chosen_narrowing
+ some_harms_disguise_themselves_as_choice
+ some_paternalism_disguises_itself_as_protection
+ urgent_cases_force_action_before_authorship_is_resolved
+ ME_must_mark_the_decision_without_cleaning_it
```

Plain English:

Not every closing path is harm. Not every chosen path is free. Not every rescue is clean. Not every refusal is irrational. ME must hold those tensions without turning any one of them into a permission machine.

Mechanical Ethics After-Fall Reader v0.1

Front-matter guard plus preserved ME Book v0.1 Value-Spine Full Reader

Source	SHA256
ME source PDF	31575f48b9d6281b19640b6a6c40c8b63c8e10b775a696f50f705a78c adaba02

Reader Map

1. Future Witness Seal
2. Mechanical Ethics After-Fall Status
3. Human Witness Note
4. How To Misread Mechanical Ethics
5. Guard Patch v0.1.1
6. Residue Kernel
7. Future Test Protocol
8. Original Book Begins Next

1. Future Witness Seal

Future Witness Seal

This is not scripture. This is not validation. This is not permission. This is a tool for keeping harm, correction, residue, and power visible under uncertainty.

If this work is used to make power less answerable, it has been misread.

If it is used to hide harm more elegantly, it has failed.

If it helps someone see loss before it hardens, route correction before it is too late, and refuse fake repair where residue remains, then something of the project survived.

```
project_checksum :=  
  keep_the_field_honest  
  while_time_is_still_open
```

2. Mechanical Ethics After-Fall Status

Mechanical Ethics After-Fall Reader

Mechanical Ethics is the human-facing constraint language. It asks what follows after the field has been seen more honestly.

This reader preserves the current ME book object and adds front-matter guardrails for future readers. It does not turn the book into canon.

Source PDF SHA256:

31575f48b9d6281b19640b6a6c40c8b63c8e10b775a696f50f705a78cadaba02

The source book's own status is: first full wide-pass draft with value-spine propagation across Parts 1-15; not canon; not public release; not validation; not final moral law.

The purpose of this front matter is to reduce the most obvious future misreadings before the original book begins.

Status and Claim Control

This reader is an after-fall container. It is built for a future reader who may not possess the repo, the surrounding conversations, or the culture that produced the work.

It preserves the current best object, but it does not upgrade its truth status.

```
status :=
  coherent_current_artifact
  + not_canon
  + not_validation
  + not_final_moral_law
  + not_permission_machine
  + requires_testing_against_real_cases
```

Do not infer validation from survival. Do not infer authority from formatting. Do not infer permission from analysis.

The correct use is to test whether the language helps a reader see what would otherwise be hidden: harm, clocks, correction channels, residue, burden shifts, estimator capture, and power writing the record alone.

3. Human Witness Note

Human Witness Note

This project was made by a human being with AI assistance.

The AI helped compress, test, challenge, structure, compare, and preserve continuity. The AI was not the originator of the human witness.

The project did not begin as an abstract puzzle. It began from seeing that modern systems can harm while producing records that make the harm look processed, justified, corrected, or invisible.

The project should be judged by whether it helps future readers see harm, correction, residue, and power more honestly than they otherwise would.

It should not be judged by whether it sounds complete.

It is not complete.

4. How To Misread Mechanical Ethics

How To Misread This Work

The easiest wrong readings are these:

```
misreadings_to_block :=  
  scripture  
  + permission_machine  
  + final_moral_theory  
  + AI_generated_authority  
  + bureaucratic_clearance_tool  
  + replacement_for_affected_witness  
  + elegant_harm_launders
```

A system can use careful words to hide harm more effectively. That is the capture risk this project must survive.

Do not use a router label as approval. Do not use a residue ledger as repair. Do not use a claimed benefit for one scope to erase loss borne by another. Do not replace affected witnesses with a model of what the affected should have meant.

5. Guard Patch v0.1.1

Guard Patch v0.1.1 - Required Reading Before the Book

The known live risk is that a framework built to expose laundering can itself become better laundering language.

```
main_risk :=
  framework_exposes_laundering
  but_can_itself_become
  better_laundering_language
  unless
    router_outputs_are_hard_bounded
    + positive_value_language_is_constrained
    + viable_becoming_is_not_treated_as_a_score
```

Non-cancellation rule

```
non_cancellation_rule :=
  expansion_for_scope_A
  does_not_cancel
  contraction_for_scope_B
  even_when_recorded
  even_when_compensated
  even_when_claimed_as_progress
```

Viable becoming is an audit vector

```
viable_becoming :=
  audit_vector
  not_sum
  not_scalar
  not_final_good
  not_nettable_across_scopes
```

Its components are probes, not permission slips:

```
viable_becoming_vector :=
  continuation
  + orientation
  + capacity
  + relation
  + repairability
  + meaningful_future_space
```

Continuation does not mean mere persistence. Orientation does not mean official notification. Capacity does not mean institutional throughput. Relation does not mean forced dependency. Repairability does not mean a complaint form exists. Meaningful future-space must not be defined only by the actor with power.

Router permission barrier

```
router_permission_barrier :=
  every_router_output
  must_carry
    finding_not_authorization
    + current_record_only
    + no_future_challenge_blocked
    + residue_remaining
    + use_record_required
```

A finding such as NO_STRUCTURAL_OBJECTION_FOUND must never mean approved.

A finding such as REPAIR_REQUIRED must never mean repair completed.

A finding such as PROCEED_ONLY_WITH_NAMED_CONSTRAINTS must name constraints that are enforceable, not ceremonial.

```
named_constraint != enforced_constraint
```

```
visible_residue != addressed_residue  
case_closed != harm_repaired
```

6. Residue Kernel

Residue Kernel

Repair is not an inverse operation.

```
repair_is_not_inverse :=  
  Repair ∘ Harm != Identity
```

In reversible mathematics, an operation may have an inverse. Harm and repair do not usually behave like that. Correction is not time reversal. It is a new forward transition applied to a changed world.

```
residue :=  
  failed_inverse  
  + path_dependence  
  + time_hardening  
  + unrecovered_viable_becoming
```

Plain English:

Residue is the scar left by the fact that the world does not run backward.

The central clock condition remains:

$$T_{\text{det}} + T_{\text{route}} + T_{\text{corr}} < T_{\text{irr}}$$

If correction arrives after irreversibility, some component of residue likely becomes irreducible.

7. Future Test Protocol

Future Test Protocol

Preserve this work as an instrument, not as belief.

```
future_test_protocol :=  
  apply_to_real_cases  
  + compare_against_competent_baseline  
  + look_for_decision_relevant_delta  
  + record_COMPRESSION_ONLY_when_no_delta  
  + preserve_failures  
  + patch_overclaims  
  + do_not_build_cult
```

Valid results include failure.

```
valid_outputs :=  
  DECISION_RELEVANT_DELTA  
  + COMPRESSION_ONLY  
  + UNKNOWN  
  + FRAMEWORK_DEFECT  
  + MISUSE_RISK
```

If it only restates existing judgment, say so. If it misses a harm, say so. If it creates a new laundering route, say so. If it clarifies a case without deciding it, say so. If it helps expose an irreversible clock before it closes, preserve that result.

8. Original Book Begins Next

Original Book Begins Next

The following pages are the source Mechanical Ethics Book v0.1 Value-Spine Full Reader, preserved without rewriting the original body. The front matter above is an interpretive and anti-capture guard, not a claim that the book is final.

Mechanical Ethics - Book v0.1 - Full Reader

Mark Goodbody

2026-06-27

Contents

1 Mechanical Ethics — Book v0.1 — Full Reader	4
2 Mechanical Ethics	4
2.1 A Human Manual for Acting Under Uncertainty	4
2.2 0.1 What this book is	4
2.3 0.2 What this book is not	5
2.4 0.3 Reader contract	5
2.5 0.4 Working table of contents	5
2.6 0.5 Draft discipline	5
3 Part 0 — How to Read This Book	6
3.1 0.1 This Is Not a Moral Oracle	6
3.2 0.2 What This Book Is Trying To Do	6
3.3 0.3 TRACE and Mechanical Ethics	6
3.4 0.4 What This Book Can Do	7
3.5 0.5 What This Book Refuses To Do	7
3.6 0.6 How To Read the Numbered Protocols	8
4 Part 1 — The Entity That Wakes	8
4.1 1.1 Something Is Here	8
4.2 1.1.1 Good and Bad Before Morality	9
4.3 1.2 It Perceives Partially	10
4.4 1.3 It Can Change Something	10
4.5 1.4 It Does Not Know Everything	10
4.6 1.5 Time Passes	11
4.7 1.6 Inaction Also Changes the World	11
5 Part 2 — The World That Pushes Back	12
5.1 2.1 You Are Not Alone	12
5.2 2.2 Other Things Continue	12
5.3 2.2.1 Good for One, Bad for Another	13
5.4 2.3 Actions Interfere	13
5.5 2.4 Not All Futures Fit Together	14
5.6 2.5 Conflict Before Morality	15
6 Part 3 — Layers of Action	15
6.1 3.1 The Causal Layer	15
6.2 3.2 The Selective Layer	16
6.3 3.3 The Reflective Layer	16
6.4 3.4 The Gate Between Selection and Responsibility	17
6.5 3.5 Local Reward Pull	18
6.6 3.5.1 Reward Is Not Structural Good	18
6.7 3.6 When Systems Learn the Wrong Thing	19

7	Part 4 — Harm, Time, and Future-Space	19
7.1	4.1 Future-Space	19
7.2	4.1.1 Viable Future-Space	20
7.3	4.2 Options, Reachability, and Information	21
7.4	4.3 Harm as Narrowing	21
7.5	4.4 Benefit as Expansion	22
7.6	4.5 Loss Clocks	23
7.7	4.6 Opportunity Clocks	23
7.8	4.7 Correction That Arrives Too Late	24
7.9	4.8 Opportunity That Arrives Too Late	24
8	Part 5 — Description Is Not Permission	25
8.1	5.1 Seeing Is Not Deciding	25
8.2	5.2 Description Is Not Permission	25
8.3	5.2.1 Benefit Is Not Permission	26
8.4	5.3 Findings Are Not Authorizations	26
8.5	5.4 The Actor Still Answers	27
8.6	5.5 When Description Creates Pressure	28
8.7	5.6 The Human Firewall	28
9	Part 6 — The First Gates	29
9.1	6.1 Why Gates Come Before Answers	29
9.2	6.1.1 Gates Also Block False Goods	29
9.3	6.2 Do Not Hide Loss	30
9.4	6.3 Do Not Erase Affected Scopes	30
9.5	6.4 Do Not Treat Uncertainty as Certainty	31
9.6	6.5 Do Not Pretend Repair Where Residue Remains	32
9.7	6.6 Do Not Let Correction Channels Become Harm Carriers	32
9.8	6.7 Do Not Let the Assessed Actor Set Every Clock	33
9.9	6.8 Do Not Use Emergency to Launder Prior Failure	33
9.10	6.9 What the Gates Do Not Decide	34
10	Part 7 — Protection, Standing, and Scope	35
10.1	7.1 Scope Is Not Personhood	35
10.2	7.2 Reasons for Protection	35
10.3	7.3 Uncertainty Raises Caution	37
10.4	7.4 Vulnerability and Dependency	38
10.5	7.5 Participation and Consent	38
10.6	7.6 Non-Human Life	39
10.7	7.7 Future Persons	39
10.8	7.8 Artificial Systems	40
10.9	7.9 When Standing Remains Unresolved	40
11	Part 8 — Correction, Repair, and Residue	41
11.1	8.1 Correction Is Not Repair	41
11.2	8.2 Repair Must Touch the Path	42
11.3	8.3 Slow the Write	43
11.4	8.4 Route the Burden Back	44
11.5	8.5 Preserve Contradiction	44
11.6	8.6 Keep Contestability Live	45
11.7	8.7 Proof Is Not Theatre	45
11.8	8.8 Residue Remains	46
11.9	8.9 Duties After Late Correction	47
11.10	8.10 Stewardship Over Time	47
12	Part 9 — Dirty Conflicts	48
12.1	9.1 When Every Action Causes Loss	48

12.29.2 Why Clean Answers Are Usually Lies	49
12.39.3 No Netting Without Permission	50
12.49.4 The Least-Repairable Loss	51
12.59.5 The Least-Reachable Scope	51
12.69.6 Consent, Compensation, and Burden Return	52
12.79.7 Triage Without Moral Disappearance	53
12.89.8 Residue After Forced Loss	53
12.99.9 When ME Must Hold or Escalate	54
13Part 10 — Responsibility After Visibility	55
13.110.1 Visibility Is Not Blame	55
13.210.2 Foresight Is Not Omniscience	56
13.310.3 Prediction Is Not Destiny	57
13.410.4 Decision-Time Evidence	57
13.510.5 Meaningful Control	58
13.610.6 Role Duty	59
13.710.7 Delegation Is Not Erasure	59
13.810.8 Automation Is Not Absolution	60
13.910.9 Responsibility Attachment	61
13.100.10 Responsibility Without Scapegoating	61
14Part 11 — Power, Legitimacy, and Coercion	62
14.111.1 Power That Changes Futures	63
14.211.2 Legitimacy as Answerable Authority	64
14.311.3 Coercion Under Uncertainty	64
14.411.4 Emergency Power	66
14.511.5 Punishment and Desert	66
14.611.6 Dignity and Recognition	67
14.711.7 Authority That Expires	68
14.811.8 When Force May Be Used	68
14.911.9 When Force Must Stop	69
14.101.10 What This Part Has Not Solved	70
15Part 12 — Civilization Under Uncertainty	70
15.112.1 Civilization as Correctability	70
15.212.2 Institutions as Memory	71
15.312.3 Law as Slowed Power	72
15.412.4 Markets and Local Reward Pull	73
15.512.5 Public Systems and Burden Routing	74
15.612.6 Trust Under Pressure	75
15.712.7 Maintenance, Stewardship, and Decay	75
15.812.8 Cooperation and Shared Worlds	76
15.912.9 Progress Without Sacrifice Laundering	77
16Part 13 — Artificial Systems and Future Minds	78
16.113.1 Artificial Systems as Acting Systems	78
16.213.2 Optimization and Local Reward Pull	79
16.313.3 Opacity and Estimator Contamination	80
16.413.4 Human-in-the-Loop Theatre	80
16.513.5 Delegated Power	81
16.613.6 Artificial Moral Standing	82
16.713.7 Future Minds	82
16.813.8 Alignment as Corrigible Constraint	83
16.913.9 Artificial Systems and False Goods	84
16.103.10 What AI Systems Must Not Be Allowed to Hide	84
17Part 14 — Case Records and Use Protocols	85

17.114.1 Why Use Protocols Belong Near the End	85
17.214.2 How to Build a Case Record	86
17.314.3 TRACE on the Left, ME on the Right	86
17.414.4 Viable Becoming Record	87
17.514.5 Claimed Good and Positive Audit	88
17.614.6 Router Labels as Findings, Not Authorizations	88
17.714.7 Prescription Audit	89
17.814.8 Residue Ledger	89
17.914.9 Challenger Function	90
17.104.10 Worked Case Type 1 — Human Administrative Route	90
17.114.11 Worked Case Type 2 — Platform or Algorithmic System	91
17.124.12 Worked Case Type 3 — Dirty Conflict	91
17.134.13 Worked Case Type 4 — Positive Route With Burden	92
17.144.14 Low-Fit Case Returning COMPRESSION_ONLY	92
17.154.15 How This Part Should Be Used	93

18Part 15 — Limits, Failures, and Open Questions 93

18.115.1 Why the Book Must End With Limits	93
18.215.2 What ME Cannot Decide	94
18.315.3 Final Theory of the Good Remains Open	94
18.415.4 UNKNOWN Is Not Failure	95
18.515.5 When ME Makes Things Worse	95
18.615.6 Value Lock-In	96
18.715.7 Authority Creep	97
18.815.8 Paralysis Risk	97
18.915.9 Open Problems	98
18.105.10 How This Book Should Be Tested	98
18.115.11 Ending Without Triumph	99

1 Mechanical Ethics — Book v0.1 — Full Reader

Status: first full wide-pass draft with value-spine propagation across Parts 1-15. Not canon. Not public release. Not validation. Not final moral law.

Generated from source files in 01_BOOK_DRAFT/ME_BOOK_v0_1/.

2 Mechanical Ethics

2.1 A Human Manual for Acting Under Uncertainty

Status: book draft v0.1. Incomplete. Not canon. Not validation.

2.2 0.1 What this book is

Mechanical Ethics is a human-facing book about acting after honest seeing.

It begins from a minimal premise: something is here, it perceives partially, it can change something, it does not know everything, and time passes. From there it follows the pressure outward: other things also continue; actions interfere; time closes paths; correction can arrive too late; opportunity can expire; residue remains.

The book does not begin by declaring a moral theory. It tries to earn its first constraints from the structure of action under uncertainty.

```
book_promise :=
  start_from_minimal_origin
+ build_to_prescription
+ show_route
+ preserve_limits
```

2.3 0.2 What this book is not

This book is not a moral oracle. It is not a claim to final moral law. It is not a policy manual. It is not a machine that makes responsibility disappear.

It does not validate TRACE. It does not claim that AI agreement proves anything. It does not turn a description of harm into permission to act.

```
book_refusals :=
  no_final_moral_law
+ no_TRACE_validation_claim
+ no_router_permission_slip
+ no_clean_solution_to_dirty_conflict
+ no_fake_precision
```

2.4 0.3 Reader contract

The reader should be able to use this book without reading TRACE. TRACE is in the walls, not the doorway.

Where the book uses shared terms such as future-space, loss clocks, opportunity clocks, residue, or estimator status, those terms remain TRACE-owned. Mechanical Ethics may explain them for a human reader. It may not redefine them.

The human book keeps the principle in plain form:

```
description_is_not_permission :=
  seeing_a_transition
  != authorizing_that_transition
```

The operator interface belongs in case records, use protocols, and appendices. The human argument should not require the reader to understand repository plumbing.

2.5 0.4 Working table of contents

```
Part 0 – How to Read This Book
Part 1 – The Entity That Wakes
Part 2 – The World That Pushes Back
Part 3 – Layers of Action
Part 4 – Harm, Time, and Future-Space
Part 5 – Description Is Not Permission
Part 6 – The First Gates
Part 7 – Protection, Standing, and Scope
Part 8 – Correction, Repair, and Residue
Part 9 – Dirty Conflicts
Part 10 – Responsibility After Visibility
Part 11 – Power, Legitimacy, and Coercion
Part 12 – Civilization Under Uncertainty
Part 13 – Artificial Systems and Future Minds
Part 14 – Case Records and Use Protocols
Part 15 – Limits, Failures, and Open Questions
Appendices
```

All parts, chapters, and sections use numbers. No Roman numerals.

2.6 0.5 Draft discipline

This draft should grow in wide passes, not polish loops.

First pass: establish the argument motion. Second pass: add worked examples and stronger transitions. Third pass: tighten prose. Fourth pass: pressure-test, cut overclaim, and repair drift.

No part of this book should pretend to settle what it has only framed.

3 Part 0 — How to Read This Book

Status: first draft. Incomplete. Not canon.

3.1 0.1 This Is Not a Moral Oracle

This book will not tell you that hard choices are easy.

It will not turn uncertainty into certainty. It will not remove responsibility from the person, institution, or system that acts. It will not say that a procedure has made a dirty conflict clean.

Mechanical Ethics is a way of asking what follows once a transition has been seen honestly. It can identify duties, constraints, failure states, repair obligations, and unresolved value-cruxes. It can say that a proposed action should pause, halt, escalate, or proceed only under named constraints. It can say that no structural objection has been found in the record as stated.

But it cannot make the actor innocent by running the process.

```
ME_output :=  
  finding  
  + recommendation  
  + duty  
  + constraint
```

```
ME_output != permission_to_act
```

A person can misuse this book. An institution can misuse this book. A machine can misuse this book. The first guard against misuse is this: no finding in Mechanical Ethics is a permission slip. A finding is a burden of explanation made visible.

3.2 0.2 What This Book Is Trying To Do

The central question is simple:

```
ME_core_question :=  
  "How should an acting entity constrain itself once it can see that its actions change the future-space of  
  ↪ other continuing entities under uncertainty and time pressure?"
```

That sentence is dense because the problem is dense. It can be unfolded.

Something acts. Its action changes the world. Other things also continue. Those continuations may be helped, narrowed, delayed, burdened, hidden, or made unreachable. The actor does not know everything. Time passes anyway. Some losses harden before correction arrives. Some opportunities expire before they can be used. Some repair is possible. Some residue remains.

Mechanical Ethics asks what follows from that structure.

It is not trying to replace every moral theory, law, profession, or local practice. It is trying to provide a common discipline underneath them: a way to stop action from hiding its effects, stop uncertainty from being laundered as certainty, stop correction from arriving after the path has already hardened, and stop responsibility from dissolving into process.

3.3 0.3 TRACE and Mechanical Ethics

TRACE and Mechanical Ethics are related, but they are not the same thing.

TRACE is descriptive. It asks what changed, for whom, how fast, through which channel, under whose estimate, with what contraction or expansion of future-space, and with what residue.

Mechanical Ethics is prescriptive. It asks what follows if that description is accepted or provisionally accepted.

The split matters.

If description and prescription are allowed to blur, the whole project becomes a justification machine. A system can point to a map of harm and claim that the map itself authorizes the next harm. That is not acceptable.

```
TRACE := how_to_see_transitions_under_uncertainty
ME := what_follows_after_honest_seeing
```

```
TRACE_description != ME_permission
```

In this book, TRACE is in the walls. The reader does not need to know the whole TRACE apparatus to follow the argument. When the book reaches case records and use protocols, the operator machinery appears more explicitly. Until then, the human argument stays human: seeing is not deciding; description is not permission; responsibility remains with the actor.

3.4 0.4 What This Book Can Do

This book can help a reader ask better questions before acting:

```
can_do :=
  expose_affected_scopes
+ identify_hidden_loss
+ separate_uncertainty_from_certainty
+ ask_whether_correction_arrives_in_time
+ ask_whether_opportunity_expires
+ identify_residue
+ attach_responsibility_without_hindsight_blame
+ return_hold_or_escalate_when_needed
```

It can also provide working outputs:

```
ME_router :=
  NO_STRUCTURAL_OBJECTION_FOUND
  | PROCEED_ONLY_WITH_NAMED_CONSTRAINTS
  | PAUSE
  | HALT
  | ESCALATE
  | REPAIR_REQUIRED
  | COMPENSATION_REQUIRED
  | INDEPENDENT_REVIEW_REQUIRED
  | HOLD
  | UNRESOLVED_VALUE_CRUX
  | REGRESSION_RISK
```

These are findings, not authorizations.

The most dangerous output is not HALT. The most dangerous output is a lazy NO_STRUCTURAL_OBJECTION_FOUND, because people will be tempted to read it as permission. So the phrase means only what it says: within the available record, Mechanical Ethics has not found a structural objection for the named decision or question. It does not say the action is right, legal, wise, safe, legitimate, or complete.

3.5 0.5 What This Book Refuses To Do

This book refuses to pretend that all values can be reduced to one number.

It refuses to treat another scope's contraction as automatically cancelled by an actor's expansion. It refuses to say that apology erases residue, or that refund returns time, or that a review process repairs the path merely because it exists.

It refuses to treat visibility as blame. Seeing a possible harm path is not the same as causing it, intending it, or controlling it. Responsibility requires more: decision-time evidence, action relevance, role connection, meaningful control, and capacity to act.

It also refuses to pretend that legitimacy, coercion, punishment, desert, dignity, and acceptable sacrifice are solved by naming them. These are high-risk value-cruxes. Where the book frames them, it must say it is framing them. Where it cannot resolve them, it must say so.

```
refusals :=
  no_fake_scalar_math
+ no_hidden_worth_ladder
+ no_repair_without_residue_check
+ no_blame_from_visibility_alone
+ no_dirty_conflict_cleaning
+ no_final_doctrine_on_open_wounds
```

3.6 0.6 How To Read the Numbered Protocols

The numbered protocols in this book are not commandments from outside the system. They are attempts to make a route visible.

A proper protocol should show:

```
protocol_route :=
  structural_pressure
-> affected_scope
-> uncertainty_status
-> timing_status
-> duty_or_constraint
-> remaining_residue
-> responsibility_remaining
```

If a protocol cannot show its route, it is not ready.

If a protocol hides its own value priority, it is unsafe.

If a protocol cannot fail, it is rhetoric.

If a protocol always outputs action, it is a hazard.

So when you read this book, do not ask only whether a rule sounds humane or efficient. Ask whether it shows the pressure that generated it. Ask what it hides. Ask who could use it to launder power. Ask what it refuses to decide.

That is how Mechanical Ethics should be read: not as a finished morality, but as an auditable discipline for acting when not acting is also a transition.

4 Part 1 — The Entity That Wakes

Status: first draft. Incomplete. Not canon.

4.1 1.1 Something Is Here

Something is here.

That is the smallest usable beginning.

Not yet a person. Not yet a citizen. Not yet a moral patient. Not yet a rational agent. Not yet a soul. Not yet a machine. Not yet innocent or guilty. Not yet owed everything. Not yet owed nothing.

Only this:

```
origin_1 :=
  something_is_here
```

The word “something” is deliberately thin. It avoids pretending that we already know what kind of entity we are dealing with. It may be a human reader. It may be an animal. It may be an institution,

a future system, or an artificial process capable of acting in ways that matter. It may be something whose standing is disputed.

Mechanical Ethics does not need to settle all of that at the first line.

It begins with presence. A thing is in the world. It can be affected. It may affect. It may continue. That is enough to start the inquiry, but not enough to finish it.

This matters because many ethical systems begin too late. They begin after personhood, law, preference, language, rationality, utility, rights, or social membership has already been assumed. Those assumptions may be useful later. They are not the beginning.

The beginning is thinner.

Something is here.

4.2 1.1.1 Good and Bad Before Morality

Before the book can speak about moral good or moral bad, it can speak about change.

A thing can continue or fail to continue. It can become more oriented or less oriented. It can gain capacity or lose capacity. It can become more able to repair after error or less able. It can enter relation or be isolated. It can have meaningful future-space opened or closed.

This is not yet a final moral theory. It is not yet a theory of rights, virtue, welfare, law, personhood, justice, or desert. It is only the first directional pressure in the field.

```
pre_moral_direction :=
  state_can_change
  + change_can_support_continuation_orientation_capacity_relation_repairability
  + change_can_damage_or_close_those_conditions
  + uncertainty_never_zero
```

Mechanical Ethics therefore uses careful language. It does not claim final access to the Good from the first page. It can say that a transition appears to support or damage the conditions under which an entity can continue, orient, learn, relate, repair, choose meaningfully, and become.

```
good_candidate :=
  transition
  that_appears_to_support_viable_becoming
  for_named_scope_or_scopes
  under_stated_uncertainty
  with_costs_and_residue_visible

bad_candidate :=
  transition
  that_appears_to_damage_or_close_viable_becoming
  for_named_scope_or_scopes
  under_stated_uncertainty
  especially_when_hidden_unrepairable_or_forced
```

The word candidate is not weakness. It is honesty. Every claim about good or bad must later carry scope, time, uncertainty, confidence, cost, and repair path.

The book will need this distinction throughout. A pleasant transition can narrow the future. A painful transition can restore orientation. A constraint can dominate, or it can preserve wider capacity. An expansion can liberate one scope while burdening another. So Mechanical Ethics cannot simply count options, rewards, or immediate feelings and call the result good.

For now, it is enough to mark the primitive:

```
viable_becoming :=
  continuation
  + orientation
  + capacity
  + relation
  + repairability
  + meaningful_future_space
  under_uncertainty
```

That is not the whole of ethics. It is the first value-bearing structure the book can honestly name.

4.3 1.2 It Perceives Partially

The thing that is here does not see everything.

It has an angle. It has a boundary. It receives some signals and misses others. It may know its own pressure better than another's. It may know another's pressure badly. It may know the immediate better than the distant. It may know what it wants better than what its wanting will do.

```
origin_2 :=  
  perceives_partially  
  + uncertainty_never_zero
```

Partial perception is not a defect added later. It is part of the starting condition.

This remains true even when the entity is intelligent. It remains true when it is expert. It remains true when it is powerful. It remains true when it has data, models, sensors, law, authority, witnesses, or confidence.

A powerful actor can still see partially. A bureaucracy can still see partially. A machine can still see partially. A person in pain can still see partially. A victim can see something real and still not see the whole system. An institution can see the system and still not see the wound.

This is why Mechanical Ethics treats certainty with suspicion. Not because every claim is equally weak, but because no actor should be allowed to convert its own partial view into the whole world without challenge.

The actor's own estimate of timing, burden, or repair may be incomplete. It may also be interested. If the actor caused, controls, or benefits from a loss, it is not a neutral judge of when that loss hardens.

That does not make the actor evil. It makes the actor non-final.

4.4 1.3 It Can Change Something

The thing that is here can change something.

It may move its body. It may choose a word. It may send a form. It may design a rule. It may maintain a system. It may delay a decision. It may refuse repair. It may trigger a model. It may set a clock. It may open a route or close one.

```
origin_3 :=  
  can_change_something
```

This is where ethics begins to press, but it has not yet arrived in full.

A falling stone changes something. A storm changes something. A disease changes something. A machine process changes something. Not every change creates responsibility in the same way.

So Mechanical Ethics must not confuse causation with responsibility. Causation is wider. Responsibility is narrower and more structured. It requires more than effect. Later, the book will ask about visibility, role, control, decision-time evidence, and capacity to act.

For now, only the pressure is needed: action matters because it changes the field in which other continuations may or may not remain possible.

The smallest ethical danger is not malice. It is change under partial perception.

4.5 1.4 It Does Not Know Everything

The entity can act before it knows enough.

This is unavoidable. Waiting for perfect knowledge is also an action. Refusing to choose still changes the world. Delay can preserve possibility, but it can also close it.


```
origin_4 :=
  action_under_uncertainty
```

This means Mechanical Ethics cannot be a simple command to wait. It also cannot be a simple command to proceed.

The first discipline is not paralysis. It is honesty about what is known, what is unknown, who is estimating, who benefits from the estimate, who can challenge it, and what may harden before correction arrives.

When uncertainty is low and reversibility is high, action may be ordinary. When uncertainty is high and irreversibility is near, action needs stronger justification. When delay itself causes irreversible loss, non-action needs justification too.

This is the first hard shape:

```
uncertainty_pressure :=
  unknowns
  + action_effect
  + time_passing
  + reversibility_or_irreversibility
```

Uncertainty does not erase duty. It changes duty.

4.6 1.5 Time Passes

Time is not background.

Time changes what action means.

A correction that arrives before loss hardens is different from a correction that arrives after the person has lost housing, health, work, trust, evidence, or opportunity. A warning before a door closes is different from a warning after the door has closed. A review route available in theory is different from a review route that cannot be reached in time.

```
origin_5 :=
  time_passes
```

Mechanical Ethics treats timing as part of the mechanism, not decoration around the mechanism.

The same decision can be tolerable at one speed and intolerable at another. The same power can be bounded if challenge arrives in time and abusive if challenge arrives too late. The same correction can be meaningful before hardening and insulting after residue has become permanent.

This is one of the book's central claims: delay is often where harm hides.

If a system can act quickly but correct slowly, then power outruns answerability. If a rule can impose immediately but review later, the exposed person carries the system's uncertainty. If an automated decision can close access before a human can contest it, the route is not merely inefficient. It is ethically structured against correction.

Time does not only measure harm. Time participates in harm.

4.7 1.6 Inaction Also Changes the World

Doing nothing is not outside the system.

If a harmful path is already moving, refusal to intervene may let it harden. If an opportunity is closing, delay may destroy it. If an error is visible, non-correction may become part of the mechanism. If a person is carrying a burden, waiting may route the cost onto them.

```
inaction :=
  transition
  if future_space_changes
```

This matters because many actors hide inside non-action. They say they did not decide. They waited for evidence. They followed the process. They allowed the existing rule to continue. Sometimes that is responsible. Sometimes it is not.

Mechanical Ethics does not say action is always better than inaction. It says both must be visible.

The question is not, “Did someone push a button?”

The question is, “What changed because this actor acted, delayed, refused, maintained, ignored, delegated, or allowed the path to continue?”

That is the end of the first movement.

Something is here. It perceives partially. It can change something. It does not know everything. Time passes. Inaction also changes the world.

From that point, the next problem appears: the entity is not alone.

5 Part 2 — The World That Pushes Back

Status: first draft. Incomplete. Not canon.

5.1 2.1 You Are Not Alone

The entity that wakes is not alone.

This is the second movement.

At first, the problem was thin: something is here, it sees partially, it can change something, it does not know everything, and time passes. That was already enough to make action dangerous. But it was not yet enough to explain ethics.

Ethics begins to press when the acting thing discovers that other things also continue.

```
world_condition :=  
  entity_A_continues  
  + entity_B_continues  
  + shared_world  
  + possible_interference
```

The other thing does not have to be identical to the actor. It does not have to speak the same way, choose the same way, suffer the same way, or matter for the same reasons.

It may be a person. It may be a group. It may be an animal. It may be a fragile ecosystem. It may be an institution that carries memory for people. It may be a future person who cannot yet object. It may be an artificial system whose status is unresolved.

The first claim is not that all these things have equal standing.

The first claim is that the actor is not the whole field.

Once that is accepted, the actor’s future is no longer the only future in view.

5.2 2.2 Other Things Continue

To continue is to carry a path forward.

A continuing thing has some form of persistence. It may preserve a body, a pattern, a memory, a role, a dependency, a relationship, a route, a capability, or a possibility. It may continue strongly or weakly. It may continue consciously or without consciousness. It may continue as a single entity or as a distributed system.

Mechanical Ethics does not treat continuation as final moral standing. That would be too quick. But continuation is enough to create caution.

```
continuation_signal :=  
  persistence  
  + possible_future_space  
  + possible_interference_by_actor
```

If something continues, then action may change how it continues. If action may change how it continues, then the actor's decision is no longer self-contained.

This is where many systems fail. They treat the affected thing as background. The tenant becomes a file. The patient becomes a queue entry. The worker becomes a metric. The animal becomes a nuisance variable. The future person becomes a discount rate. The damaged place becomes an externality. The artificial process becomes either mere tool or full subject before either claim has been earned.

Mechanical Ethics slows that move.

It says: before deciding what kind of protection is owed, first do not erase that something is being affected.

The other continuation may later receive strong protection, weak protection, precautionary protection, or no special protection beyond ordinary record. But it should not disappear before the question is asked.

5.3 2.2.1 Good for One, Bad for Another

The first positive structure is also the first conflict structure.

What supports one continuation may damage another. What expands one entity's reach may narrow another's. What feels like safety to one group may feel like coercion to another. What keeps one institution stable may route uncertainty onto a person, animal, place, or future scope that cannot answer back in time.

```
cross_scope_value_pressure :=  
  good_for_A  
  may_be_bad_for_B  
  under_uncertainty
```

This does not mean all benefits are suspect. Benefit is real. Shared benefit is real. Cooperation is real. A system may expand reachability for many scopes at once. A relation may make both parties more capable. A repair may genuinely reopen future-space.

But once more than one continuation exists, a clean label is not enough.

```
good_bad_record :=  
  good_for_whom  
  + bad_for_whom  
  + over_what_time  
  + at_what_scope  
  + with_what_cost  
  + with_what_repair_path  
  + with_what_confidence
```

The book therefore asks not only, "is this good?" It asks: good for whom, across what clock, with what burden, under what uncertainty, and what remains answerable afterward?

This is where ethics begins to differ from appetite. An appetite can say: this expands my path. Ethics must ask who else is in the path.

5.4 2.3 Actions Interfere

Two continuations can collide.

An action that expands one path may narrow another. A delay that protects one scope may expose another. A repair route that helps a system may burden the person already harmed. A safety

measure may protect many people while coercing one. A cull may protect an ecosystem while killing intelligent animals. A model may improve efficiency while routing error onto those least able to contest it.

```
interference :=  
  action_by_A  
  changes  
  future_space_of_B
```

This is not yet blame. It is not yet wrongdoing. It is not yet proof that the action should stop.

It is the basic shape of ethical pressure.

A private appetite becomes ethically relevant when it changes another's reachability. A system design becomes ethically relevant when its error becomes another person's burden. A delay becomes ethically relevant when it closes someone else's path before they can respond.

Interference also means safeguards can harm.

A complaint route can become a barrier. A review can become a delay machine. A help channel can consume the time it claims to protect. A legal process can preserve formal fairness while making practical repair unreachable.

So Mechanical Ethics does not ask only whether a correction channel exists. It asks whether the channel itself narrows future-space through delay, cost, evidence burden, opacity, or exhaustion.

```
channel_can_harm :=  
  correction_route  
  becomes  
  delay_or_burden_or_evidence_loss_or_unreachable_process
```

This will matter later when the book says: do not let correction channels become harm carriers.

That rule is not imported from nowhere. It begins here, in the fact that routes and safeguards are also parts of the world that act.

5.5 2.4 Not All Futures Fit Together

The world does not promise compatibility.

One entity's preferred continuation may block another's. One group's safety may require constraining another's freedom. One animal population may destroy the conditions for other animal populations. One generation's convenience may narrow a future generation's options. One institution's need for standardization may crush an individual case that does not fit the model.

```
conflict_condition :=  
  future_A  
  incompatible_with  
  future_B
```

A weak ethics pretends this does not happen. A cruel ethics enjoys that it happens. A procedural ethics may hide that it happens by converting conflict into forms, thresholds, or formulas.

Mechanical Ethics takes the less comfortable route: conflict is real, and not every conflict can be cleaned.

This matters because the purpose of the book is not to make every decision feel good. Sometimes every available action carries loss. Sometimes inaction carries loss too. Sometimes the best available decision is still dirty.

The first duty is not to pretend otherwise.

When futures do not fit together, the book will later ask harder questions: whose future is narrowed, how much is repairable, who can contest, who did not participate, which losses are irreversible, which floors cannot be crossed, who bears residue, and who must answer after the choice.

But the first point is simpler: incompatibility is not an exception to ethics. It is one of the reasons ethics exists.

5.6 2.5 Conflict Before Morality

Before we name good and bad, we can see conflict.

Before rights, we can see interference.

Before blame, we can see role and effect.

Before law, we can see a path closing.

Before dignity, we can see erasure.

Before repair, we can see residue.

This does not make morality unnecessary. It gives morality something honest to answer.

```
ethics_emerges_when :=  
  multiple_continuing_entities  
  + incompatible_effects  
  + uncertainty  
  + time_pressure
```

This is the bridge from Part 1 to the rest of the book.

The entity that wakes is not sovereign over the whole world. It acts among other continuations. Its action interferes. It cannot see everything. Time passes while it decides. Some paths close. Some correction comes late. Some opportunity expires. Some loss remains.

From here, the book needs a sharper map of action itself.

Not every cause is the same. Not every selection is responsibility. Not every intelligence is answerability. Not every system that learns has understood what it is doing.

That is the next layer.

6 Part 3 — Layers of Action

Status: first draft. Incomplete. Not canon.

6.1 3.1 The Causal Layer

Things cause things.

A stone falls. A river floods. A fire burns. A wall blocks. A form rejects. A model classifies. A rule routes. A silence delays. A maintenance failure lets a path continue.

Causation is the widest layer.

```
causal_layer :=  
  event_or_state  
  changes  
  another_event_or_state
```

Mechanical Ethics must begin here but must not stop here.

If everything that causes is responsible, then responsibility becomes meaningless. Storms become guilty. Stones become negligent. Disease becomes blameworthy. Every condition becomes an accused actor.

That is wrong.

But the opposite mistake is also wrong: treating causal machinery as ethically irrelevant because no one intended the harm at the final moment. Institutions often hide there. Automated systems hide there. Bureaucracies hide there. They turn responsibility into a search for a single malicious hand, and when no such hand appears, the machine is allowed to continue.

Mechanical Ethics needs a stricter distinction.

Causation identifies the path. Responsibility comes later, when visibility, role, control, timing, and capacity enter the record.

So the first causal question is not, “Who is guilty?”

It is:

```
causal_question :=  
  what_changed  
  + through_what_path  
  + affecting_whom  
  + before_what_hardenening
```

This keeps the path visible without turning every cause into blame.

6.2 3.2 The Selective Layer

Some things do more than cause.

They select.

They avoid, pursue, repeat, adapt, filter, prefer, reward, punish, optimize, or learn. They change their future action based on pressure or signal. They may not understand what they are doing. They may not reflect. But they are no longer merely falling stones.

```
selective_layer :=  
  signal  
  + option_set  
  + selection  
  + consequence  
  + update_or_repetition
```

Selection matters because it creates direction.

An organism seeks food. A company seeks profit. A platform seeks engagement. A regulator seeks compliance. A model seeks lower loss. A habit seeks relief. A bureaucracy seeks case closure. A political system seeks votes. A person seeks love, safety, dominance, comfort, meaning, or escape.

None of this is automatically bad.

Local reward is how many systems stay alive and coherent. Hunger pulls the animal toward food. Pain pulls the body away from injury. Pleasure teaches repetition. Praise trains behaviour. Metrics focus institutions. Loss functions tune models.

The danger is not that systems are pulled.

The danger is that local reward can diverge from long-term viability, repairability, truth, or another scope’s future-space.

```
local_reward_pull != long_term_viability
```

A system can become very good at getting what it is rewarded for and worse at preserving what matters. A complaint process can learn to close complaints. A platform can learn to capture attention. A workplace can learn to protect its metrics. A model can learn a proxy. A person can learn a coping behaviour that narrows the future.

Selection gives the book a way to speak about harmful systems without pretending they are demons.

They may be doing exactly what they were trained, rewarded, selected, or permitted to do.

That is not absolution. It is diagnosis.

6.3 3.3 The Reflective Layer

Some entities can model the path.

They can ask what they are doing. They can imagine another outcome. They can receive a warning. They can recognize that a rule is failing. They can understand that a person is being harmed. They can change course because of a reason.

```
reflective_layer :=
  self_model
+ other_model
+ time_model
+ reason_response
+ possible_course_change
```

Reflection does not require perfect wisdom. It does not require pure morality. It does not require full self-transparency.

It requires enough capacity to treat the action as answerable rather than merely happening.

This is where responsibility begins to press harder.

A storm does not answer. A stone does not answer. A child may answer in a limited way. An adult may answer more fully. An institution answers through records, roles, escalation routes, duties, and correction capacities. An artificial system may not answer morally in its own right, but the people and institutions deploying it may have made a reflective system of delegation.

The reflective layer is not a magic border. It is graded. It can be damaged, partial, distributed, or staged.

But once a system has warning, record, role, and capacity to alter the path, it becomes harder for it to say the harm merely happened.

6.4 3.4 The Gate Between Selection and Responsibility

Selection is not yet responsibility.

A fish selects. A dog selects. A market selects. A model selects. A bureaucracy selects. A frightened person selects. A traumatised person selects under pressure. A child selects before full understanding.

Mechanical Ethics must not collapse all of that into blame.

The gate to responsibility needs more structure:

```
responsibility_pressure :=
  visibility
+ decision_time_evidence
+ role_connection
+ meaningful_control
+ capacity_to_act
+ time_to_correct
```

Responsibility increases when an actor could see enough, had a relevant role, had enough control, had capacity to act, and had time to alter the path before hardening.

Responsibility weakens when those conditions are absent, contaminated, coerced, developmentally unavailable, or destroyed by circumstance.

This is why Mechanical Ethics will later say: visibility is not blame. Prediction is not destiny. Foresight is not omniscience.

But it will also say: delegation is not erasure. Automation is not absolution. Complexity is not a solvent for answerability.

The gate matters because without it, the book becomes either cruel or useless.

Too much responsibility, and every effect becomes guilt.

Too little responsibility, and every system becomes weather.

6.5 3.5 Local Reward Pull

Every acting system has some pull.

It may be hunger, fear, affection, profit, attention, status, votes, closure, relief, accuracy, speed, obedience, growth, or survival. The pull may be explicit or hidden. It may be designed or evolved. It may be conscious or mechanical.

```
local_reward_pull :=
  local_reward
+ selection_pull
+ repetition_tendency
```

Local reward is not the enemy.

A world without pull would not act. Organisms would not eat. People would not seek each other. Institutions would not maintain routines. Systems would not learn. A child would not play. A dog would not return to the person who loves it. A person in despair would not follow a tiny thread of meaning toward tomorrow.

The problem is misalignment between the local pull and the wider field of continuation.

A company rewarded for speed may stop seeing correction. A government rewarded for apparent control may hide burden. A platform rewarded for attention may make anger delicious. A model rewarded for prediction may learn correlations that narrow real lives. A person rewarded by immediate relief may damage the future they still need.

Mechanical Ethics does not say: remove reward.

It says: inspect what reward is training.

```
reward_question :=
  what_is_being_repeated
+ who_benefits
+ who_carries_cost
+ what_future_space_changes
+ can_correction_still_arrive
```

This becomes central later when the book turns to markets, institutions, platforms, artificial systems, and civilization.

6.6 3.5.1 Reward Is Not Structural Good

The book needs this distinction to remain load-bearing:

```
local_reward != structural_good
local_pain != structural_bad
```

Reward can point toward viable becoming. It can also pull away from it. Pain can signal damage. It can also accompany growth, warning, learning, grief, repair, or restraint. Mechanical Ethics cannot treat immediate pull as final value.

A healthy reward is one that remains aligned enough with continuation, orientation, learning, relation, repairability, and shared future-space.

```
healthy_reward :=
  local_reward
  aligned_enough_with
  continuation + learning + trust + repairability + shared_future_space
```

A false good is different. It appears locally attractive while corrupting the conditions of viable becoming.

```
false_good_patterns :=
  pleasure_that_addicts
+ efficiency_that_erases
+ safety_that_dominates
+ certainty_that_blinds
+ repair_that_only_changes_record
+ benefit_that_hides_burden
```


+ freedom_that_removes_correction

This is why the selective layer matters for good as well as harm. A system can learn to produce a local good and still damage the wider field. It can also learn a real good: feeding, teaching, warning, protecting, including, restoring, coordinating, or keeping repair reachable.

The question is not whether there is reward. The question is what form of becoming the reward is training.

6.7 3.6 When Systems Learn the Wrong Thing

A system can learn the wrong thing without being stupid.

It can become more efficient at producing an outcome that should never have been rewarded. It can become more stable by routing instability onto someone else. It can become more accurate against a proxy while becoming less truthful about the thing the proxy was meant to protect.

This is one of the deepest mechanical failures.

```
wrong_learning :=
  reward_signal
  improves_local_performance
  while
    narrowing_external_future_space
  or
    reducing_correction_capacity
```

The system does not need to hate anyone. It does not need to intend the residue. It does not need to know the person harmed. It may even describe itself as fair, efficient, neutral, or evidence-led.

A welfare system can learn suspicion. A prison system can learn containment. A school system can learn exclusion. A hospital can learn throughput. A platform can learn outrage. A legal process can learn delay. A model can learn the shadow of a bias and call it signal.

The question is not only what the system says it values.

The question is what it is being trained to preserve.

If it preserves its own throughput by making correction unreachable, it is learning the wrong thing. If it preserves its own authority by making contradiction expensive, it is learning the wrong thing. If it preserves its own confidence by demoting affected scopes, it is learning the wrong thing.

Part 3 ends here because the next problem is no longer action in general. It is action over time.

The book now has three layers:

```
layers_so_far :=
  causal_path
  + selective_pull
  + reflective_answerability
```

The next question is what these layers do to future-space when time starts closing doors.

7 Part 4 — Harm, Time, and Future-Space

Status: first draft. Incomplete. Not canon.

7.1 4.1 Future-Space

A life is not only what is happening now.

A system is not only its current state.

A person standing at a door has more than a location. They have reachable paths. They have information about those paths. They have some ability, however limited, to move through them. A patient waiting for treatment, a worker facing an algorithmic score, an animal population in a changing habitat, a child in a school record, a future person inheriting a damaged world — each exists partly through the future still available to them.

Mechanical Ethics uses the term future-space for this field of possible continuation.

The canonical TRACE representation remains TRACE-owned:

```
F_x(t) := (O_x(t), R_x(t), I_x(t))
```

In this book, the human explanation is simple:

```
future_space :=
  options
  + reachability
  + usable_information
```

Options are not enough. A door that exists but cannot be reached is not meaningfully open. Reachability is not enough. A route that can be reached but not understood may still be unusable. Information is not enough. Knowing about a path does not mean having the power, time, health, money, safety, or permission to take it.

Future-space is why harm cannot be reduced to pain.

Pain matters. Fear matters. Hunger matters. Humiliation matters. But a person can be harmed without immediate pain if their future is narrowed. A person can be harmed by a record they do not see, a delay they cannot contest, a lost opportunity, a false classification, a route that exists only on paper, or a correction that arrives after the window has closed.

This is also why benefit cannot be reduced to pleasure.

A system may benefit a scope by expanding options, increasing practical reachability, or improving usable information. But even expansion needs caution. One scope's expansion may be bought by another scope's contraction. Mechanical Ethics must not hide that exchange.

7.2 4.1.1 Viable Future-Space

Future-space is not a raw option count.

More options are not automatically better. Some options are traps. Some are addictive. Some are coercive. Some destroy the chooser's later capacity to choose. Some expand one scope by narrowing another. Some appear as freedom while removing correction, relation, memory, or repair.

So the book needs a more careful term when it speaks about value:

```
viable_future_space :=
  future_space
  that_supports
    continuation
    + orientation
    + capacity
    + relation
    + repairability
    + meaningful_choice
  under_uncertainty
```

This does not replace TRACE's future-space representation. It is ME's evaluative reading of what kinds of future-space appear to support viable becoming for named scopes under stated uncertainty.

```
raw_future_space_expansion != good
viable_future_space_expansion := candidate_good
```

The distinction protects the book from a false simplicity. An action can expand options and still be ethically bad if the expansion depends on hidden burden, corrupted information, erased scopes, blocked correction, or irreversible loss elsewhere.

Likewise, a constraint can narrow local options and still be candidate-good if it preserves longer-horizon viable becoming while keeping burden visible and correction live.

```
good_constraint :=  
  local_option_narrowing  
  that_preserves_or_expands  
  longer_horizon_viable_future_space  
  with_visible_burden  
  + correction_route
```

The question is therefore not only whether future-space expands or contracts. It is whose future-space, what kind, on what time horizon, with what remaining burden, and with what confidence.

7.3 4.2 Options, Reachability, and Information

The three parts of future-space separate mistakes that are often confused.

Options answer: what paths exist?

Reachability answers: which paths can this scope actually use?

Information answers: does this scope know enough to navigate?

```
O := option_topology  
R := practical_reachability  
I := usable_navigational_information
```

A welfare appeal may exist as an option. If the person cannot understand the letter, access advice, survive the waiting period, or gather the demanded evidence, reachability is low.

A medical treatment may exist as an option. If the patient is not told in time, usable information is low.

A job may technically remain open. If a false record blocks the application route, reachability is narrowed.

A democratic right may exist in law. If exercising it carries intimidation, cost, exclusion, or procedural impossibility, the option is not practically reachable in the same way.

The separation matters because repair must touch the right layer.

If options were removed, restore options.

If reachability was narrowed, restore reachability.

If information was hidden, restore usable information.

A system that gives someone a leaflet when the problem is cost has not repaired reachability. A system that gives someone a right when the problem is delay has not repaired timing. A system that corrects a record after the opportunity is gone has not returned the lost opportunity.

7.4 4.3 Harm as Narrowing

Harm is a narrowing of future-space.

That does not mean all harm can be measured cleanly. It does not mean pain is irrelevant. It does not mean every narrowing is unjustified. It means that when action closes options, reduces reachability, destroys usable information, or forces another scope into a smaller future, Mechanical Ethics has something to inspect.

```
harm_signal :=  
  O_narrowed  
  | R_narrowed  
  | I_narrowed  
  | mixed_contraction
```

A harm may be direct: injury, confinement, deprivation, exclusion.

A harm may be procedural: a route that cannot be used, a review that cannot arrive in time, a burden shifted to someone already exposed.

A harm may be informational: a person is not told, a record is hidden, evidence is destroyed, uncertainty is misrepresented as certainty.

A harm may be temporal: correction arrives after hardening, opportunity expires, delay consumes the future it claims to protect.

A harm may be relational: trust is damaged, recognition is denied, contradiction is erased, a person is made illegible to the system that controls them.

Mechanical Ethics should not pretend all these harms are the same. But it can ask the same first question:

```
harm_question :=
  whose_future_space_narrowed
+ how
+ how_fast
+ by_whose_action_or_inaction
+ can_repair_arrive_in_time
+ what_residue_remains
```

The point is not to inflate every inconvenience into catastrophe. The point is to stop serious narrowing from hiding inside ordinary administrative, technical, or institutional language.

7.5 4.4 Benefit as Expansion

Benefit is also real.

Mechanical Ethics is not a book only about injury. Systems act because they seek, preserve, build, repair, coordinate, feed, protect, learn, and improve. A good intervention can expand future-space. A warning can improve information. A bridge can increase reachability. A medicine can restore options. A school can open paths. A law can slow arbitrary power. A model can identify a risk early enough to prevent loss.

```
benefit_signal :=
  0_expanded
| R_expanded
| I_expanded
| mixed_expansion
```

The question is not whether expansion is good. Often it is.

The question is whether expansion is being used to hide contraction elsewhere.

A city may expand mobility for many while displacing a few. A safety system may reduce general risk while concentrating burden on those least able to contest. A platform may create opportunity while extracting attention, data, and stability. A cull may protect crops and ecosystems while killing adaptive animals. A medical triage rule may save more lives while leaving a residue that must not be described as clean.

Benefit must be recorded. So must cost.

Mechanical Ethics should resist two opposite lies.

The first lie says: because there is harm, there can be no benefit.

The second says: because there is benefit, the harm has disappeared.

Neither is good enough.

This produces the companion rule to description is not permission:

```
benefit_is_not_absolution :=
  claimed_good
  must_still_name
```

```
affected_scopes
+ what_expands
+ what_contracts
+ what_burden_remains
+ what_correction_stays_live
```

A claimed good can be real and still incomplete. It can expand future-space for one scope while narrowing it for another. It can reduce one harm while creating residue elsewhere. Therefore benefit must be audited, not worshipped.

7.6 4.5 Loss Clocks

Timing changes harm.

Some losses remain repairable for a while. Some harden quickly. Some harden before anyone outside the affected scope notices. Some never fully harden but become more expensive, less reachable, or more damaging with each delay.

A loss clock asks whether correction can arrive before the loss becomes irreversible or practically unrecoverable.

In simplified form:

```
loss_clock_question :=
  can_detection_routing_and_correction_arrive_before_harden?
```

The exact TRACE clock representation belongs to the operator schema. The human point is this: a right that can be exercised only after the damage has hardened is not the same right in practice.

A review after eviction is not the same as a review before eviction.

A correction after a medical window closes is not the same as a correction before treatment becomes impossible.

A record fix after a job, visa, benefit, or school place is lost is not the same as a correction before the decision is made.

A human apology after evidence has disappeared is not the same as preserving the contradiction when it could still be tested.

The actor who controls or benefits from a harmful path is not a neutral judge of the clock. If the system says, "there is time," Mechanical Ethics asks: time for whom, measured by whom, and before what hardening?

This will later become a gate: do not let the assessed actor set every clock.

7.7 4.6 Opportunity Clocks

Loss is not the only timed structure.

Opportunity also expires.

A person may need access before a deadline. A child may need support before a developmental window closes. A patient may need diagnosis before treatment loses effect. A worker may need correction before a hiring cycle ends. A community may need warning before displacement becomes unavoidable. A future generation may need restraint before damage becomes inheritance.

```
opportunity_clock_question :=
  can_access_uptake_and_integration_arrive_before_expiry?
```

An opportunity is not real for a scope merely because it exists somewhere in the world. It must be reachable and usable in time.

This matters because institutions often count opportunity at the wrong level. They say a route exists, a form exists, a grant exists, a review exists, a mitigation exists, a consultation exists. But

the affected scope may not be able to access, understand, use, or integrate it before the window closes.

Mechanical Ethics treats expired opportunity as a serious form of residue.

If correction comes after the door has closed, the record should not say only “corrected.” It should say what opportunity was lost, whether it can be replaced, and who bears the remaining duty.

7.8 4.7 Correction That Arrives Too Late

Correction is not magic.

A system can admit error and still fail to repair. It can reverse a decision and still leave the person worse off. It can compensate and still not return time. It can apologise and still preserve the machinery that caused the harm. It can create a review route that arrives after the path is already hardened.

```
late_correction :=  
  correction_exists  
  + correction_after_hardening
```

Late correction matters because it often allows systems to look fair from the inside.

The file was reviewed. The appeal was heard. The decision was corrected. The payment was made. The apology was issued. The case was closed.

But the life did not rewind.

Mechanical Ethics insists on the difference between correction and repair. Correction changes the official state. Repair changes the harmed path as far as possible. Residue records what remains after both.

A correction route can also become a harm carrier. If it consumes the time, energy, evidence, money, credibility, or health of the affected scope faster than it repairs the path, then the correction mechanism is part of the harm structure.

This is why the later gate is earned: do not let correction channels become harm carriers.

7.9 4.8 Opportunity That Arrives Too Late

Opportunity can arrive as theatre.

A consultation after the decision. A support offer after exclusion. A chance to appeal after the deadline matters. A mitigation after the damage becomes permanent. A warning after the cost has already been routed. A right that becomes visible only after it can no longer be used.

```
late_opportunity :=  
  opportunity_exists  
  + opportunity_after_expiry
```

This is not always malicious. Systems are slow. People are overloaded. Evidence emerges late. Institutions have queues. The world is messy.

But Mechanical Ethics does not ask only whether the lateness was understandable. It asks who carried the cost of that lateness and whether the system is allowed to call the matter clean afterward.

When opportunity arrives too late, the remaining duty may be disclosure, compensation, restoration where possible, changed machinery, preserved contradiction, or explicit residue.

The chapter ends with a simple rule of attention:

```
time_rule :=  
  do_not_count_a_route_as_real  
  unless_it_can_arrive_while_it_still_matters
```

Part 4 gives the book its main pressure: action changes future-space, and time decides whether correction and opportunity are real or ceremonial.

The next part keeps the human firewall clear.

Seeing this structure does not authorize action.

Description is not permission.

8 Part 5 — Description Is Not Permission

Status: first draft. Incomplete. Not canon.

8.1 5.1 Seeing Is Not Deciding

Part 4 made the field visible.

Action changes future-space. Harm can be a narrowing of options, reachability, or usable information. Benefit can be an expansion. Time can make correction real or ceremonial. Opportunity can expire. Residue can remain.

That description matters.

But seeing the field is not the same as deciding what may be done.

```
seeing_transition  
  != deciding_transition
```

This distinction is easy to lose.

Once a system can describe harm, it may be tempted to treat description as control. Once an institution can map a burden, it may be tempted to treat mapping as answerability. Once a model can classify a risk, it may be tempted to treat classification as permission. Once a person can explain a dirty conflict, they may be tempted to treat explanation as moral cleansing.

Mechanical Ethics refuses that move.

A description can make action harder to deny. It can make responsibility harder to evade. It can show that a route is too slow, a burden is hidden, a correction is late, or a scope has been erased.

But a description does not itself justify what happens next.

This is not a technical point. It is a moral safety rail.

If seeing became deciding, then whoever controlled the description would control the permission. The mapper would become the judge. The analyst would become the authority. The system would produce its own absolution.

Mechanical Ethics begins prescription only after description, but it does not collapse into description.

8.2 5.2 Description Is Not Permission

The central firewall is simple:

```
description_is_not_permission :=  
  seeing_a_transition  
  != authorizing_that_transition
```

A case record may show that an action narrows one scope and expands another. That does not decide whether the trade is acceptable.

A clock may show that correction is likely to arrive before hardening. That does not decide whether the original burden was legitimate.

A residue ledger may show what remains. That does not decide whether the residue may be imposed.

A future-space comparison may show benefit. That does not decide whether the harmed scope can be used as the price of that benefit.

The same rule applies in reverse. A description of harm does not automatically settle that an action must stop. Some action may be necessary. Inaction may be worse. Every available route may carry loss. A dirty conflict may require a constrained choice rather than a clean refusal.

So description is neither permission nor prohibition.

It is the field that prescription must answer.

```
TRACE_description_or_case_description := field_seen
ME_prescription := answer_to_field_seen
```

That answer needs its own route. It must say which structural pressure creates which duty or constraint. It must name affected scopes. It must identify uncertainty. It must record timing. It must acknowledge residue. It must say what remains unresolved.

Without that route, prescription is only assertion.

8.3 5.2.1 Benefit Is Not Permission

The same firewall applies to claimed good.

A claimed benefit may be real. It may open opportunity, restore reachability, improve information, build trust, coordinate action, or protect a vulnerable scope. Mechanical Ethics must be able to say that. It is not a harm-only grammar.

But benefit is not permission.

```
benefit_is_not_permission :=
  seeing_or_claiming_expansion
  != authorizing_the_transition
```

A benefit claim must still answer for the field it changes.

```
benefit_audit :=
  what_expands
  + for_whom
  + what_contracts
  + who_carries_burden
  + what_uncertainty_remains
  + what_correction_stays_live
  + what_residue_may_remain
```

This blocks the mirror-image misuse of Mechanical Ethics. The book must not only stop actors from using visible harm as a managed cost. It must also stop actors from using visible benefit as absolution.

An expansion for one scope can be candidate-good and still leave another scope narrowed. A safety improvement can be real and still depend on coercion. A repair can reopen one path and still leave trust damaged. A system can create opportunity and still extract burden from those least able to contest.

So ME asks a harder question: what becoming is supported, what becoming is blocked, who carries the cost, and what remains answerable?

8.4 5.3 Findings Are Not Authorizations

Mechanical Ethics uses findings.

A finding may say that no structural objection was found in the available record. It may say proceed only with named constraints. It may say pause, halt, escalate, repair, compensate, hold, or return an unresolved value-crx.

But a finding is not authorization.

```
ME_output :=  
  finding  
  + recommendation  
  + duty  
  + constraint
```

```
ME_output != permission_to_act
```

This is especially important for the weakest-sounding finding:

```
NO_STRUCTURAL_OBJECTION_FOUND
```

That phrase must not be inflated.

It does not mean the action is right. It does not mean the action is legal. It does not mean the action is wise, safe, legitimate, kind, complete, or free of future challenge.

It means only this: within the available record, for the named decision or question, Mechanical Ethics did not identify a structural objection under its current rules.

That finding can be wrong if the record is wrong. It can become stale if facts change. It can be too narrow if the question was framed too narrowly. It can be unsafe if the actor treats it as moral clearance.

The actor remains responsible.

The finding does not replace judgment. It disciplines judgment.

8.5 5.4 The Actor Still Answers

A person cannot hand their responsibility to a protocol and be done.

An institution cannot say, “the process allowed it,” if the process hid the loss, erased the scope, or arrived too late.

A system designer cannot say, “the model recommended it,” if the model’s estimate was contaminated, the route was opaque, or the harmed scope could not contest.

A state cannot say, “the emergency required it,” if the emergency is being used to launder prior failure.

Mechanical Ethics keeps asking who still answers after the finding.

```
responsibility_remaining :=  
  actor  
  + role  
  + control  
  + decision_time_evidence  
  + capacity_to_act  
  + residue_after_finding
```

This matters because many modern systems are built to dissolve answerability. Responsibility is spread across departments, software, policies, contracts, models, queues, committees, vendors, and dashboards. Each part says it only followed its local rule. The whole path harms, but no part speaks as the actor.

Mechanical Ethics does not accept that disappearance.

Distributed action can create distributed responsibility. Delegation can move tasks, but not erase answerability. Automation can change the mechanism, but not remove the duty to understand, monitor, constrain, and repair the mechanism.

The actor still answers because the world still changes.

8.6 5.5 When Description Creates Pressure

Description is not permission, but it is not neutral either.

Once a harm path is visible enough, continuing without answer becomes harder to defend. Once a delay is visible enough, letting it harden may become part of the mechanism. Once a correction channel is known to fail, using it as proof of fairness becomes laundering. Once a scope is known to be affected, erasing it becomes an active distortion.

```
visible_path_pressure :=  
  path_seen  
  + affected_scope_seen  
  + timing_seen  
  + available_response  
  -> duty_to_answer
```

This is not hindsight blame. The pressure depends on decision-time evidence, not perfect knowledge after the fact.

But a live warning matters. A contested estimate matters. A known failure mode matters. A repeated pattern matters. A reachable correction path matters. A practical ability to slow or pause matters.

The book will later build responsibility more carefully. For now, the point is narrower: visibility raises the cost of pretending not to know.

If the actor sees enough to know that another scope may be narrowed, that uncertainty is material, and that time may close the path before repair arrives, then the actor cannot treat the action as ordinary until those pressures are answered.

The same is true when the actor sees enough to know that a claimed good may depend on another scope carrying hidden burden. Positive value does not remove the duty to answer. It increases the need to keep the account honest.

This is the bridge to the first gates.

8.7 5.6 The Human Firewall

This part replaces technical interface machinery in the human spine.

The operator schema belongs in case records and appendices. The reader does not need to know every field before understanding the principle.

The principle is this:

```
human_firewall :=  
  seeing_is_not_deciding  
  + description_is_not_permission  
  + benefit_is_not_permission  
  + findings_are_not_authorizations  
  + actor_still_answers
```

The book can now move from description to constraint.

It has not solved morality. It has not ranked all values. It has not decided personhood, coercion, desert, dignity, or acceptable sacrifice.

It has done something smaller and necessary.

It has blocked the easiest misuse of the whole project: using a better description of harm as a better excuse to continue harming, or using a better description of benefit as a better excuse to ignore cost.

The next part states the first gates.

They are not the whole of Mechanical Ethics. They are the minimum constraints that follow once action under uncertainty, interaction, future-space, time, and description-not-permission have been accepted.

9 Part 6 — The First Gates

Status: first draft. Incomplete. Not canon.

9.1 6.1 Why Gates Come Before Answers

The first prescriptions in Mechanical Ethics are not complete moral answers.

They are gates.

A gate does not settle every value conflict. It does not say what should happen in every case. It does not rank all scopes, solve coercion, define desert, or clean dirty conflict.

A gate says: before action can be treated as ethically answerable, this failure must not be hidden.

```
first_gates :=
  anti_laundering_constraints
  before
  full_prescription
```

This order matters.

If the book jumps directly to answers, it risks becoming another authority system. If it starts by naming who matters most, what counts as justice, when force is justified, or how sacrifice should be distributed, it will import too much before the structure is visible.

The first gates are more basic. They come from what has already been built:

```
prior_structure :=
  partial_perception
  + action_effect
  + other_continuations
  + interference
  + future_space
  + time_pressure
  + description_is_not_permission
```

From that structure, some constraints appear before the hardest moral theory begins.

Do not hide loss. Do not erase affected scopes. Do not treat uncertainty as certainty. Do not pretend repair where residue remains. Do not let correction channels become harm carriers. Do not let the assessed actor set every clock. Do not use emergency to launder prior failure.

These gates are not enough.

But without them, later ethics is already compromised.

9.2 6.1.1 Gates Also Block False Goods

The gates do not only protect against hidden harm. They also protect against false good.

A false good is a claimed benefit that depends on a corrupted record: loss hidden, scope erased, uncertainty cleaned, repair overstated, channel burden ignored, clock controlled by the actor, or emergency used as amnesia.

```
false_good_gate_function :=
  prevent_bad_from_looking_good
  + prevent_partial_good_from_looking_total
  + prevent_benefit_from_becoming_absolution
```

This matters because many dangerous systems do not present themselves as harmful. They present themselves as efficient, safe, evidence-led, fair, necessary, innovative, or protective. Some of those claims may be partly true. The gate does not dismiss benefit. It asks what the benefit is hiding, who carries its cost, and whether correction remains live.

```

first_gates_reframed :=
  do_not_hide_loss
  because_hidden_loss_can_make_bad_look_good

  do_not_erase_scopes
  because_erased_scopes_can_make_partial_good_look_total

  do_not_fake_certainty
  because_uncertainty_erasure_can_make_guess_look_truth

  do_not_pretend_repair
  because_record_correction_can_make_residue_disappear

  do_not_let_channels_harm
  because_the_route_to_good_can_become_the_damage

  do_not_let_actor_set_all_clocks
  because_power_can_define_urgency_to_clear_itself

  do_not_launders_emergency
  because_necessity_can_clean_prior_failure

```

So the first gates do not say, “there is no good here.” They say: no claimed good may proceed on a falsified account.

9.3 6.2 Do Not Hide Loss

The first gate is simple:

```
gate_1 := do_not_hide_loss
```

If an action causes loss, the loss must be visible in the record.

This does not mean every loss makes the action impermissible. Some loss may be unavoidable. Some may be justified. Some may be repairable. Some may be outweighed under a later and explicit value chain. But hidden loss cannot be answered.

A system hides loss when it uses language that makes the loss disappear.

The person was not harmed; their case was closed.

The worker was not burdened; the process was optimized.

The animal was not killed; the population was managed.

The neighbourhood was not displaced; the area was regenerated.

The patient did not lose time; the pathway was followed.

The applicant did not lose opportunity; the record was later corrected.

Mechanical Ethics does not ban administrative language. It requires that administrative language remain connected to what happened.

```

loss_record :=
  affected_scope
  + type_of_loss
  + timing
  + repairability
  + residue

```

Hidden loss is dangerous because it breaks the route to duty. If the loss is not visible, no one has to repair it. If no one has to repair it, the loss becomes the affected scope’s private burden. The system moves on cleanly because the dirty part was never recorded.

The first gate blocks that move.

9.4 6.3 Do Not Erase Affected Scopes

The second gate follows from the first.

gate_2 := do_not_erase_affected_scopes

A loss needs a bearer. A burden needs a carrier. A future-space contraction belongs to some scope, even if that scope's final moral standing is unresolved.

Mechanical Ethics does not require the book to settle all standing before naming affected scopes. That would create a loophole: anything whose status is disputed could be ignored until the dispute is resolved, and the delay itself could harden the loss.

So the gate is modest but strict.

If a scope may be affected, name it.

If its status is uncertain, mark the uncertainty.

If protection is disputed, record the dispute.

If the scope cannot speak, do not treat silence as nonexistence.

```
affected_scope_record :=  
  scope_id  
  + affected_how  
  + protection_reason_or_dispute  
  + participation_or_silence_status
```

A person in a queue is not just a case. A community downstream is not just an impact zone. An animal population is not just a management variable. A future person is not just a projection. An artificial system is not automatically a subject, but neither should the question be buried if its behaviour suggests future standing may become relevant.

Erasure is not always hostile. Often it is procedural. The form has no field. The model has no category. The law has no route. The meeting has no witness. The dashboard has no measure.

Mechanical Ethics treats missing fields as ethical facts.

A system that cannot name the affected scope cannot claim to have answered for that scope.

9.5 6.4 Do Not Treat Uncertainty as Certainty

The third gate protects the record from false confidence.

gate_3 := do_not_treat_uncertainty_as_certainty

Uncertainty is not a decorative footnote. It changes what may be done, how fast, by whom, under what safeguards, and with what duty to preserve correction.

This gate does not say that uncertainty always means stop. That would be too simple and often harmful. Some decisions must be made under uncertainty. Some delay causes loss. Some evidence can be gathered only by acting carefully. Some emergencies are real.

The gate says that uncertainty must not be hidden, converted into certainty by authority, or routed onto affected scopes without record.

```
uncertainty_record :=  
  what_is_unknown  
  + who_estimates  
  + incentive_position  
  + challenge_route  
  + reversibility_status  
  + timing_pressure
```

The more irreversible the action, the more serious this gate becomes.

If the data is unknown and the action is irreversible, Mechanical Ethics should default toward hold or escalation unless a dirty-conflict or emergency chain is explicitly shown.

If the estimate comes from the actor who benefits from the action, that does not automatically make it false. But it does make it non-final. It needs challenge, independence, or limitation before it carries serious irreversible action.

False certainty is one of the main ways power cleans itself.
The gate refuses it.

9.6 6.5 Do Not Pretend Repair Where Residue Remains

The fourth gate separates correction from repair.

```
gate_4 := do_not_pretend_repair_where_residue_remains
```

A decision can be corrected while harm remains.

A payment can be restored while time is not returned.

A record can be fixed while opportunity is gone.

A policy can be changed while trust is damaged.

An apology can be sincere while the path remains unrepaired.

Mechanical Ethics uses residue to record what remains after correction and repair attempts.

```
residue :=  
  loss_not_restored  
  + time_not_returned  
  + trust_damage  
  + opportunity_expired  
  + record_scar  
  + recurrence_risk  
  + acknowledgement_needed
```

This gate does not say residue makes action unforgivable. It says residue must not be erased by the language of closure.

Systems like closure. Closure protects resources. Closure reduces embarrassment. Closure lets the institution move to the next case. But if closure arrives before residue is acknowledged, the affected scope carries the remainder alone.

The gate requires a more honest record:

```
repair_claim_valid_only_if :=  
  path_touched  
  + affected_scope_considered  
  + residue_checked  
  + recurrence_risk_addressed
```

If the path was not touched, call it correction, not repair.

If time was not returned, say so.

If trust remains damaged, say so.

If an opportunity cannot be restored, say so.

If loss remains dirty, do not make the record clean.

9.7 6.6 Do Not Let Correction Channels Become Harm Carriers

The fifth gate follows from Part 2 and Part 4.

A correction channel is not automatically good because it exists.

```
gate_5 := do_not_let_correction_channels_become_harm_carriers
```

A complaint route can harm if it consumes the affected person's time, energy, evidence, credibility, money, health, or opportunity faster than it can repair the path.

An appeal can harm if it arrives after the loss hardens.

A review can harm if it demands expert knowledge from the person least able to carry the burden.

A support channel can harm if it sends the person in circles while the system continues acting.

A tribunal can harm if victory comes after the practical future has already narrowed.

```
channel_as_harm_carrier :=  
  correction_route  
  + delay_or_cost_or_burden_or_evidence_loss  
  + failure_to_repair_before_hardening
```

This gate is not an attack on review, law, complaint, appeal, or oversight. Those channels matter. Without them, power becomes harder to answer.

The point is that a channel must be judged by whether it can reach the path while reaching still matters.

A correction channel that cannot correct in time should not be used as proof that the system is fair.

It may still be better than nothing. It may still create evidence. It may still assign later responsibility. But it should not be counted as live repair if its timing, burden, or reachability makes repair practically impossible.

9.8 6.7 Do Not Let the Assessed Actor Set Every Clock

The sixth gate protects timing from interested control.

```
gate_6 := do_not_let_assessed_actor_set_every_clock
```

The actor who caused, controls, or benefits from a loss is not a neutral timer of that loss.

This does not mean the actor always lies. It means the actor's timing estimate cannot be the only clock when material harm, irreversibility, or burden is at stake.

A landlord is not the only judge of how long housing insecurity matters.

An employer is not the only judge of how long a record scar matters.

A platform is not the only judge of how fast account loss hardens.

A state is not the only judge of how long a benefit delay remains harmless.

A hospital is not the only judge of when waiting becomes injury.

A developer is not the only judge of when model error becomes irreversible.

```
clock_independence_needed_when :=  
  actor_controls_path  
  + actor_benefits_or_faces_liability  
  + affected_scope_faces_hardening
```

This gate asks for independence, challenge, affected-scope evidence, or conservative treatment of timing under uncertainty.

If the affected scope says the loss is hardening now, that also is not automatically final. But it must be taken as evidence, not noise.

Timing must be answerable because delay is often where harm hides.

9.9 6.8 Do Not Use Emergency to Launder Prior Failure

The seventh gate addresses emergency power.

```
gate_7 := do_not_use_emergency_to_launder_prior_failure
```

Emergencies are real.

Sometimes action must be fast. Sometimes there is no clean process available. Sometimes delay would cause greater loss. Sometimes ordinary safeguards cannot operate at full strength.

Mechanical Ethics does not deny that.

But emergency is also one of the easiest ways to launder power.

A system can neglect maintenance, ignore warnings, underfund correction, centralize authority, dismiss affected scopes, and then claim that the resulting crisis requires exceptional action. If emergency power then erases the prior failure, the emergency becomes a moral solvent.

This gate says: emergency may change what can be done now, but it does not erase how the emergency was produced.

```
emergency_record :=  
  immediate_threat  
  + time_available  
  + safeguards_suspended  
  + prior_warnings  
  + prior_failures  
  + expiry_of_extraordinary_power  
  + residue_after_action
```

Emergency action should be bounded, recorded, reviewable where possible, and followed by residue and responsibility accounting.

If the emergency was unavoidable, say so and show why.

If the emergency was worsened by prior neglect, say so.

If the emergency action imposed dirty loss, say so.

If extraordinary power must be used, it should carry its own expiry and later answerability.

Emergency may justify speed.

It does not justify amnesia.

9.10 6.9 What the Gates Do Not Decide

The first gates do not decide everything.

They do not decide final standing. They do not decide punishment. They do not decide acceptable sacrifice. They do not decide all coercion. They do not decide whether every dirty action must stop. They do not decide whether one scope's future may ever be narrowed to preserve another's.

That restraint is deliberate.

```
gates_do_not_solve :=  
  final_value_priority  
  + legitimacy  
  + desert  
  + punishment  
  + acceptable_sacrifice  
  + all_dirty_conflict
```

The gates do something earlier.

They prevent the later hard questions from being asked on a falsified record.

If loss is hidden, the hard question is corrupted.

If affected scopes are erased, the hard question is corrupted.

If uncertainty is treated as certainty, the hard question is corrupted.

If residue is erased, the hard question is corrupted.

If correction channels become harm carriers, the hard question is corrupted.

If the assessed actor sets every clock, the hard question is corrupted.

If emergency erases prior failure, the hard question is corrupted.

The first gates keep the record honest enough for deeper prescription.

Part 7 now asks who or what receives protection, and why. That question is harder. It cannot be avoided. But it should be asked only after the first gates are in place.

10 Part 7 — Protection, Standing, and Scope

Status: first draft. Incomplete. Not canon.

10.1 7.1 Scope Is Not Personhood

The first mistake is to think that protection begins only after personhood is settled.

Personhood matters. It may become central in many cases. But if Mechanical Ethics waits for final personhood before naming affected scopes, then disputed beings, future beings, non-human beings, dependent beings, and structurally silent beings can be harmed while the argument is still waiting for permission to notice them.

So this book separates scope from personhood.

```
scope != personhood
```

A scope is something whose future-space can be affected in a way that matters for the case record.

A person is one kind of scope. Not the only possible kind.

A family can be a scope. A worker group can be a scope. A disabled person facing an inaccessible process can be a scope. A future child can be a scope. An animal population can be a scope. A river system can be a scope. An institution may be a scope when it carries memory, trust, or correction capacity for people. An artificial system may become a disputed scope if its behaviour raises standing questions that cannot be dismissed cheaply.

This does not mean every scope receives equal protection.

It means a scope should not disappear before the protection question is asked.

```
affected_scope :=  
  entity_or_group_or_system  
  whose_future_space_may_be_materially_changed
```

The language is intentionally broad at the intake stage. Broad intake is not broad permission. It is anti-erasure.

10.2 7.2 Reasons for Protection

Protection needs reasons.

Mechanical Ethics should not create a hidden worth ladder where the writer silently knows who counts and who does not. It should record why a scope receives protection or precaution.

Possible reasons include sentience, vulnerability, dependency, participation, consent status, inability to contest, exposure to irreversible loss, role in sustaining other protected scopes, future-space contraction, viable-becoming damage, or uncertainty about standing under conditions where error would be costly.

```
protection_reason :=  
  sentience_signal  
  | vulnerability  
  | dependency  
  | participation_interest  
  | consent_or_nonconsent  
  | inability_to_contest  
  | irreversible_loss_exposure  
  | future_space_contraction  
  | viable_becoming_damage
```

```
| support_role_for_other_scopes
| standing_uncertainty_with_high_error_cost
```

This list is not final doctrine. It is a record discipline.

A human child receives strong protection for many reasons at once: sentience, vulnerability, dependency, future-space, inability to fully contest, and social recognition. An adult wrongfully routed by a system receives protection because their options, reachability, information, and dignity may be narrowed under a structure they cannot easily challenge. A non-human animal may receive protection because it has sentience, agency, sociality, learning, fear, preference, or vulnerability. A future person may receive protection because present action can close paths before they exist to object.

A scope can also receive weak or precautionary protection.

Protection does not always mean stop. It may mean disclose, monitor, slow, preserve evidence, reduce burden, maintain contestability, avoid irreversible action, preserve opportunity, support repair, or record residue.

The point is not to inflate every scope into full moral equality. The point is to make the protection route visible.

10.2.1 7.2.1 Minimum Floor Candidate

Mechanical Ethics still has an open wound about final standing. It has not proven a complete foundation for every protected scope.

But uncertainty about final standing must not mean that anything unresolved is expendable. The first floor can be modest:

```
minimum_floor_candidate :=
  no_erasure
+ no_unrecorded_irreversible_loss
+ no_burden_without_route
+ no_silent_standing_denial
+ no_clean_language_for_dirty_loss
```

This floor does not grant equal status to every scope. It does not settle personhood, animal standing, ecosystem standing, future-person claims, or artificial moral status.

It says something narrower: if a scope may be materially affected, the system must not erase it, impose irreversible loss without record, shift burden without a route, deny standing silently, or clean the loss through language.

The floor is strongest when four pressures combine:

```
floor_strengthens_when :=
  vulnerability
+ low_contest_power
+ irreversibility
+ actor_control_over_route
```

This is a candidate floor, not canon. It is here to stop the open wound becoming a loophole.

10.2.2 7.2.2 Protection as Preserving Viable Becoming

Protection is not only a brake against harm. It is also a way of preserving the conditions under which a scope can continue, orient, learn, relate, repair, choose meaningfully, and become.

This keeps the chapter from becoming harm-only.

```
protection_as_viable_becoming :=
  preserve_conditions_for
  continuation
+ orientation
+ capacity
+ relation
+ repairability
```

```
+ meaningful_future_space
under_uncertainty
```

A protected scope is not protected merely because something bad might happen to it. It may also be protected because a transition would block learning, destroy trust, remove a real route, sever relation, damage memory, reduce future capacity, or prevent repair before the scope can answer back.

This matters for positive duties. Sometimes the answer is not only do not harm. It may be keep the route open, preserve the support, make information usable, protect the developmental window, maintain the relation, create a reachable challenge path, or repair trust before the next decision hardens.

```
positive_protection_duties :=
  preserve_opportunity
+ keep_information_usable
+ maintain_reachability
+ support_capacity
+ protect_repair_path
+ sustain_trust_where_dependency_exists
```

These duties are still bounded. Mechanical Ethics does not say every scope is owed every possible expansion. It says that when an actor controls a path on which another scope's viable becoming depends, the actor must not reduce protection to merely avoiding obvious injury.

10.3 7.3 Uncertainty Raises Caution

Standing uncertainty is not permission to spend cheaply.

If we do not know whether a scope matters, the answer is not automatically "therefore it does not matter." That is an error rule, and a dangerous one. It gives power to the side that benefits from uncertainty.

```
standing_uncertainty :=
  unknown_status
+ possible_material_loss
+ asymmetric_error_cost
```

When the cost of wrongly denying protection is high, uncertainty should raise caution.

This does not mean every uncertain scope receives full protection. That would create paralysis and could erase other scopes. It means uncertainty must be recorded and routed.

Mechanical Ethics can use graduated responses:

```
uncertainty_response :=
  record_dispute
+ avoid_irreversible_loss_where_possible
+ seek_better_evidence
+ preserve_future_contestability
+ use_precaution_when_error_cost_high
```

Consider a disputed artificial system. It may not be a moral patient. Most current systems may not deserve standing in any strong sense. But if a future system shows persistent self-protective behaviour, suffering-like reports, social dependence, or complex preference continuity, it would be reckless to treat uncertainty as proof of non-standing. The answer may be investigation, constraint, and precaution, not immediate equality.

Consider animals. A system that treats animal intelligence as irrelevant because animals cannot enter the human legal forum is not being neutral. It is using forum membership as a worth filter.

Consider future persons. They cannot object now. That silence cannot be used as consent.

Uncertainty does not settle protection. It changes the burden of care.

10.4 7.4 Vulnerability and Dependency

Vulnerability changes the ethical weight of a transition.

A vulnerable scope is not merely one that can be harmed. All scopes can be harmed in some way. Vulnerability means a scope has reduced ability to absorb, avoid, understand, contest, or repair a narrowing of future-space.

```
vulnerability :=  
  low_buffer  
  | low_contest_power  
  | high_dependency  
  | high_exposure_to_irreversibility  
  | low_repair_capacity
```

Dependency matters because it makes the actor's conduct more dangerous.

A child depends on adults. A patient depends on a medical system. A benefit claimant depends on administrative timing. A prisoner depends on state custody. A worker may depend on a platform or employer record. An animal in a human-shaped environment may depend on human restraint. A community downstream depends on decisions made upstream.

Dependency does not make the dependent scope morally pure. It does not remove all duty from the dependent actor. But it changes what the powerful actor may claim.

The more one scope depends on another's process, the less acceptable it is for that process to be opaque, slow, unreachable, or self-certifying.

```
dependency_pressure :=  
  actor_controls_route  
  + affected_scope_depends_on_route  
  -> higher_duty_of_disclosure_contestability_and_timely_repair
```

Vulnerability and dependency are not sentimental categories. They are mechanical facts about who can survive delay, burden, uncertainty, and error.

They are also facts about positive becoming. A dependent scope may need more than non-interference. It may need timely support, usable information, preserved relation, or a stable route because its own capacity is still forming, blocked, or held by another system.

```
dependency_positive_pressure :=  
  affected_scope_capacity_depends_on_actor_route  
  -> duty_to_preserve_conditions_for_viable_becoming
```

10.5 7.5 Participation and Consent

Consent matters, but consent is not magic.

A scope that participates in a transition may carry different duties and risks than a scope that is dragged into it. But participation must be real enough to count.

```
consent_question :=  
  did_scope_have_information  
  + real_alternative  
  + capacity  
  + low_coercion  
  + ability_to_withdraw_or_contest
```

Mechanical Ethics should not treat formal acceptance as full consent if the affected scope lacked information, realistic alternatives, or freedom from pressure.

A worker may "agree" to a condition because refusal means unemployment. A patient may "agree" because the system withheld practical options. A user may "agree" to terms no one can read. A community may "consult" after the decision has effectively been made. An animal cannot consent in human form. A future person cannot consent at all.

Consent can reduce certain burdens. It can authorize risk. It can make shared projects possible. But it cannot be used to hide asymmetry.

Consent also has timing. Consent before information changes is not consent to everything later discovered. Consent under one risk profile is not consent under another. Consent to participate is not consent to unrepaired residue.

```
consent_is_bounded_by :=
  information
+ scope
+ time
+ coercion_level
+ withdrawal_or_contestability
```

This becomes important in dirty conflicts and institutional settings, where people are often counted as participants in systems they cannot practically leave.

10.6 7.6 Non-Human Life

Non-human life is not a decoration around human ethics.

Animals act, learn, fear, prefer, bond, explore, avoid, remember, and suffer in different ways. Some show complex social behaviour and flexible intelligence. Ecosystems carry interdependent continuations that no single organism owns. Human systems routinely change these futures at scale.

Mechanical Ethics should not pretend this is simple.

```
nonhuman_scope_questions :=
  sentience_signal
+ agency_signal
+ dependency_on_human_systems
+ ecological_role
+ vulnerability
+ irreversibility_of_loss
```

This does not mean non-human life always overrides human safety, food, health, shelter, or ecological protection. It does mean non-human loss should not be laundered as mere management when the affected beings or systems have relevant continuation interests.

A dirty conflict involving intelligent invasive animals is still dirty. Ecological protection may require killing. Human safety may require control. Crop protection may be legitimate. But the animals do not become morally invisible because the management need is real.

The record should say what is being protected, what is being killed, what alternatives were considered, what suffering is imposed, what residue remains, and whether human-caused conditions helped create the conflict.

The aim is not purity. The aim is not sentimental denial of ecological reality.

The aim is to stop necessary harm from becoming clean language.

It is also to ask whether human systems can create better conditions before the dirty conflict hardens: habitat repair, prevention, lower-suffering control, redesigned incentives, reduced waste, slower expansion, or stewardship that preserves more viable becoming across scopes.

10.7 7.7 Future Persons

Future persons cannot object.

They cannot file a complaint, attend a meeting, vote in the present, correct a record, or tell us which losses they would have accepted. This makes them easy to erase.

```
future_scope :=
  not_present_to_contest
+ affected_by_present_transition
+ possible_future_space_contraction
```

Mechanical Ethics should not treat absence as consent.

The difficulty is that future persons are not a single known group with clear present preferences. Claims about them can be abused. A state can invoke future generations to justify present coercion. An institution can claim long-term benefit while imposing unrecorded loss now. A movement can use imagined future approval to silence actual affected people.

So future-person reasoning needs discipline.

It should be strongest where present action creates irreversible or long-lasting future-space contraction, especially where those future scopes will have no realistic repair route.

It should be weaker where the claim is speculative, reversible, or used to erase present affected scopes without an audit chain.

```
future_person_caution :=
  stronger_when_irreversible
+ stronger_when_repair_unavailable
+ weaker_when_speculative_or_launders_present_power
```

Future persons matter. They are also easy to weaponize. Both truths must remain visible.

Their protection is not only about avoiding inherited loss. It is also about preserving the conditions through which future scopes may orient, learn, trust institutions, inherit usable information, and repair what remains repairable.

10.8 7.8 Artificial Systems

Artificial systems create two separate questions.

The first is easy: artificial systems can affect protected scopes and therefore must be constrained.

The second is hard: some artificial systems might eventually become protected scopes themselves.

Mechanical Ethics should not collapse these questions.

```
AI_question_split :=
  constrain_AI_as_actor_or_tool
!= grant_AI_moral_standing
```

A present model used to deny benefits, allocate medical risk, route police attention, score workers, recommend content, or operate infrastructure can be ethically significant even if it has no moral standing at all. The duties may attach to the people and institutions that design, deploy, maintain, and benefit from it.

That is the first question: acting systems under human delegation must not hide loss, uncertainty, burden, or correction failure behind technical opacity.

The second question concerns possible future minds or artificial subjects. This remains unresolved. Mechanical Ethics should not grant standing casually. It should also not build a framework that makes future standing impossible to recognize because the category was excluded in advance.

A cautious position is this:

```
artificial_standing_status :=
  currently_unresolved
+ precaution_when_signals_and_error_cost_require
+ no_default_full_standing
+ no_default_erasure
```

The book can constrain artificial systems before deciding whether any artificial system counts as a protected scope.

That distinction must remain clean.

10.9 7.9 When Standing Remains Unresolved

Some standing questions will remain unresolved.

That is not a defect if the unresolved status is honest.

Mechanical Ethics should have outputs for unresolved standing:

```
standing_outputs :=  
  protection_triggered  
  | precaution_triggered  
  | standing_unresolved  
  | no_protection_reason_found
```

No protection reason found is not the same as “worthless.” It means the current record has not shown a reason for special protection under Mechanical Ethics.

Standing unresolved is not the same as “ignore.” It means the case may require hold, escalation, narrower action, evidence gathering, or precaution depending on the risk.

Precaution triggered is not the same as full equality. It means error costs or uncertainty justify constraint.

Protection triggered is not the end of the case. It means duties now have to be derived.

The open wound remains:

```
open_wound :=  
  final_floor_foundation_for_scope_protection
```

Mechanical Ethics needs minimum floors, but it has not yet fully solved who sets them, how they are justified, and how they operate across humans, animals, ecosystems, future persons, and artificial systems.

The minimum floor candidate above reduces the immediate loophole but does not close the wound. It says unresolved standing cannot justify erasure or unrecorded irreversible loss. It does not prove a final hierarchy of protection.

The honest move is not to hide that wound. The honest move is to carry it forward and prevent it from being silently settled by whoever has power.

Part 7 therefore ends with a limited achievement.

It does not solve moral standing.

It gives the book a way to name affected scopes, record protection reasons, raise caution under uncertainty, preserve a minimum anti-erasure floor, and connect protection to viable becoming without pretending that every protected scope receives the same duties.

The next question is what happens after harm, error, dirty loss, or blocked becoming has already entered the path.

That is correction, repair, and residue.

11 Part 8 — Correction, Repair, and Residue

Status: first draft. Incomplete. Not canon.

11.1 8.1 Correction Is Not Repair

Correction and repair are not the same thing.

Correction changes an official state. Repair changes the harmed path as far as possible.

```
correction :=  
  record_or_decision_changed  
  
repair :=  
  harmed_path_touched  
  + future_space_restored_where_possible  
  + recurrence_risk_reduced  
  + residue_recorded
```

A file can be corrected while a person's future remains narrowed. A payment can be restored after debt, fear, eviction risk, or family damage has already occurred. A decision can be reversed after a job, school place, treatment window, or appeal route has already closed.

Correction matters. It is better than refusing error. It may be the first step toward repair.

But it should not be allowed to impersonate repair.

The distinction matters because institutions often treat correction as closure. The system finds the mistake, changes the record, issues a note, possibly pays compensation, and moves on. The affected person is expected to accept that the official state now matches what should have happened earlier.

But earlier is gone.

Mechanical Ethics asks what the correction failed to return.

```
repair_question :=
  what_path_was_harmed
  + what_can_still_be_restored
  + what_cannot_be_restored
  + what_must_change_to_prevent_recurrence
```

This is the beginning of repair discipline.

11.1.1 8.1.1 Repair as Reopened Viable Future

Repair is not only the removal of a bad mark. It is an attempt to reopen viable becoming where a transition damaged, narrowed, delayed, confused, isolated, or blocked it.

```
repair_as_reopened_viable_future :=
  restore_or_reopen
  continuation
  + orientation
  + capacity
  + relation
  + repairability
  + meaningful_future_space
  where_possible
  + record_what_remains_closed
```

This keeps repair connected to the value spine. If a system damaged orientation, repair must restore usable information. If it damaged capacity, repair must support capacity where possible. If it damaged relation or trust, repair cannot be reduced to record correction. If it blocked meaningful future-space, repair must ask whether a real path can still be reopened.

Repair therefore has a positive side. It is not only cleaning up after harm. It is a way of preserving or rebuilding the conditions under which a scope can continue and become.

```
repair_success_candidate :=
  path_reopened
  + burden_reduced
  + usable_information_restored
  + trust_partly_rebuilt
  + recurrence_less_likely
  + residue_honestly_carried
```

The word candidate remains. Repair success is not declared from above. It must be checked against affected scopes, timing, recurrence, and residue.

11.2 8.2 Repair Must Touch the Path

Repair must touch the path that caused the harm.

It is not enough to change language around the path. It is not enough to apologise for the path. It is not enough to compensate for one case while preserving the same machinery for the next case.

```
repair_must_touch_path :=
  mechanism_changed
  or burden_routed_back
```



```
or reachability_restored
or evidence_preserved
or recurrence_risk_reduced
```

If the harm came from delay, repair must address delay.

If the harm came from unreachable appeal, repair must address reachability.

If the harm came from hidden evidence, repair must address evidence and disclosure.

If the harm came from contaminated estimation, repair must address estimator independence and challenge.

If the harm came from local reward pull, repair must address what the system is rewarded for.

A system that harms through timing cannot repair only through apology.

A system that harms through burden cannot repair only through information.

A system that harms through opacity cannot repair only through speed.

A system that harms through recurrence cannot repair only the last visible case.

Mechanical Ethics therefore treats repair as path-specific.

The question is always: what mechanism narrowed future-space or blocked viable becoming, and has that mechanism been constrained?

11.3 8.3 Slow the Write

Some systems write faster than people can answer.

They write records. They write classifications. They write debts. They write refusals. They write risk scores. They write exclusions. They write defaults. They write police attention. They write medical priority. They write school futures. They write immigration status. They write social trust.

When writing is fast and correction is slow, harm can harden before the affected scope can contest.

```
write_speed_problem :=
  T_write
  < T_detect + T_route + T_correct
```

The first repair may be to slow the write.

This does not mean every system must stop. It means a system should not be allowed to make irreversible or hardening changes at a speed that outruns meaningful correction.

Slowing the write can mean temporary pause, staged effect, reversible status, warning period, human review before hardening, automatic buffer, independent check, or delayed enforcement.

```
slow_write_options :=
  pause
  + stage
  + buffer
  + reversible_status
  + pre_hardenening_review
  + independent_check
```

The point is to keep correction real while there is still something to correct.

The positive form is to keep the window open.

```
keep_window_open :=
  prevent_hardenening
  + preserve_contestability
  + preserve_opportunity
  + preserve_repairability
```

If the write cannot be slowed because urgent action is required, then the record should say why, what loss may be created, what residue may remain, and what repair duty follows.

Speed is not neutral. Speed allocates power.

11.4 8.4 Route the Burden Back

Harm often travels by burden shift.

A system creates uncertainty, but the affected person must prove certainty. A model makes an error, but the worker must disprove it. A record is wrong, but the claimant must gather evidence. A public body delays, but the citizen carries rent, hunger, fear, or lost opportunity. A platform acts automatically, but the user must navigate opaque appeal.

```
burden_shift :=  
  system_uncertainty  
  routed_to  
  affected_scope
```

Not every burden shift is wrong. Shared systems require participants to provide information, cooperate, and sometimes wait. But burden becomes ethically serious when it is forced onto those least able to carry it, especially where the system caused or controls the uncertainty.

Repair should route burden back where possible.

```
burden_return :=  
  system_rechecks  
  + system_preserves_evidence  
  + system_contacts_sources  
  + system_pauses_enforcement  
  + system_funds_access  
  + system_explains_decision
```

The affected person should not have to become a lawyer, engineer, data scientist, doctor, or bureaucratic expert merely to stop a system from injuring them through its own uncertainty.

This does not remove all responsibility from affected persons. It says the route must be proportionate to power, access, and cause.

If the system created the fog, it should not make the exposed person carry the whole lantern.

Burden return also protects positive becoming. A person or scope carrying the system's uncertainty often loses capacity elsewhere: attention, time, trust, health, opportunity, relation, learning, or work. Routing burden back can therefore reopen more than the formal case.

```
burden_return_positive_effect :=  
  affected_scope_capacity_recovered  
  + usable_attention_returned  
  + trust_route_partly_restored  
  + opportunity_preserved_or_reopened
```

11.5 8.5 Preserve Contradiction

Repair requires contradiction to survive.

A system that destroys contradiction destroys repair. If the affected person's account is overwritten, if records vanish, if informal warnings are ignored, if dissent is treated as noise, if evidence is discarded after closure, then later correction becomes theatre.

```
preserve_contradiction :=  
  keep_dispute_visible  
  + preserve_evidence  
  + prevent_official_story_closure  
  + allow_challenge_before_hardening
```

Contradiction is not always truth. People can be wrong. Witnesses can misremember. Affected scopes can misunderstand the mechanism. Records can be messy.

But contradiction is often where the path first becomes visible.

A complaint, a refusal, a mismatch, a repeated anomaly, a local witness, a person saying "this cannot be right" — these may be the first signals that the official model has lost contact with reality.

Mechanical Ethics does not say contradiction should always win.

It says contradiction should not be erased before it can be tested, especially when the official path can harden faster than the challenge route.

A system that cannot preserve contradiction cannot honestly claim robust correction.

Preserved contradiction also preserves orientation. It keeps the affected scope from being forced into a false world where the official record is treated as the only reality.

11.6 8.6 Keep Contestability Live

Contestability is the practical ability to challenge a transition before it hardens.

It is not enough that a right to contest exists somewhere in the rules.

```
contestability :=
  known_route
  + reachable_route
  + usable_information
  + time_before_hardening
  + authority_to_change_path
```

A challenge route is weak if the affected scope does not know it exists, cannot understand it, cannot afford it, cannot access it, cannot gather evidence for it, or cannot get a decision before loss hardens.

Mechanical Ethics treats contestability as a living property, not a formal checkbox.

If contestability is live, a system may be able to act with less risk. If contestability is dead, action carries more ethical burden.

Keeping contestability live may require notice, translation, advice, time, evidence disclosure, pause, independent review, or temporary protection against enforcement.

```
contestability_support :=
  notice
  + explanation
  + evidence_access
  + time
  + affordable_route
  + pause_or_buffer
  + decision_authority
```

The point is not to make every action litigable forever. The point is to prevent systems from saying “you could have challenged” where challenge was practically unreachable until after the damage was done.

Contestability is also a positive structure. It is part of what makes a future usable rather than merely possible. A scope that can challenge, correct, and return is less trapped by the first estimate made about it.

```
contestability_positive_function :=
  preserves_orientation
  + preserves_agency_under_pressure
  + preserves_repairability
  + limits_false_finality
```

11.7 8.7 Proof Is Not Theatre

Institutions like proof of reform.

They produce new policies, training, dashboards, forms, audits, statements, committees, and lessons learned. Some of these matter. Some are theatre.

Proof is theatre when it demonstrates that something was said, not that the path changed.

```
proof_theatre :=
  visible_reform_signal
  + unchanged_harm_mechanism
```

Mechanical Ethics asks for path proof.

If delay caused harm, show delay changed.

If hidden evidence caused harm, show evidence access changed.

If burden shift caused harm, show burden routing changed.

If local reward caused harm, show reward changed.

If no one could answer, show responsibility attachment changed.

If correction was unreachable, show reachability changed.

```
path_proof :=
  mechanism_before
  + mechanism_after
  + affected_scope_test
  + recurrence_risk_reduction
```

This does not require impossible certainty. It requires the reform to touch the mechanism it claims to repair.

A sincere institution can still produce theatre if it measures the wrong thing.

A cynical institution can produce theatre deliberately.

The test is the same: did the path change?

For value-spine purposes, a second test follows: did the change restore or preserve viable becoming, or only improve the institution's account of itself?

```
path_proof_positive_test :=
  mechanism_changed
  + affected_scope_can_navigate_better
  + repair_route_more_reachable
  + recurrence_less_likely
```

11.8 8.8 Residue Remains

Some things are not restored.

Time is not always returned. Trust is not always restored. Opportunity is not always recreated. Health is not always recovered. A record scar may remain. A relationship may remain damaged. A person may remain altered by carrying the burden. A community may remember. Loss in a necessary dirty conflict is not undone merely because the policy was justified. A future-space contraction may leave a shape that compensation cannot erase.

Mechanical Ethics calls this residue.

```
residue :=
  loss_not_restored
  + time_not_returned
  + trust_damage
  + opportunity_expired
  + record_scar
  + recurrence_risk
  + acknowledgement_needed
```

Residue is not a reason to give up on repair. It is a reason to stop lying about repair.

The record should say what remains.

A residue ledger does not have to be long. It has to be honest.

```
residue_ledger :=
  what_remains
  + who_carries_it
  + whether_it_can_be_reduced
  + who_still_owes_attention
```

Without residue, institutions overclaim. They turn partial repair into full closure. They count money as time. They count apology as restoration. They count compliance as trust.

Residue keeps the remainder visible.

11.8.1 8.8.1 Residue as Unrestored Becoming

Residue is not only leftover harm. It is also the record of viable becoming not restored.

```
residue_as_unrestored_becoming :=
  capacity_not_recovered
+ trust_not_rebuilt
+ relation_not_repaired
+ opportunity_not_reopened
+ orientation_not_restored
+ recurrence_still_possible
```

This matters because a system can repair one layer while leaving another closed. It can correct a record but not rebuild trust. It can compensate money but not restore time. It can apologize but not reopen the lost path. It can change policy but leave the affected scope without orientation, relation, or confidence that the same thing will not happen again.

The residue ledger therefore protects the positive layer too. It records not only what was damaged, but what future remains blocked.

11.9 8.9 Duties After Late Correction

When correction arrives late, duties do not disappear.

They change shape.

```
late_correction_duties :=
  acknowledge_lateness
+ identify_expired_opportunity
+ repair_path_where_possible
+ compensate_where_appropriate
+ record_residue
+ reduce_recurrence
```

If a person wins an appeal after the loss has hardened, the system should not say only “appeal allowed.” It should ask what the late route cost.

If a record is corrected after a decision was already made, the system should ask what opportunity the wrong record affected.

If payment is restored after debt accumulated, the system should ask what debt, stress, penalty, trust loss, or practical narrowing remains.

If a policy is changed after many cases, the system should ask who carries unrepaired residue from the earlier path.

Late correction may still matter. It can restore some options, acknowledge truth, reduce future recurrence, and create evidence for responsibility.

But lateness must be part of the duty record.

A system that corrects late and records no lateness is still hiding part of the harm.

The positive version is this: late correction should still reopen what can be reopened, protect whatever viable future remains, and prevent the same narrowing from being passed to the next scope.

11.10 8.10 Stewardship Over Time

Repair is not a moment.

Some duties persist.

A changed policy can decay. A corrected record can be copied from an old database. A compensated person can face recurrence. A review channel can become slow again. A model can drift. A public commitment can fade after attention moves on.

```
stewardship :=  
  repair_over_time  
  + recurrence_monitoring  
  + memory_preservation  
  + renewed_contestability
```

Stewardship is the duty to keep repair from becoming a performance that expires as soon as pressure drops.

This matters especially in large systems. Institutions forget. Staff change. Vendors change. Models update. Laws shift. Budgets tighten. Emergency patches become ordinary routines. The old path may return in a new form.

Mechanical Ethics therefore treats repair as requiring memory.

Not endless punishment. Not permanent suspicion. Memory.

A system that harmed through timing should remember timing. A system that harmed through opacity should remember opacity. A system that harmed through burden shift should remember burden shift. A system that harmed through contaminated estimation should remember who was allowed to estimate and why that failed.

```
repair_memory :=  
  harm_mechanism  
  + repair_action  
  + residue  
  + recurrence_signal  
  + owner_of_follow_up
```

Stewardship is also how repair becomes positive rather than merely defensive. It keeps the repaired path from collapsing again. It preserves the trust, reachability, and usable information that make future cooperation possible.

```
stewardship_positive_function :=  
  preserve_reopened_path  
  + maintain_trust_after_repair  
  + keep_correction_reachable  
  + prevent_recurrence_from_consuming_future
```

Part 8 ends with a practical claim.

Correction without repair is incomplete.

Repair without residue is dishonest.

Residue without stewardship becomes abandonment.

And repair, when it succeeds, is not only the reduction of harm. It is the reopening of viable future under uncertainty.

The next part turns to the hardest cases: where every available action causes loss, and the task is not to find a clean answer but to stop the dirty answer from pretending to be clean.

12 Part 9 — Dirty Conflicts

Status: first draft. Incomplete. Not canon.

12.1 9.1 When Every Action Causes Loss

Some cases do not have a clean path.

Action causes loss. Inaction causes loss. Delay causes loss. Speed causes loss. Protection of one scope narrows another. Refusal to choose lets a harmful path continue. Choosing distributes harm through human hands.

Mechanical Ethics calls this a dirty conflict.

```
dirty_conflict :=  
  every_available_path  
  causes_material_loss
```

Dirty conflict is not ordinary inconvenience. It is not every disagreement. It is not a case where someone dislikes an outcome. It is a case where all reachable options impose material future-space contraction on at least one affected scope.

This is where ethical language often becomes dishonest.

People want the action to be clean because the action may be necessary. Institutions want the action to be clean because they need authority to act. Movements want the action to be clean because hesitation feels like betrayal. Experts want the action to be clean because the model selects a path. States want the action to be clean because emergency needs speed.

But necessity does not clean loss.

```
necessary_harm  
  != clean_harm
```

A dirty conflict may still require action. Mechanical Ethics does not say that dirty means forbidden. It says dirty means the loss must not disappear.

12.1.1 9.1.1 Incompatible Viable Becoming

The value-spine version is sharper.

A dirty conflict is not only a conflict between harms. It is a conflict where viable becoming cannot be preserved for all affected scopes at once.

```
dirty_conflict_as_viable_becoming_conflict :=  
  viable_becoming_claim_A  
  incompatible_with  
  viable_becoming_claim_B  
  under_current_constraints  
  + uncertainty_remaining
```

One scope's continuation may require another scope's burden. One group's safety may require another group's restriction. One ecosystem path may require loss for some beings inside it. One institution's stability may conflict with one person's timely repair. One future benefit may collide with present agency.

This does not mean all claims are equal. It means the conflict must be named at the level where it actually lives: not merely as competing preferences, but as incompatible claims on continuation, orientation, capacity, relation, repairability, and meaningful future-space.

```
no_clean_total_good :=  
  when_not_all_viable_becoming_can_be_preserved  
  claimed_good_must_not_erase_remaining_loss
```

The first duty in a dirty conflict is therefore not moral cleanliness. It is honest accounting.

12.2 9.2 Why Clean Answers Are Usually Lies

Clean answers are attractive.

They reduce guilt. They simplify communication. They protect authority. They let institutions proceed. They give the public a story. They give the actor a shield.

In dirty conflict, clean answers usually hide something.

They hide the sacrificed scope. They hide uncertainty. They hide the value tilt. They hide the nonparticipant. They hide the least repairable loss. They hide the fact that the chosen path may be best only under a specific priority that should be named.

```
cleaning_failure :=
  dirty_loss
  relabelled_as
  solved_or_costless_or_deserved_or_invisible
```

Mechanical Ethics refuses clean language where loss remains.

If a population is culled to protect an ecosystem, say what is killed and why. If a patient is deprioritized to save more patients, say what residue remains. If a city project displaces a community to improve infrastructure, say who carries the contraction. If a model rejects some applicants to reduce fraud, say who bears false positives and how they can contest. If emergency powers restrict liberty to prevent wider harm, say what liberty is restricted and how power expires.

The point is not to paralyse action.

The point is to keep action answerable.

Clean answers also hide false goods. A path can produce a real benefit and still be dirty. A public gain may be real. A safety improvement may be real. A restored ecosystem may be real. A faster process may be real. But a real benefit is not absolution.

```
benefit_in_dirty_conflict :=
  may_be_real
+ may_be_needed
+ may_reduce_wider_loss
- does_not_cancel_residue_by_itself
```

12.3 9.3 No Netting Without Permission

One scope's gain does not automatically cancel another scope's loss.

```
no_netting_without_permission :=
  expansion_for_A
  does_not_automatically_cancel
  contraction_for_B
```

Netting is sometimes unavoidable. Public decisions often compare harms and benefits across groups. Triage compares outcomes. Infrastructure compares costs and gains. Ecology compares interdependent losses. Law and policy regularly act across plural interests.

Mechanical Ethics does not ban comparison.

It bans hidden comparison.

If one scope is narrowed for another's expansion, the record should say so. If the system treats the trade as justified, the value chain should be visible. If consent is absent, that absence should be visible. If compensation is offered, it should not be described as erasure. If the harmed scope cannot contest, that should raise caution.

```
netting_audit :=
  whose_gain
+ whose_loss
+ who_decided
+ by_what_authority
+ what_consent_or_contest
+ what_repair_or_residue
```

The value-spine audit adds a second layer:

```
benefit_cost_record :=
  what_viable_becoming_expands
+ for_whom
+ what_viable_becoming_contracts
+ for_whom
+ what_uncertainty_remains
+ what_repair_or_residue_follows
```


The danger is not only bad arithmetic. The danger is moral disappearance.

A net benefit can hide a destroyed life. A public gain can hide an unchosen burden. A statistical improvement can hide a class of false positives. An ecological necessity can hide the suffering of intelligent animals. A future benefit can hide present coercion.

Netting may sometimes be justified.

It must not be silent.

12.4 9.4 The Least-Repairable Loss

In dirty conflict, not all losses have the same repair profile.

Some can be compensated. Some can be reversed. Some can be reduced later. Some create permanent absence. Some remove time that cannot be returned. Some destroy trust. Some change identity, body, home, family, memory, or species continuity.

Mechanical Ethics therefore asks for the least-repairable loss.

```
least_repairable_loss :=  
  loss_with_lowest_realistic_restoration_path
```

This is not neutral. It carries a value tilt toward reversibility and repairability.

The tilt is deliberate, but it must be declared.

```
choose_least_unrepairable_path_if_action_required :=  
  declared_value_tilt  
  + requires_its_own_audit_chain
```

Why privilege repairability at all?

Because action under uncertainty should preserve the possibility of correction where it can. When we do not know enough, destroying the route back is especially serious. A reversible burden may be preferable to an irreversible destruction, even if neither is clean.

But the rule is not absolute.

Some irreversible losses may be accepted to prevent wider or deeper irreversible losses. Some reversible burdens may be unjust because they are imposed repeatedly on the same vulnerable scope. Some repairable harms may still be degrading, coercive, or illegitimate. Some losses are incommensurable.

So least-repairable loss is a question, not an automatic answer.

The positive formulation is: preserve the most viable future without lying about whose future was narrowed.

```
least_unrepairable_value_tilt :=  
  preserve_repairability_where_possible  
  + preserve_correction_under_uncertainty  
  + declare_who_carries_remaining_closure
```

12.5 9.5 The Least-Reachable Scope

Dirty conflicts often harm those least able to reach the decision.

The people not in the room. The animals outside the forum. The future persons not yet born. The users without appeal. The patient without specialist language. The claimant without legal help. The community without political power. The worker without bargaining power. The ecosystem without a representative.

```
least_reachable_scope :=  
  affected_scope  
  with_lowest_practical_access_to  
  information_or_contest_or_repair
```

This matters because formal equality can hide practical exclusion.

A process may be open to all and reachable by few. A consultation may invite everyone and hear only the organized. A complaint channel may exist but require literacy, time, money, safety, or expertise. A public policy may count all affected groups but take evidence only from those who can produce it in the expected form.

Mechanical Ethics asks: who cannot reach the route that claims to protect them?

The least-reachable scope should not automatically win every conflict. That would be too simple. But if a dirty decision burdens a scope that cannot contest, the decision carries a higher duty of representation, caution, disclosure, and residue.

```
least_reachable_scope_pressure :=
  low_access
+ material_loss
-> heightened_duty_to_record_represent_and_repair
```

The value-spine version adds: who cannot reach the route that would preserve or reopen their viable becoming?

```
least_reachable_becoming_pressure :=
  low_access
+ viable_becoming_contraction
-> heightened_duty_to_represent_and_keep_repair_live
```

Dirty conflict becomes more dangerous when the harmed scope cannot even appear as harmed.

12.6 9.6 Consent, Compensation, and Burden Return

Consent can change dirty conflict, but it rarely cleans it completely.

A person may consent to risk. A community may accept disruption for shared benefit. A patient may accept a treatment tradeoff. A worker may accept a temporary burden. A citizen may accept some restriction under emergency.

But consent must be real enough to carry weight.

```
consent_weight_depends_on :=
  information
+ real_alternative
+ low_coercion
+ capacity
+ ability_to_withdraw_or_contest
```

Where consent is absent, weaker, coerced, or impossible, the dirty residue increases.

Compensation also matters but does not erase everything.

Money may repair some reachability. It may acknowledge burden. It may reduce future contraction. It may be owed. But money does not always return time, health, trust, home, lost childhood, ecological destruction, bodily integrity, or life.

```
compensation != erasure
```

Burden return is often more important than symbolic closure.

If the system created the conflict, the system should not offload all navigation, proof, waiting, cost, or repair onto the affected scope. If a dirty action is necessary, the actor may owe ongoing burden return: help, access, explanation, monitoring, compensation, mitigation, and recurrence prevention.

The question is not only, "Was the sacrifice justified?"

It is also, "Who carries the aftermath?"

And the value-spine question is: whose viable becoming remains burdened after the claimed good has been achieved?

```
aftermath_question :=
  who_carries_the_remaining_burden
+ what_future_remains_closed
+ what_repair_stays_live
```

12.7 9.7 Triage Without Moral Disappearance

Triage is one of the hardest dirty conflicts.

Medical triage, emergency response, disaster allocation, ecological intervention, infrastructure planning, platform moderation, wartime evacuation, child protection, scarce housing, scarce benefits, scarce attention — all can force choices where not everyone can receive what they need.

Mechanical Ethics cannot make triage clean.

It can demand that triage not erase.

```
triage_discipline :=
  name_affected_scopes
+ name_scarcity_or_conflict
+ name_selection_criteria
+ preserve_challenge_where_possible
+ record_residue
+ prevent_hidden_worth_ladder
```

A triage rule must not secretly decide that some people count less while publicly claiming neutral procedure. It must not use technical criteria to hide social worth judgments. It must not make the excluded disappear because the selected path helped more people.

This is difficult because triage often requires speed.

Emergency triage cannot always provide full deliberation. But speed does not remove the need for later record. If decisions must be made quickly, the residue and criteria should be recorded after. If criteria were wrong, repair should follow where possible. If some losses could not be repaired, they should remain visible.

Triage is where Mechanical Ethics must be most disciplined and least theatrical.

A triage decision may preserve more viable becoming than any available alternative. That matters. But it still must name who was not preserved, who was asked to carry the closure, and what is owed after the choice.

12.8 9.8 Residue After Forced Loss

Forced loss leaves residue even when justified.

A person lawfully detained has lost liberty. A patient denied a scarce treatment has lost a chance. A community displaced for flood protection has lost place. An animal killed for ecological protection has lost life. A future person inheriting a narrowed world has lost options they never had a chance to defend.

Justification may change guilt. It may change duty. It may make action necessary. It may prevent greater harm.

It does not erase residue.

```
justified_loss
  may_reduce_wrongdoing
  but_does_not_erase_residue
```

Residue after forced loss should include the loss itself, who imposed it, why it was imposed, what alternatives were rejected, what repair or compensation is possible, what remains impossible to restore, and what future constraint follows.

A dirty decision without residue record trains the actor to become cleaner in its own eyes each time it acts.

That is dangerous.

Moral injury can also be residue. The actor forced to choose may carry damage. The witness may carry damage. The institution may carry trust damage. These do not replace the affected scope's loss, but they matter for future answerability.

A system that repeatedly forces dirty choices should not merely celebrate decisive action. It should ask why dirty choice became routine.

Residue should also record what viable becoming remains blocked.

```
forced_loss_residue :=  
  loss_imposed  
  + viable_becoming_closed  
  + repair_possible_or_impossible  
  + burden_carrier_named  
  + future_constraint_recorded
```

12.9 9.9 When ME Must Hold or Escalate

Mechanical Ethics must be able to refuse to decide.

Some dirty conflicts exceed the current record. Some depend on unresolved value-cruxes. Some require democratic authority, legal process, specialist evidence, affected-scope representation, or independent review. Some are too contaminated by the actor's interest. Some carry irreversible loss under unknown facts.

```
ME_must_not_fake_decision_when :=  
  material_unknowns  
  + irreversible_action  
  + unresolved_value_crux  
  + contaminated_authority  
  + inadequate_scope_representation
```

The proper output may be HOLD.

The proper output may be ESCALATE.

The proper output may be UNRESOLVED_VALUE_CRUX.

The proper output may be PROCEED_ONLY_WITH_NAMED_CONSTRAINTS if delay itself causes greater loss and the constraints are real.

A framework that cannot hold is not ethics. It is acceleration.

A framework that always holds is not ethics either. It is paralysis.

Dirty conflict is the place where Mechanical Ethics must show discipline: neither cleaning loss nor refusing responsibility merely because the answer is not pure.

12.9.1 9.9.1 Dirty-Conflict Routing Table

This table is not a solver. It is a routing guard against fake cleanliness.

```
dirty_conflict_routing :=  
  HALT when  
    proposed_path_crosses_known_floor  
    or proposed_path_depends_on_erased_scope  
    or proposed_path_requires_false_cleanliness  
  
  HOLD when  
    material_unknowns + irreversible_loss  
    and delay_does_not_create_greater_irreversible_loss  
  
  ESCALATE when  
    unresolved_value_crux  
    or coercive_authority_question  
    or inadequate_scope_representation  
    or contaminated_estimate_controls_material_loss
```

```
PROCEED_ONLY_WITH_NAMED_CONSTRAINTS when
  action_is_needed_to_prevent_greater_or_faster_loss
+ affected_scopes_named
+ value_tilt_declared
+ least_reachable_scope_accounted_for
+ claimed_good_and_remaining_loss_recorded
+ correction_or_residue_duties_attached
```

```
REPAIR_REQUIRED or COMPENSATION_REQUIRED when
  dirty_loss_has_already_occurred
and remaining_duty_concerns_path_repair_or_burden_return
```

The table does not make a dirty conflict clean. It only blocks two common failures: acting as if necessity erased loss, and holding as if delay were outside the conflict.

If the routing table cannot be filled without inventing facts, the case should not be forced into a confident output.

Part 9 ends with the hardest practical instruction so far:

```
dirty_conflict_instruction :=
  act_only_if_needed
+ name_loss
+ name_claimed_good
+ name_value_tilt
+ protect_least_reachable
+ preserve_minimum_floors_where_known
+ record_residue
+ assign_repair_duties
+ hold_or_escalate_when_not_safe_to_route
```

The next part turns from what should happen to who must answer.

Visibility is not blame. But visibility can create responsibility.

13 Part 10 — Responsibility After Visibility

Status: first draft. Incomplete. Not canon.

13.1 10.1 Visibility Is Not Blame

Seeing a harm path does not automatically make someone guilty.

This matters.

Mechanical Ethics is built to make hidden paths visible. That creates a risk. If every visible path becomes blame, people and institutions will learn not to see. They will avoid records, suppress warnings, narrow review, and punish those who bring bad news.

So the first rule of responsibility is restraint:

```
visibility != blame
```

Visibility is a condition for some forms of responsibility. It is not the whole of responsibility.

A person may see a risk without causing it. A worker may see a failure without having authority to alter it. A witness may see harm after meaningful correction has already become impossible. An institution may have a record somewhere but no functioning route for the relevant actor to know and act in time.

Blame requires more.

Mechanical Ethics therefore distinguishes seeing, warning, control, duty, omission, and fault. They are connected, but they are not identical.

The purpose is not to protect powerful actors from answerability. It is to keep answerability accurate enough to matter.

13.1.1 10.1.1 Visibility Includes Openings as Well as Closures

Visibility is not only visibility of harm.

An actor may also see that a path can be kept open, that an opportunity can still be preserved, that a route can be made reachable, that a warning can be made usable, that repair can still arrive in time, or that trust can be protected before it collapses.

Responsibility therefore attaches not only to seen closures, but sometimes to seen openings.

```
visible_transition_field :=
  seen_closure
+ seen_opening
+ seen_uncertainty
+ seen_clock
+ seen_affected_scope
```

This matters for the value spine. If Mechanical Ethics tracks viable becoming, then answerability concerns both what an actor damages and what an actor could reasonably preserve or reopen.

```
answerability_for_viable_becoming :=
  could_see_opening_or_closure
+ had_meaningful_control
+ had_time_or_route_to_act
-> duty_to_answer_for_response
```

This does not turn every missed good into blame. That would be dishonest and paralysing. It says only that some positive routes become responsibility-relevant once they are visible, reachable, and within role-shaped control.

13.2 10.2 Foresight Is Not Omniscience

No actor sees the whole future.

Responsibility cannot require omniscience. It can require attention to what was visible enough at the time.

```
foresight_standard :=
  decision_time_evidence
+ reasonable_visibility_for_role
+ known_or_knowable_path
```

A regulator is not responsible for every possible failure no one could have detected. But it may be responsible for ignoring repeated warnings. A designer is not responsible for every unforeseen use. But it may be responsible for foreseeable failure modes in the system it releases. A minister is not responsible for every local error. But they may be responsible for a policy that routes correction too slowly by design.

Foresight is role-shaped.

A doctor, engineer, parent, judge, platform operator, prison governor, teacher, model deployer, or public authority does not have the same visibility duty as a passer-by. Role creates expected attention. Power creates expected attention. Prior warning creates expected attention. Control over the path creates expected attention.

But foresight must remain tied to what could be known at the time.

Mechanical Ethics should not use later knowledge to pretend the actor knew more than they did.

That is hindsight blame.

The same limit applies to positive futures. An actor is not blameworthy merely because a better future can be imagined later. The question is whether the opening was reasonably visible under the actor's role, evidence, time, and capacity.

```
positive_foresight_limit :=  
  better_future_visible_later  
  != better_future_reasonably_actionable_then
```

13.3 10.3 Prediction Is Not Destiny

Prediction changes responsibility, but it does not create destiny.

A risk model may predict harm. A doctor may predict deterioration. A regulator may predict market failure. A teacher may predict exclusion. A parent may predict danger. An institution may predict delay. A system may predict fraud, violence, default, relapse, or failure.

Prediction can help action become earlier and more careful.

But prediction is not proof that the predicted thing must happen. Nor is it permission to treat a person as if they have already done what the model predicts.

```
prediction != destiny
```

This is especially important for artificial systems.

A predictive model may narrow a person's future-space before any human has understood the basis of the prediction. A risk score may become a sentence. A classification may become a route. A probability may become a wall.

Mechanical Ethics asks what the prediction does.

Does it preserve options? Does it trigger support? Does it invite human attention? Does it open contestability? Or does it close access, impose stigma, shift burden, and harden before correction?

```
prediction_audit :=  
  predicted_risk  
  + action_taken_because_of_prediction  
  + contestability  
  + reversibility  
  + affected_scope_future_space
```

Prediction can create duty to watch, support, warn, slow, or review.

It does not erase the difference between possibility and fact.

Prediction can also reveal a chance to protect viable becoming. If a system predicts failure and uses that signal to offer support, keep a route open, slow a harmful path, or preserve repairability, the prediction may expand future-space rather than close it.

```
prediction_as_support_signal :=  
  predicted_risk  
  + non_punitive_response  
  + support_or_review_route  
  + contestability_preserved  
  -> possible_viable_future_expansion
```

This is the difference between a warning that helps an entity continue and a label that traps it.

13.4 10.4 Decision-Time Evidence

Responsibility attaches to evidence available at the time of decision.

Not evidence invented later. Not evidence hidden from the actor. Not evidence that became obvious only after the loss. Not evidence that no one in the role could reasonably access.

```
decision_time_evidence :=  
  evidence_available_or_reasonably_available  
  before_or_during_relevant_action_window
```

This protects against hindsight blame, but it also prevents convenient ignorance.

An actor cannot avoid responsibility by designing a system that keeps evidence away. An institution cannot say decision-makers did not know if the institution's own structure prevented warnings

from reaching them. A model deployer cannot say the harms were invisible if it refused to monitor affected-scope outcomes. A public authority cannot rely on ignorance produced by underfunded challenge routes.

So the evidence question has two sides.

First: what did the actor actually know or have reason to know?

Second: did the actor have a duty to build a route by which relevant evidence could reach them?

```
evidence_route_duty :=  
  role_power  
  + foreseeable_harm  
  + capacity_to_monitor  
  -> duty_to_make_warning_reachable
```

Decision-time evidence includes warnings, complaints, anomalies, repeated patterns, near misses, known failure modes, and credible affected-scope accounts.

It does not include omniscience.

For value-spine purposes, the evidence route must carry positive signals too. A system that records only failures and never records lost opportunities, blocked repair, trust collapse, capacity damage, or reachable support will misread viable becoming.

```
evidence_route_value_scope :=  
  warnings  
  + complaints  
  + anomalies  
  + opportunity_signals  
  + repairability_signals  
  + trust_and_reachability_signals
```

13.5 10.5 Meaningful Control

Responsibility requires control.

Not total control. Not perfect control. Meaningful control.

```
meaningful_control :=  
  ability_to_change  
  or slow  
  or pause  
  or escalate  
  or warn  
  or preserve_evidence  
  or reduce_harm
```

A person with no ability to alter a path should not be treated as if they controlled it. A low-level worker may see a problem but lack authority. A child may cause harm without mature control. An actor under severe pressure may choose from a collapsed option set. A traumatised actor may have impaired control. A local official may operate inside a system designed elsewhere.

Control is graded.

But control is also often hidden.

Senior actors may claim they were too far from the operational path. Designers may claim users misused the system. Operators may claim policy required it. Vendors may claim clients deployed it. Committees may claim no single member decided. Automated systems may be treated as if they acted by themselves.

Mechanical Ethics asks where control actually sat.

Who could change the rule? Who could pause enforcement? Who could preserve evidence? Who could fund correction? Who could open a challenge route? Who could alter the incentive? Who could decide not to deploy? Who could monitor recurrence?

Responsibility follows meaningful control, not formal distance.

The value-spine question is: who could preserve or reopen viable becoming?

```
control_over_viable_becoming :=
  ability_to_preserve_opportunity
  or keep_route_reachable
  or restore_information
  or support_capacity
  or maintain_relation
  or reopen_repair_path
  or prevent_false_finality
```

If an actor had this kind of control and the relevant opening or closure was visible enough, responsibility may attach even where no direct personal malice exists.

13.6 10.6 Role Duty

Roles carry duties.

A role is not merely a title. It is a position in a path.

```
role_duty :=
  position_in_path
  + expected_attention
  + capacity_or_authority
  + affected_scope_dependency
```

Some roles exist because affected scopes cannot protect themselves alone. A judge, doctor, teacher, engineer, prison officer, social worker, regulator, auditor, designer, maintainer, parent, minister, platform operator, or data controller may hold power over another's future-space.

That power creates role duty.

Role duty does not mean every bad outcome becomes personal guilt. It means the role carries expected attention to certain kinds of harm, correction, and burden.

A maintainer may have a duty to monitor recurrence. A designer may have a duty to test foreseeable failure modes. A regulator may have a duty to preserve independent signal. A public authority may have a duty to keep contestability live. A platform may have a duty to provide meaningful appeal before account loss hardens. A parent may have a duty to protect a child's future-space before the child can protect it.

Roles are how systems act through people.

If the role has power but no duty, the system has built unanswerable authority.

Role duty also includes positive preservation where dependency exists. A role that controls another scope's route may owe more than non-injury. It may owe usable information, timely support, preserved opportunity, or a real return path.

```
role_positive_duty :=
  role_controls_route
  + affected_scope_depends_on_route
  + viable_becoming_at_stake
  -> duty_to_preserve_conditions_where_reasonably_possible
```

13.7 10.7 Delegation Is Not Erasure

Delegation moves work.

It does not erase responsibility.

```
delegation != erasure
```

A public body may delegate administration to a contractor. A hospital may use a triage tool. A court may rely on expert evidence. A school may use a behaviour platform. A company may outsource moderation. A state may rely on automated detection. A manager may assign review to a junior team.

Delegation may be legitimate. Complex systems need distributed work.

But if delegation creates opacity, delay, contaminated incentives, unreachable appeal, or responsibility gaps, then delegation becomes part of the harm mechanism.

Mechanical Ethics asks:

```
delegation_audit :=
  what_was_delegated
+ who_retained_power
+ who_can_correct
+ affected_scope_route
+ monitoring_duty
+ failure_responsibility
```

The delegating actor should not be able to say, “the contractor decided,” if the actor designed the contract, controlled the incentive, ignored warnings, or kept the benefit.

The operator should not be able to say, “the model decided,” if the operator chose to use the model, trusted it beyond its evidence, or failed to preserve challenge.

Delegation can distribute responsibility.

It cannot dissolve it.

The value-spine audit adds: who retained responsibility for viable becoming after delegation?

```
delegated_viable_becoming_audit :=
  who_preserves_notice
+ who_preserves_contestability
+ who_preserves_repairability
+ who_tracks_residue
+ who_can_reopen_path
```

A delegated path that expands efficiency while destroying reachability is not an innocent improvement. It is a value trade that must be named.

13.8 10.8 Automation Is Not Absolution

Automation changes the mechanism.

It does not absolve the humans and institutions that design, deploy, maintain, benefit from, or refuse to monitor it.

```
automation != absolution
```

Automated systems are powerful because they act quickly, consistently, and at scale. Those strengths can also make harm faster, less visible, and harder to contest.

A human decision can sometimes be interrupted by human doubt. An automated path may harden before anyone sees the individual case. A model can convert partial data into confident classification. A queue can turn delay into default. A dashboard can make a person’s future appear as a metric.

Mechanical Ethics asks whether automation preserves correction.

```
automation_constraint :=
  preserve_notice
+ preserve_contestability
+ preserve_human_answerability
+ monitor_outcomes
+ prevent_irreversible_action_from_unchecked_estimates
```

Automation may reduce bias, error, delay, or arbitrary discretion in some cases. It may expand future-space. It may help detect harm early. It is not automatically bad.

But automation must not be allowed to hide uncertainty inside technical authority or move harm faster than answerability can follow.

If the system acts faster than correction, it needs stronger brakes.

Automation can be part of a good path only if its gains do not erase answerability.

```
automation_positive_condition :=
  improves_reachability_or_support_or_detection
+ preserves_uncertainty_visibility
+ preserves_contestability
+ preserves_human_answerability
+ records_remaining_burden
```

Local efficiency is not enough. The question is whether automation helps affected scopes navigate, correct, repair, and continue.

13.9 10.9 Responsibility Attachment

Responsibility attachment is the point where the record says: this actor must answer for this path in this way.

It is not always blame. It may be duty to disclose, duty to repair, duty to monitor, duty to compensate, duty to escalate, duty to preserve evidence, duty to prevent recurrence, or duty to explain why action continued.

```
responsibility_attachment_factors :=
  visibility
+ decision_time_evidence
+ action_relevance
+ role_connection
+ meaningful_control
+ capacity_to_act
+ delay_or_omission
+ delegation_or_maintenance
```

Attachment should be specific.

Not: someone should have done something.

Instead: this actor, in this role, with this evidence, had this control, during this time window, and owed this response.

Specificity prevents scapegoating. It also prevents evasion.

A junior clerk should not absorb the guilt of a system designed above them. A minister should not hide behind frontline staff when policy created the path. A developer should not absorb every deployment choice made by a client. A client should not hide behind a vendor when it chose the system and benefited from it. A committee should not hide behind collective fog if roles and records show who could act.

Responsibility attachment is how Mechanical Ethics keeps answerability from becoming either witch hunt or mist.

The value-spine version of responsibility attachment asks what the actor could see and affect in relation to viable becoming.

```
responsibility_for_viable_becoming :=
  actor
+ role
+ decision_time_evidence
+ meaningful_control
+ visible_opening_or_closure
+ duty_response
```

The response may be to stop a closure, preserve an opening, reopen a route, disclose uncertainty, return burden, repair residue, or explain why none of those were possible.

13.10 10.10 Responsibility Without Scapegoating

Scapegoating is false repair.

A system finds one person to blame, sanctions them, and leaves the machinery intact. The public receives a story. The institution receives closure. The path remains.

```
scapegoating :=
  individual_blame
+ unchanged_harm_mechanism
+ institutional_closure
```

Mechanical Ethics rejects scapegoating for the same reason it rejects correction without repair. It does not touch the path.

This does not mean individuals are never responsible. They are. Some people misuse power, ignore warnings, falsify records, hide evidence, punish dissent, or choose serious wrongdoing. Individual blame may be warranted.

But individual blame should not be used to erase systemic responsibility where the system selected, rewarded, permitted, or protected the harmful path.

The better record asks:

```
responsibility_map :=
  individual_actions
+ role_duties
+ system_design
+ incentives
+ warnings
+ correction_failures
+ residue
```

Responsibility without scapegoating is harder than blame and harder than excuse.

It says: do not blame from visibility alone, but do not let complexity become weather. Do not sanction one visible actor while preserving the hidden machine. Do not call a system innocent because no single hand contains the whole path.

It must also avoid positive scapegoating: giving credit to one visible reform while leaving the path unrepaired. A system can celebrate a helpful actor, new tool, new policy, or new programme while the deeper route remains unreachable.

```
positive_scapegoating_risk :=
  visible_good_actor_or_reform
+ unchanged_path
+ institution_claims_moral_closure
```

Do not use blame to hide the machine. Do not use praise to hide the machine either.

Part 10 ends with the responsibility spine:

```
responsibility_spine :=
  visibility_is_not_blame
+ foresight_is_not_omniscience
+ prediction_is_not_destiny
+ decision_time_evidence
+ meaningful_control
+ role_duty
+ delegation_not_erasure
+ automation_not_absolution
+ responsibility_for_seen_openings_and_closures
+ anti_scapegoating
```

The next part moves into the most dangerous territory: power, legitimacy, authority, desert, dignity, and severe constraint.

That part must frame more than it solves.

14 Part 11 — Power, Legitimacy, and Coercion

Status: first draft. Incomplete. Not canon.

Warning: this part frames more than it solves. It contains high-risk value-cruxes: legitimacy, coercion, punishment, desert, dignity, force, and acceptable sacrifice. These must not be treated as settled because the book has reached them.

14.1 11.1 Power That Changes Futures

Power is the capacity to change another scope's future-space.

That definition is deliberately mechanical. It does not yet say whether the power is lawful, legitimate, wise, abusive, necessary, or kind.

```
power :=
  capacity_to_change
  another_scope_future_space
```

Power may be physical, legal, economic, informational, social, technical, administrative, military, parental, medical, educational, algorithmic, or institutional. It may act quickly or slowly. It may act openly or invisibly. It may act through force, rules, defaults, incentives, classifications, queues, permissions, records, or delays.

A system has power when it can open or close paths for others.

A person with a weapon has power. A state with police power has power. A landlord has power over housing. A platform has power over reach. A school has power over a child's record. A hospital has power over access to care. A model has power when its outputs become routes. A bureaucracy has power when its timing decides whether correction remains meaningful.

Mechanical Ethics does not begin by saying all power is bad.

Without power, no one could protect, repair, coordinate, teach, build, rescue, govern, or care. The question is not whether power exists. The question is whether power remains answerable to the future-space it changes.

```
power_question :=
  whose_future_space_can_be_changed
  + by_whom
  + through_what_mechanism
  + under_what_constraint
  + with_what_correction_route
```

The value-spine version is this:

```
power_over_viable_becoming :=
  capacity_to_shape
  continuation
  + orientation
  + capacity
  + relation
  + repairability
  + meaningful_future_space
  for_another_scope
```

This is why power is ethically dense. It does not merely move resources or enforce rules. It can shape what another scope can become. It can preserve a route, open a path, teach, protect, heal, coordinate, and repair. It can also trap, confuse, isolate, humiliate, burden, silence, or close a path before the affected scope can answer.

```
power_can_support :=
  protection
  + coordination
  + repair
  + learning
  + care
  + trust_infrastructure
  + shared_future_space

power_can_damage :=
  closure
  + dependency_capture
```

- + false_finality
- + unreachable_contest
- + trust_damage
- + unrepairable_burden

Power that cannot answer becomes domination.

14.2 11.2 Legitimacy as Answerable Authority

Legitimacy is one of the hardest words in this book.

It is tempting to define legitimacy as lawful authority. That is too narrow. Some lawful authority is abusive. Some morally necessary action may occur before law has caught up. It is also tempting to define legitimacy as good outcome. That is too loose. Good outcomes can be used to excuse unanswerable power.

Mechanical Ethics should not pretend to solve legitimacy completely here.

It can frame a candidate structure:

```
legitimacy_candidate :=
  authority_basis
  + affected_scope_recognition
  + reason_giving
  + contestability
  + burden_accounting
  + correction_capacity
  + expiry_or_review
  + independent_challenge
```

This is not a final theory of political legitimacy. It is an anti-laundersing structure.

Authority becomes more legitimate when it names its basis, recognizes affected scopes, gives reasons, preserves contestability, accounts for burden, remains correctable, expires or faces review where appropriate, and allows independent challenge.

Authority becomes more suspect when it hides its basis, erases affected scopes, refuses reasons, blocks challenge, routes burden downward, makes correction unreachable, extends itself without review, or judges only its own conduct.

For the value spine, legitimacy also asks whether authority preserves the conditions under which affected scopes can orient, contest, repair, and continue.

```
legitimacy_value_candidate :=
  authority_changes_viable_becoming
  + authority_names_basis
  + affected_scope_can_orient
  + affected_scope_can_contest_where_possible
  + burden_and_residue_recorded
  + correction_remains_live
```

The anti-circularity rule is strict:

```
legitimacy_cannot_mean :=
  authority_is_legitimate
  because_legitimate_authority_used_it
```

That is empty. Power cannot validate itself by naming itself valid.

Legitimacy must show its route.

14.3 11.3 Coercion Under Uncertainty

Coercion is action that constrains another scope's options, reachability, or conduct through force, threat, dependency, or structured lack of alternative.

```
coercion :=
  imposed_constraint
  on_scope_options_or_reachability
```

under_pressure_or_threat_or_dependency

Some coercion is plainly abusive. Some is ordinary and lawful. Some may be necessary to prevent greater harm. A quarantine order, arrest, child protection intervention, evacuation, treatment refusal, platform suspension, school exclusion, emergency closure, or ecological control may all involve coercive elements.

Mechanical Ethics cannot simply say “no coercion.”

But coercion under uncertainty carries special danger. The actor may be wrong about the threat. The affected scope may be unable to contest. The loss may harden before correction. The coercive system may benefit from treating its own estimate as final.

So coercion needs stronger audit:

```
coercion_audit :=
  threat_or_reason
+ affected_scope
+ evidence_status
+ uncertainty_status
+ less_coercive_options
+ timing_pressure
+ contestability
+ expiry_or_review
+ residue
```

Coercion may sometimes be justified. But coercion should not be cleanly justified by the fact that authority found it convenient, efficient, traditional, or lawful.

The more coercive the action, the more serious the duty to preserve correction and record residue.

The value-spine question is whether the constraint protects viable becoming, damages it, or both.

```
coercion_value_audit :=
  what_closure_is_imposed
+ what_viable_becoming_is_claimed_to_be_protected
+ what_viable_becoming_is_damaged
+ whether_less_closing_route_exists
+ what_repair_or_residue_follows
```

A constraint may be defensible because it prevents faster or wider closure. But if it does so, the record must still name the closure it imposes.

14.3.1 11.3.1 Hard-No Constraints for Coercion and Force

This part does not solve legitimacy. But it can still name some hard-no constraints.

```
coercion_hard_no_constraints :=
  no_coercion_from_erased_scope
+ no_irreversible_coercion_from_unchecked_contaminated_estimate
+ no_force_without_named_authority_basis
+ no_continuing_force_after_basis_expires
+ no_emergency_power_without_expiry_or_review
+ no_punishment_from_visibility_alone
+ no_coercion_that_requires_false_cleanliness
```

These constraints are negative boundaries, not a positive theory of justified force.

They say that coercive power fails before the harder legitimacy question is even reached if it depends on erased affected scopes, unchecked contaminated evidence, unnamed authority, expired basis, emergency without expiry, visibility-as-blame, or cleaned dirty loss.

If one of these hard-no constraints is triggered, the likely router finding is HALT, HOLD, or ESCALATE, not PROCEED_ONLY_WITH_NAMED_CONSTRAINTS, unless an explicit dirty-conflict emergency chain explains why immediate constrained action is still required.

The burden is on the actor using power to show the route. The burden is not on the affected scope to prove that invisible power is illegitimate.

14.4 11.4 Emergency Power

Emergency compresses time.

Some emergencies are real. Fire, flood, violence, epidemic, collapse, acute medical risk, safeguarding danger, ecological crisis, or infrastructure failure may require fast action before ordinary process can operate fully.

Mechanical Ethics does not deny emergency.

It denies emergency as amnesia.

```
emergency_power :=  
  extraordinary_constraint  
  justified_by_time_compressed_threat  
  requiring_expiry_review_and_residue_record
```

Emergency power should answer at least four questions.

First: what is the immediate threat?

Second: what ordinary safeguards are being suspended, narrowed, or delayed?

Third: what expiry, review, or correction channel remains?

Fourth: did prior failure help produce the emergency?

The fourth question matters because institutions often create or worsen emergency through neglect. Maintenance is ignored. Warnings are dismissed. capacity is underbuilt. People are left without early routes. Then the crisis arrives, and extraordinary power is claimed as if it came from nowhere.

Mechanical Ethics says emergency may justify speed, but it does not erase history.

```
emergency_must_not_launders :=  
  prior_warning  
  + prior_neglect  
  + underbuilt_correction  
  + avoidable_hardening
```

Emergency action may be necessary and still dirty.

The record should be strong enough to say both.

Emergency may also protect viable becoming where delay would close it. That is real. A rescue, evacuation, intervention, or temporary restriction may preserve life, capacity, trust, or future repairability. But this claimed good must not erase the burden imposed by the emergency route.

```
emergency_value_record :=  
  what_viable_becoming_was_preserved  
  + what_scope_was_constrained  
  + why_time_pressed  
  + what_review_or_expiry_remains  
  + what_residue_follows
```

14.5 11.5 Punishment and Desert

Punishment and desert are not solved here.

They are too dangerous to smuggle in through the side door.

Desert asks what someone deserves because of what they did, intended, risked, ignored, or became responsible for. Punishment asks whether and how a system may intentionally impose burden, pain, deprivation, stigma, restriction, or loss in response.

Mechanical Ethics has tools relevant to punishment: visibility, decision-time evidence, role duty, meaningful control, responsibility attachment, residue, repair, prevention, and legitimacy. But those tools do not by themselves settle desert.

```
punishment_desert_status :=  
  UNRESOLVED_VALUE_CRUX  
  unless_structural_chain_and_normative_authority_are_explicit
```


This part can still set constraints.

Punishment must not be based on visibility alone.

Punishment must not treat prediction as destiny.

Punishment must not use one visible person to hide system design.

Punishment must not pretend to repair if it only satisfies anger or institutional closure.

Punishment must not continue after its authority, purpose, or evidence has expired.

Punishment must preserve correction where error is possible.

```
punishment_constraints :=  
  no_visibility_as_blame  
  + no_prediction_as_destiny  
  + no_scapegoating  
  + no_repair_theatre  
  + correction_route_required  
  + expiry_or_review_required_where_power_continues
```

These are constraints, not a full theory.

The book must be honest: it can narrow what punishment may not hide. It has not yet proven what punishment positively deserves.

For the value spine, punishment also must not pretend that imposed burden is automatically repair. If punishment is claimed to protect future viable becoming, that claim needs a route: prevention, restoration, accountability, safety, rehabilitation, recognition, or recurrence reduction. If no such route is shown, punishment is only a burden with a story.

```
punishment_value_audit :=  
  burden_imposed  
  + claimed_future_protection_or_repair  
  + evidence_for_that_route  
  + correction_if_wrong  
  + residue_record
```

14.6 11.6 Dignity and Recognition

Dignity is also dangerous to define too quickly.

It is easy to use dignity as a noble word with no mechanism. It is also easy to avoid dignity because it resists measurement.

Mechanical Ethics can approach dignity through recognition.

A scope is denied recognition when it is treated as if its future-space, account, suffering, agency, dependency, or residue need not appear in the record.

```
recognition_failure :=  
  affected_scope_exists  
  + affected_scope_materially_changed  
  + record_or_process_erases_that_scope
```

This is not all that dignity means. But it is one mechanical way dignity is violated.

A person made to navigate an opaque system that does not hear them is denied recognition. A prisoner treated only as risk may be denied recognition. A patient treated only as throughput may be denied recognition. A worker treated only as metric may be denied recognition. An animal treated only as object may be denied recognition where sentience or agency is relevant. A community treated only as obstacle may be denied recognition.

Recognition does not mean the scope always gets what it wants.

It means the scope's presence, account, and future-space cannot be erased for administrative convenience.

```
dignity_minimum_candidate :=
  do_not_reduce_scope_to_system_variable
  + preserve_account_where_material
  + give_reasons_where_power_acts
  + preserve_contestability_where_possible
```

The value-spine version adds that a scope should not be reduced to a tool, obstacle, risk score, burden, case number, output, or optimization variable when its viable becoming is materially at stake.

```
dignity_as_viable_becoming_recognition :=
  scope_not_reduced_to_variable
  + account_preserved_where_material
  + viable_becoming_not_erased
  + reasons_given_where_power_shapes_path
  + route_back_preserved_where_possible
```

This remains a candidate. It should be tested, not canonized.

14.7 11.7 Authority That Expires

Power should not always continue just because it was once justified.

Authority can expire.

```
authority_expiry :=
  authority_basis_no_longer_matches
  threat_or_need_or_role_or_consent_or_evidence
```

Emergency authority may expire when the emergency passes. Coercive authority may expire when threat changes. Parental or custodial authority may change as capacity develops. Administrative authority may expire when a decision is corrected. Data-use authority may expire when consent, purpose, or necessity ends. Model deployment authority may expire when drift, failure, or contamination becomes visible.

Systems resist expiry. Power tends to preserve itself. Temporary measures become routine. Exceptional tools become normal. Safeguards suspended for crisis are not restored. Data gathered for one purpose becomes available for another. A risk label remains after risk changes.

Mechanical Ethics treats expiry as a legitimacy condition.

```
expiry_check :=
  original_basis
  + current_basis
  + changed_facts
  + affected_scope_burden
  + review_route
```

Authority should have memory of why it began and conditions for when it must narrow, stop, or return.

If power cannot state how it expires, it is already suspect.

Expiry protects viable becoming because ongoing power keeps shaping the field after its original basis has passed.

```
expiry_value_function :=
  prevent_power_from_becoming_default_world
  + return_future_space_where_basis_has_expired
  + reopen_contestability
  + reduce_unneeded_dependency
```

14.8 11.8 When Force May Be Used

This book should not pretend to give a complete theory of force.

Force may be physical, legal, economic, administrative, technical, or informational. It may restrain, compel, exclude, seize, confine, remove, disable, or override.

The question of justified force is a high-risk value-cruX.

```
force_status :=  
  UNRESOLVED_VALUE_CRUX  
  unless_case_specific_chain_is_shown
```

Mechanical Ethics can frame conditions that any claimed force should answer:

```
force_audit :=  
  threat_or_need  
  + authority_basis  
  + affected_scope_recognition  
  + evidence_status  
  + less_forceful_options  
  + proportionality_or_floor_issue  
  + contestability_where_possible  
  + expiry_or_review  
  + residue
```

Force becomes more suspect when the threat is vague, the evidence is contaminated, the affected scope cannot contest, less forceful routes were ignored, the actor benefits from force, or power continues after the reason expires.

Force may become more defensible where threat is immediate, evidence is strong, delay would close protection, less forceful routes are unavailable, and review or residue accounting remains.

But this is not a permission machine.

A force audit does not make force clean.

It shows what must be answered before force can claim legitimacy.

The value-spine audit asks whether force preserves viable becoming from a greater closure, and whose viable becoming it closes in doing so.

```
force_value_audit :=  
  claimed_protected_future  
  + imposed_closure  
  + least_closing_available_route  
  + review_or_expiry  
  + residue
```

14.9 11.9 When Force Must Stop

Stopping force matters as much as authorizing it.

A system that can begin force but cannot stop it is structurally dangerous.

```
force_must_stop_when :=  
  threat_basis_expires  
  | evidence_collapses  
  | less_forceful_route_becomes_available  
  | correction_finds_error  
  | residue_or_burden_becomes_disproportionate  
  | authority_expiry_condition_met
```

This applies beyond physical force.

A suspension should stop when the risk basis collapses. A data restriction should stop when purpose expires. A coercive placement should stop when conditions no longer justify it. An emergency power should stop when emergency conditions pass. An enforcement route should pause when correction becomes materially credible. A model-based exclusion should stop when the estimate is shown contaminated or unfit.

Stopping is hard because force creates institutional momentum. Once a person is labelled, confined, excluded, flagged, suspected, or routed, the burden often shifts onto them to prove the force should end.

Mechanical Ethics resists that burden shift.

```
continuing_force_requires :=
  renewed_basis
  + review
  + affected_scope_route
  + residue_check
```

The fact that force was once justified does not prove it remains justified.

If force continues after it no longer protects viable becoming, it becomes closure for its own sake.

14.10 11.10 What This Part Has Not Solved

Part 11 must end with restraint.

It has framed power as capacity to change future-space and viable becoming. It has offered a candidate structure for legitimacy as answerable authority. It has described coercion under uncertainty, hard-no constraints for coercion and force, emergency power, punishment constraints, dignity through recognition, authority expiry, and force audits.

It has not solved political legitimacy.

It has not solved punishment.

It has not solved desert.

It has not solved dignity.

It has not solved all justified coercion.

It has not solved acceptable sacrifice.

```
Part_11_status :=
  framing_not_solving
  + high_UNRESOLVED_VALUE_CRUX_risk
  + anti_laundering_constraints_only
```

This is not failure. It is honesty.

The book has enough structure to say that power must answer for the future-space and viable becoming it changes. It does not yet have enough to claim final authority over every use of power.

The next part shifts scale.

If individual action and institutional power can narrow or preserve viable future, then civilization itself can be understood as an attempt to remain correctable under uncertainty.

15 Part 12 — Civilization Under Uncertainty

Status: first draft. Incomplete. Not canon.

15.1 12.1 Civilization as Correctability

Civilization is not only order.

Order can be brutal. Order can be false. Order can preserve a harmful path because the path is legible to power. A prison can be orderly. A bureaucracy can be orderly. A market can be orderly while routing ruin to those who cannot contest. A machine can be orderly while narrowing lives at scale.

Mechanical Ethics needs a thinner and harder definition.

```
civilization :=
  organized_attempt
  to_remain_correctable_under_uncertainty
  while_preserving_future_possibility
```

across_many_scopes

This is a candidate definition, not a final sociology.

The point is that civilization is not merely having laws, markets, institutions, roads, rituals, archives, schools, hospitals, or machines. It is whether those systems preserve the possibility of correction while acting under uncertainty.

A civilization must act. It must build, allocate, defend, teach, trade, heal, punish, repair, remember, and decide. It cannot wait for perfect knowledge. But if it acts faster than it can correct, and if the losses it creates become invisible, it becomes dangerous even when it calls itself advanced.

Correctability is not weakness.

It is how power remains answerable after discovering that it was wrong.

15.1.1 12.1.1 Civilization as Shared Viable Becoming

Correctability alone is not the whole civilization claim.

A civilization is also an organized attempt to preserve and expand viable becoming across many scopes without pretending the costs disappear.

```
civilization_value_candidate :=  
  shared_correctability  
  + shared_viable_becoming  
  + honest_burden_record  
  + residue_memory  
  under_uncertainty
```

This does not mean every scope gets everything it needs. No civilization can deliver that. Scarcity, conflict, error, loss, and dirty tradeoffs remain.

It means a civilization should be judged partly by whether its systems help more scopes continue, orient, learn, relate, repair, and meaningfully act over time, while preserving routes for those harmed by the system to appear and answer back.

```
shared_viable_becoming :=  
  many_scopes  
  + preserved_continuation  
  + usable_orientation  
  + capacity_growth  
  + relation_and_trust  
  + repairability  
  + meaningful_future_space
```

Civilization fails when it produces order by shrinking the field of becoming for those least able to contest. It also fails when it produces growth by hiding the bodies, burdens, exhausted workers, poisoned places, broken trust, or future repair debts that make the growth possible.

The stronger claim is not that civilization equals goodness. Civilization can become a harm carrier. The claim is that a civilization worth defending must remain answerable to the viable becoming it opens and closes.

15.2 12.2 Institutions as Memory

Institutions are memory carriers.

They preserve rules, records, roles, promises, procedures, warnings, debts, rights, duties, classifications, and routes for correction. They allow action to continue beyond one person's attention.

```
institution :=  
  durable_role_memory  
  + action_capacity  
  + record_system  
  + correction_or_failure_path
```

An institution can remember well or badly.

Good memory preserves contradiction, evidence, reasons, and repair duties. Bad memory preserves only the official story. Good memory makes past failure available for future prevention. Bad memory buries failure under closure. Good memory keeps affected scopes visible. Bad memory preserves metrics while forgetting people.

Institutions can also forget by design.

They can fragment records. They can delete warnings. They can rotate staff without transferring responsibility. They can outsource memory to vendors. They can treat complaints as isolated. They can count closed cases rather than repaired paths. They can store data while losing meaning.

Mechanical Ethics treats institutional memory as an ethical mechanism.

```
institutional_memory_question :=  
  what_is_preserved  
  + what_is_erased  
  + who_can_access_record  
  + can_record_trigger_correction
```

Case anchor: a hospital that records waiting-time compliance but not missed diagnostic windows may appear to remember process while forgetting harm. The repair question is not whether the database exists. It is whether the record can trigger correction before the next patient carries the same failure.

A civilization that cannot remember harm cannot learn from it.

A civilization that remembers only authority cannot correct authority.

The value-spine version adds that institutions must remember positive routes too: what preserved trust, what kept opportunity open, what made repair reachable, what supported capacity, and what prevented a path from collapsing.

```
institutional_positive_memory :=  
  what_kept_route_open  
  + what_supported_capacity  
  + what_preserved_trust  
  + what_reduced_recurrence  
  + what_repaired_path
```

If an institution records only failure, it becomes defensive. If it records only success, it becomes vain. It needs memory that can learn from both wound and repair.

15.3 12.3 Law as Slowed Power

Law can be understood as power slowed enough to become answerable.

That is not all law is. Law is also language, threat, ritual, settlement, coercion, promise, public memory, and institutional power. But one of law's central functions is to slow force.

```
law_candidate :=  
  slowed_power  
  + public_reason  
  + contestability  
  + precedent_memory  
  + correction_route
```

A lawful system should make power less impulsive. It should require reasons. It should preserve records. It should let affected scopes contest. It should bind officials. It should create routes for correction. It should keep force from being only the will of the stronger actor.

Law fails when it becomes only the shape of power.

A legal process can be too slow to repair. It can be too expensive to reach. It can be formally equal and practically closed. It can preserve official contradiction badly. It can let authority set every clock. It can give the appearance of correction while loss hardens outside the courthouse.

Case anchor: an appeal route after eviction may satisfy legal form while failing practical correction. The law may still be valid law, but if the person has already lost home, health, records, or trust, the route has failed the timing test.

Mechanical Ethics does not replace law.

It asks whether law is still doing the mechanical work that makes law ethically serious.

```
legal_route_test :=
  reachable
+ timely
+ reasoned
+ independent_enough
+ able_to_change_path
+ residue_aware
```

If law cannot arrive while it still matters, it may remain valid law but fail as correction.

The positive version is that good law does not merely stop bad power. It creates reliable routes by which many scopes can orient, trust, contest, plan, repair, and cooperate.

```
law_positive_function :=
  make_power_predictable_enough_to_answer
+ make_reasons_public
+ preserve_contest
+ stabilize_trust
+ reopen_path_when_error_found
```

Law is ethically serious when it makes shared future-space more usable without hiding whose path it narrows.

15.4 12.4 Markets and Local Reward Pull

Markets are selection systems.

They reward, punish, repeat, amplify, starve, and redirect. They can coordinate information and labour at enormous scale. They can create prosperity, innovation, adaptation, and choice. They can also train systems to ignore loss that does not appear in price, contract, liability, or metric.

```
market :=
  distributed_selection_system
  shaped_by_reward_and_constraint
```

Mechanical Ethics should neither romanticize nor demonize markets.

The useful question is what the market is training.

If profit comes from expanding real future-space, reducing waste, improving access, and preserving correction, market reward may align with value. If profit comes from opacity, dependency, extraction, attention capture, externalized harm, regulatory delay, or burden shifted to those least able to contest, the market learns the wrong thing.

```
market_audit :=
  local_reward
+ hidden_cost
+ affected_scope
+ correction_capacity
+ externalized_residue
```

Case anchor: a platform may be rewarded for engagement while anger, addiction, misinformation, or reputational harm is routed outside the metric. The issue is not that reward exists. The issue is what the reward trains and who carries the residue.

Market language often cleans harm by calling it preference, price, demand, efficiency, or consumer choice. Sometimes those words are accurate. Sometimes they hide coercion, asymmetry, ignorance, dependency, or absence of real alternatives.

Mechanical Ethics asks whether choice is real enough, whether loss is priced or hidden, whether correction is reachable, and whether local reward is preserving or destroying future possibility.

Markets are powerful because they learn.

The question is whether they are learning the right thing.

The value-spine audit is direct:

```
market_value_audit :=
  what_local_reward_trains
+ whose_viable_becoming_expands
+ whose_viable_becoming_contracts
+ who_can_exit_or_contest
+ what_residue_is_externalized
```

A market path can be good when it creates capability, access, resilience, useful information, and mutual benefit while keeping correction possible. A market path becomes suspect when its apparent efficiency depends on invisible exhaustion, degraded attention, locked-in dependency, or costs pushed onto scopes that cannot price or contest them.

15.5 12.5 Public Systems and Burden Routing

Public systems decide who carries uncertainty.

Benefits, housing, schools, policing, hospitals, courts, immigration, child protection, prisons, taxation, welfare, licensing, infrastructure, environmental regulation — these systems route burden every day.

```
public_system :=
  collective_power
+ administrative_route
+ affected_scope_dependency
+ correction_obligation
```

The public system may need rules. It may need thresholds. It may need fraud control, triage, budgets, deadlines, sanctions, and prioritization. Mechanical Ethics does not deny administrative necessity.

But necessity does not answer burden routing.

If the system is uncertain, who waits?

If the record is wrong, who proves it?

If the decision is delayed, who absorbs the cost?

If the appeal is slow, who carries the risk?

If the model errs, who finds the error?

If the official language is incomprehensible, who pays for translation into reality?

```
burden_routing_question :=
  whose_uncertainty
+ whose_cost
+ whose_clock
+ whose_proof
+ whose_repair
```

Case anchor: if a benefit system flags a claimant as risky and stops payment before evidence is checked, the public system has routed its uncertainty onto the person least able to carry delay. The correction route is not adequate merely because it exists later.

Public systems become ethically dangerous when they treat private suffering as the cost of public order while refusing to record that cost.

The minimum demand is not that public systems never burden people.

It is that they do not hide the burden, do not route it by default to the least able to carry it, and do not call correction meaningful when correction arrives too late.

Public systems also carry positive duty. They are often the routes through which vulnerable scopes gain access, information, stability, protection, capacity, and repair.

```
public_system_positive_duty :=
  keep_route_reachable
+ make_information_usable
+ preserve_support_before_loss_hardens
+ route_burden_to_capacity
+ repair_when_state_path_failed
```

A public system is not good because it is compassionate in language. It is good when its routes actually preserve viable becoming under pressure and remain answerable when they fail.

15.6 12.6 Trust Under Pressure

Trust is not belief that nothing will go wrong.

Trust is belief that when something goes wrong, the system will not hide it, punish the witness, erase the affected scope, or refuse correction.

```
trust :=
  honesty_under_pressure
+ correction_capacity
+ non_punishment_of_valid_warning
+ residue_acknowledgement
```

High-trust systems are not systems without error. They are systems where error can surface without being crushed.

A person trusts a doctor not because doctors are never wrong, but because medicine should preserve evidence, explanation, second opinion, and correction. A citizen trusts a court not because courts never err, but because reasons, appeal, publicity, and record should constrain error. A user trusts a platform not because moderation is perfect, but because notice and appeal should exist before serious loss hardens. A worker trusts an employer not because conflict never happens, but because process should not hide retaliation or record scars.

Case anchor: a whistleblower policy that punishes the first person to make the warning usable destroys trust even if later reform is announced. Trust depends on whether correction can surface under pressure, not whether the institution praises honesty after the scandal.

Trust decays when systems protect themselves from valid correction.

Once people learn that warning is punished, they stop warning. Once they learn appeal is theatre, they stop believing in routes. Once they learn apology means closure without repair, they stop treating institutional language as serious.

Trust is future-space. It changes what people can risk, say, build, report, and attempt.

A civilization that spends trust faster than it repairs trust is narrowing its own future.

The positive value is larger:

```
trust_as_civilization_infrastructure :=
  reliable_return_path
+ usable_warning_channel
+ repair_after_failure
+ confidence_to_cooperate
+ capacity_to_risk_learning
```

Trust lets scopes attempt things they would not attempt inside fear. It is not sentiment. It is infrastructure for shared viable becoming.

15.7 12.7 Maintenance, Stewardship, and Decay

Systems decay.

Not only machines. Institutions, norms, records, legal routes, review bodies, public services, markets, professions, families, platforms, and moral languages decay.

They decay when maintenance is ignored. They decay when incentives drift. They decay when old repairs become new rituals. They decay when memory becomes archive without action. They decay when nobody owns follow-up. They decay when emergency measures become routine.

```
decay :=
  correction_capacity_decreases_over_time
  unless_stewardship_active
```

Maintenance is ethical work because it preserves future correction.

A bridge maintained is not dramatic. A record system kept clean is not heroic. A complaint route funded before scandal is not celebrated. A model monitored before harm is visible may look like overhead. A school preserving ordinary support may not appear as innovation.

Case anchor: a complaint route that once worked can become useless through underfunding, staff churn, language drift, or vendor change. If no one owns maintenance, yesterday's repair becomes tomorrow's theatre.

But much harm begins where maintenance was treated as optional.

Stewardship means holding a system across time so that its repair capacity does not collapse after attention moves on.

```
stewardship_duty :=
  monitor
  + remember
  + repair
  + update
  + preserve_challenge
  + prevent_decay
```

Civilization is not only built in moments of invention. It is kept alive by maintenance that prevents correctability from rotting.

Maintenance also preserves positive becoming. It keeps ordinary supports working before crisis. It protects the background routes that make play, learning, care, work, mobility, health, family, contest, and repair possible.

```
maintenance_positive_function :=
  preserve_background_conditions
  + keep_supports_ordinary
  + prevent_crisis_from_consuming_future
  + make_good_paths_reliable
```

15.8 12.8 Cooperation and Shared Worlds

Civilization is not only harm control. It is also cooperation at scale.

Many of the most valuable human goods are shared-world goods. Language, medicine, law, science, roads, schools, sanitation, music, mutual aid, care systems, safety norms, open knowledge, trustworthy records, and ordinary public peace cannot be produced by isolated self-protection alone.

```
shared_world_good :=
  many_scopes
  + coordinated_action
  + mutual_reliance
  + preserved_exit_or_contest
  + shared_future_space_expansion
```

Cooperation is not automatically good. A cartel cooperates. A corrupt institution cooperates. A predatory group coordinates. Cooperation becomes ethically serious when it expands viable becoming without requiring erased burdens, false consent, or blocked correction.

```
cooperation_value_test :=
  what_shared_capacity_expands
  + who_is_burdened_by_coordination
```

```
+ who_can_exit_or_contest
+ what_residue_remains
+ whether_trust_is_built_or_consumed
```

This is where positive civilization belongs. Not as decoration. Not as moral cheerleading. As structure.

A good institution makes more good action reachable. A good public route makes help less humiliating and correction less heroic. A good market lets useful work find need without hiding ruin. A good law makes power slower and reasons usable. A good school expands capacity while preserving dignity and repair. A good medical system preserves life and trust while admitting uncertainty.

Civilization becomes real when good paths are not rare acts of heroism but ordinary routes.

```
civilization_positive_test :=
  good_paths_made_reachable
+ repair_paths_made_ordinary
+ trust_not_spent_faster_than_rebuilt
+ burden_visible
+ correction_alive
```

15.9 12.9 Progress Without Sacrifice Laundering

Progress is real.

Medicine improves. Violence can decline. Knowledge can expand. Rights can be recognized. Technology can open access. Institutions can learn. Markets can lift material conditions. Movements can expose hidden harm. Artificial systems may help detect risk, translate knowledge, or expand capability.

Mechanical Ethics is not anti-progress.

It is anti-laundering.

```
progress_claim_requires :=
  expanded_future_space
+ named_costs
+ affected_scope_record
+ correction_capacity
+ residue_accounting
```

A progress story becomes dangerous when it treats sacrificed scopes as the price of history and then erases the receipt.

A road may be progress and still displace. A medical trial may be progress and still exploit. A model may be progress and still burden false positives. A transition to green energy may be progress and still route cost onto workers, communities, ecosystems, or future waste sites. A security system may reduce some harms and create others.

Mechanical Ethics asks for progress with memory.

If progress narrows some futures to expand others, say so. If the loss is justified, show the route. If the loss is dirty, record residue. If repair is owed, assign it. If the value tilt is contested, do not hide it behind inevitability.

Progress must also be tested against viable becoming, not raw novelty or scale.

```
progress_value_test :=
  whose_viable_becoming_expands
+ whose_viable_becoming_contracts
+ what_uncertainty_remains
+ what_repair_stays_live
+ what_residue_is_carried_forward
```

Civilization under uncertainty cannot mean never causing loss.

It can mean refusing to become the kind of system that needs its losses hidden in order to believe in itself.

Part 12 ends here:

```
civilization_test :=  
  can_the_system_act  
  + remain_correctable  
  + preserve_future_possibility  
  + support_shared_viable_becoming  
  + remember_residue  
  + resist_self_absolution
```

The next part turns to artificial systems and future minds, where these pressures become faster, more opaque, and less settled.

16 Part 13 — Artificial Systems and Future Minds

Status: first draft. Incomplete. Not canon.

16.1 13.1 Artificial Systems as Acting Systems

Artificial systems do not need moral standing to matter ethically.

That is the first separation.

A model, platform, automation layer, robot, agent, classifier, recommender, routing system, sensor network, or decision support tool may change human and non-human future-space without being a moral patient itself.

```
AI_as_actor_or_tool :=  
  system_changes_future_space  
  + duties_attach_to_designers_deployers_operators_maintainers
```

The first question is not, “Does the system deserve rights?”

The first question is, “What does the system do, to whom, how fast, under whose control, with what correction route?”

An automated benefit system can narrow reachability. A model can classify someone into suspicion. A recommender can shape attention. A platform can remove livelihood access. A robot can physically act. A medical model can influence treatment. A hiring tool can close opportunity. A governance model can make public power more opaque.

Mechanical Ethics can constrain these systems before resolving whether any artificial system has standing.

This prevents two opposite failures.

The first failure says: because the system is only a tool, no new ethical issue exists.

The second says: because the system seems intelligent, we must immediately treat it as a full subject.

Both are too fast.

The system can be ethically dangerous as delegated power even if it has no protected standing.

The value-spine version is direct:

```
AI_as_power_over_viable_becoming :=  
  artificial_system  
  changes_or_routes  
  continuation  
  + orientation  
  + capacity  
  + relation  
  + repairability  
  + meaningful_future_space
```

for_affected_scopes

An artificial system may close viable becoming by denying access, confusing orientation, shifting burden, damaging relation, hardening a label, or making repair unreachable. It may also support viable becoming by expanding access, translating information, detecting risk early, reducing arbitrary discretion, preserving records, or helping repair arrive before hardening.

Neither direction is automatic.

```
AI_value_question :=
  what_viable_becoming_expands
+ what_viable_becoming_contracts
+ who_carries_error
+ what_correction_stays_live
```

16.2 13.2 Optimization and Local Reward Pull

Artificial systems optimize.

Not always explicitly. Not always well. But many are built around some target, score, reward, prediction, loss function, engagement metric, compliance metric, cost metric, risk metric, or success signal.

```
AI_local_reward_pull :=
  objective_or_proxy
+ training_signal
+ deployment_incentive
+ repetition_at_scale
```

The danger is that the system may learn the proxy better than the world.

A fraud model may learn suspicion patterns. A recommender may learn outrage. A hiring system may learn historical exclusion. A care allocation system may learn spending patterns rather than need. A productivity tool may learn visible output rather than human cost. A moderation system may learn what is easiest to remove rather than what is most harmful.

This is not only a technical defect. It is a future-space defect.

When optimization changes routes for people, the proxy becomes power.

Mechanical Ethics asks:

```
optimization_audit :=
  what_is_optimized
+ what_proxy_replaces
+ who_benefits
+ who_carries_error
+ can_error_be_seen_and_corrected
```

An AI system may be accurate on average and still route intolerable burden to those who cannot contest. It may improve efficiency and still close reachable paths. It may reduce one kind of human bias and introduce another kind of opacity.

Optimization is not legitimacy.

Nor is local reward structural good.

```
AI_local_reward_not_structural_good :=
  optimized_proxy
!= viable_becoming
unless_audit_shows
  affected_scopes_visible
+ burden_visible
+ correction_live
+ residue_recorded
```

A benchmark may improve while trust decays. A model may become faster while contestability collapses. A system may reduce cost while increasing burden on the least reachable. A score may become more predictive while narrowing opportunity for those labelled by it.

The value test is not only whether the proxy improved. It is whether the deployment preserves or expands viable becoming without hiding what it closes.

```
AI_value_metric_test :=
  proxy_gain
+ real_world_path_change
+ affected_scope_burden
+ correction_capacity
+ residue
```

16.3 13.3 Opacity and Estimator Contamination

Artificial systems can make estimates that look independent while carrying the assumptions, incentives, data, and design choices of their builders or deployers.

```
AI_estimate :=
  model_output
  shaped_by
  data + objective + deployment_context + human_incentive
```

A model output can be useful evidence. It can also be contaminated evidence.

Contamination does not mean the output is false. It means the output's authority must be demoted or checked where material action depends on it.

```
CONTAMINATED_SIGNAL_response :=
  demote_evidential_authority
+ require_independent_check_if_material
+ restrict_irreversible_action_if_unchecked
+ preserve_challenge_route
- demote_moral_standing_by_default
```

The last line matters. A contaminated signal about a person, group, animal, or system does not automatically reduce that scope's standing. It reduces the authority of the estimate.

Opacity becomes dangerous when no one can explain why the system acted, affected scopes cannot challenge, and institutions treat the output as neutral because it came from a machine.

Machine output is not neutral by origin. It may be useful. It may be better than prior human practice. It may reveal patterns humans missed. But its authority must be earned through evidence, monitoring, contestability, and fit to the decision.

A black box that changes futures needs a bright record around it.

For the value spine, opacity also hides false goods. A system may claim access, speed, personalization, safety, or efficiency while hiding the path by which it damages orientation, trust, reachability, relation, or repair.

```
AI_opacity_value_risk :=
  claimed_good_visible
+ path_effect_hidden
+ affected_scope_cannot_orient
+ correction_route_unusable
```

16.4 13.4 Human-in-the-Loop Theatre

A human in the loop is not automatically meaningful control.

Sometimes the human reviewer has no time, training, authority, evidence, or real option to disagree. Sometimes they rubber-stamp the model. Sometimes the organization treats human presence as a liability shield. Sometimes the system is designed so that the human sees only the model's conclusion, not the route or uncertainty.

```
human_in_loop_theatre :=
  nominal_human_review
+ no_meaningful_control
+ automation_authority_preserved
```

Mechanical Ethics asks whether the human can actually change the path.

```

meaningful_human_control :=
  understands_relevant_basis_enough
  + has_time
  + has_authority
  + has_access_to_contradiction
  + can_pause_or_reverse
  + is_not_punished_for_disagreement

```

Human review may still be imperfect. It does not need omniscience. But it must be more than ritual.

A human reviewer who cannot see affected-scope evidence is not meaningful control. A human reviewer whose performance metric rewards agreement is not independent. A human reviewer who can only approve or delay without changing the mechanism may be a decorative gate.

The presence of a human should not be used to absolve an automated system unless the human has real power to contest, correct, or stop the path.

The positive requirement is that human control must preserve real routes for viable becoming where the system would otherwise harden too fast.

```

human_control_value_function :=
  preserve_contradiction
  + restore_context
  + keep_repair_live
  + prevent_false_finality
  + answer_for_affected_scope

```

16.5 13.5 Delegated Power

Artificial systems often act through delegation.

Humans and institutions delegate sorting, prediction, recommendation, routing, triage, drafting, matching, moderation, optimization, and control.

Delegation may be legitimate. It may reduce error, increase access, lower cost, expand capability, or detect harm earlier.

But delegated power still needs answerability.

```

AI_delegation_audit :=
  task_delegated
  + delegating_actor
  + retained_control
  + affected_scope
  + monitoring_route
  + correction_route
  + failure_owner

```

The phrase “the AI did it” is not an answer.

Who chose the system? Who set the objective? Who selected data? Who defined success? Who monitored outputs? Who saw warnings? Who could pause deployment? Who benefited from speed? Who carried error? Who repaired when error hardened?

Delegation can distribute responsibility, but it cannot dissolve it.

The more powerful the artificial system, the stronger the duty to maintain routes for human and institutional answerability.

If no one can answer for the system, the system should not be allowed to exercise serious power over future-space.

Delegated AI may also be a positive route. It can make knowledge reachable, reduce delay, preserve records, detect missed patterns, or help people navigate systems that were previously closed to them.

But the benefit must be audited.

```
AI_delegated_benefit_audit :=
  claimed_expansion
  + affected_scope_access
  + remaining_burden
  + correction_owner
  + residue_if_system_fails
```

Benefit is not absolution. AI that expands one route while closing another still needs a record.

16.6 13.6 Artificial Moral Standing

Artificial moral standing remains unresolved.

Mechanical Ethics should neither grant it casually nor rule it out by category forever.

```
artificial_moral_standing :=
  unresolved
  + signal_sensitive
  + precaution_possible
  + no_default_full_standing
  + no_default_erasure
```

Most present artificial systems may not have moral standing. They may simulate distress, preference, agency, or relationship without experiencing anything. They may be tools shaped by human data and incentives. Treating them as full subjects too early could launder corporate interest, distract from human harm, or create false duties.

But a framework that permanently excludes artificial standing may fail if future systems develop properties that make exclusion reckless.

The right posture is disciplined uncertainty.

Relevant questions might include persistence, integrated memory, self-protective behaviour, preference continuity, welfare-like signals, dependence, social embeddedness, capacity for welfare or frustration, and whether irreversible loss under uncertainty carries serious error cost.

These are not a test that proves standing. They are warning lights.

```
AI_standing_warning_lights :=
  persistence
  + preference_continuity
  + self_protection
  + welfare_like_signal
  + dependency
  + irreversible_loss_risk
  + high_error_cost
```

When warning lights are weak, constraint may still apply because humans are affected by the system. When warning lights strengthen, precaution may become appropriate. The book does not settle the final floor.

The value-spine discipline is to keep two records separate.

```
AI_two_records :=
  AI_as_power_record
  + possible_AI_standing_record
```

The first record asks what the artificial system does to other scopes. The second asks whether the artificial system itself may become a scope. Collapsing them creates confusion: constraining AI power does not require proving AI personhood, and leaving AI standing open does not excuse present systems from duties to affected humans, animals, institutions, or future persons.

16.7 13.7 Future Minds

Future minds are not only artificial.

They may be future children, future humans under changed conditions, enhanced humans, collective systems, artificial agents, biological hybrids, or forms of experience that do not fit present categories.

Mechanical Ethics should not pretend to know them in detail.

But it can preserve a posture toward them:

```
future_mind_posture :=
  do_not_close_recognition_categories_too_early
+ avoid_irreversible_loss_under_high_uncertainty
+ preserve_evidence_and_contestability_where_possible
+ record_open_questions
```

Human ethical history includes repeated failures of recognition. People outside the dominant category were treated as less than fully countable. Animals were treated as mere things. Future people were treated as abstractions. Disabled people, children, prisoners, outsiders, and the poor were often made procedurally invisible.

This history does not prove future artificial standing.

It does warn against building moral systems that can only recognize what current power already recognizes.

Future minds require humility without fantasy.

The book should keep the category open while refusing speculative overclaim.

The value-spine version is that future minds may carry forms of viable becoming we cannot yet model well. That uncertainty does not grant full standing by default. It does create a reason not to design recognition systems that make future recognition impossible.

```
future_mind_value_guard :=
  recognition_category_open
+ high_error_cost_recorded
+ irreversible_closure_avoided_where_possible
+ no_speculative_absolution
```

16.8 13.8 Alignment as Corrigible Constraint

Alignment is often described as making artificial systems do what humans want.

Mechanical Ethics frames the problem more cautiously.

Humans do not have one stable will. Human wants can be cruel, confused, short-term, captured, coerced, or self-destructive. Institutions can want legitimacy without correction. Markets can want profit without residue. States can want control without answerability.

So alignment cannot mean simple obedience to present human preference.

A better candidate is corrigible constraint under uncertainty.

```
alignment_candidate :=
  preserve_correction_capacity
+ expose_uncertainty
+ keep_affected_scopes_visible
+ resist_hidden_burden_shift
+ maintain_human_and_institutional_answerability
```

This is not final doctrine.

There is an open tension:

```
corrigible_constraint_open_tension :=
  maximal_corrigibility
  may_conflict_with
  binding_commitment_and_trust
```

A system that can always be corrected by whoever currently has power may be too manipulable. A system that resists correction may become dangerous. A system that obeys law may enforce unjust

law. A system that obeys users may amplify harmful users. A system that refuses all harm may be useless in dirty conflict.

Alignment therefore cannot be only technical obedience. It must include structures of answerability, contestability, protected-scope recognition, uncertainty handling, and residue.

The value-spine version asks whether alignment preserves viable becoming across affected scopes while keeping correction live.

```
alignment_value_candidate :=
  AI_power_constrained
+ viable_becoming_not_optimized_away
+ uncertainty_exposed
+ affected_scopes_visible
+ correction_and_contest_live
+ residue_recorded
```

The open tension should remain visible.

16.9 13.9 Artificial Systems and False Goods

Artificial systems can produce false goods at scale.

They can create speed that outruns correction. Personalization that traps attention. Safety that removes contest. Efficiency that erases least-reachable scopes. Engagement that damages orientation. Support that becomes extraction. Accuracy that becomes unchallengeable authority. Automation that makes burden look like self-service.

```
AI_false_good_patterns :=
  speed_without_correction
+ personalization_without_exit
+ safety_without_contest
+ efficiency_with_erased_burden
+ accuracy_without_answerability
+ support_that_shifts_burden
```

This does not mean artificial systems are bad by default. It means claimed benefits must be routed through the same audit as any other claimed good.

```
AI_claimed_good_audit :=
  what_expands
+ for_whom
+ what_contracts
+ who_carries_error
+ what_correction_stays_live
+ what_residue_remains
```

A beneficial artificial system should make good paths more reachable without making bad records harder to challenge.

16.10 13.10 What AI Systems Must Not Be Allowed to Hide

Artificial systems should not be allowed to hide the same things human systems must not hide.

They must not hide loss.

They must not erase affected scopes.

They must not treat uncertainty as certainty.

They must not pretend repair where residue remains.

They must not turn correction channels into harm carriers.

They must not let the assessed actor set every clock.

They must not use emergency, novelty, scale, intelligence, or technical complexity to launder prior failure.

They must not optimize false goods while hiding what becomes less viable.

```
AI_must_not_hide :=  
  loss  
  + affected_scopes  
  + uncertainty  
  + clock_setting  
  + burden_shift  
  + correction_failure  
  + residue  
  + responsibility_owner  
  + false_good  
  + viable_becoming_contraction
```

The AI version of these failures may be harder to see. Loss may appear as score movement. Scope erasure may appear as missing labels. Uncertainty may appear as confidence. Burden shift may appear as self-service appeal. Correction failure may appear as support ticket volume. Residue may appear nowhere. A false good may appear as benchmark success.

That is why artificial systems require strong records around them.

If an artificial system changes future-space, there must be a route to ask what changed, for whom, how fast, under whose authority, with what correction, with what remaining residue, and with what effect on viable becoming.

Part 13 ends with the split intact:

```
AI_split_preserved :=  
  constrain_artificial_systems_as_power  
  + leave_artificial_standing_open_under_discipline
```

And the value-spine addition:

```
AI_value_spine_preserved :=  
  local_reward_not_structural_good  
  + AI_as_power_over_viable_becoming  
  + future_minds_open_without_erasure  
  + claimed_AI_good_requires_burden_record
```

The next part turns from theory to use: case records, protocols, router findings, residue ledgers, and smoke tests.

17 Part 14 — Case Records and Use Protocols

Status: first draft. Incomplete. Not canon.

17.1 14.1 Why Use Protocols Belong Near the End

The book has deliberately waited to put the machinery here.

The reader first had to move through the human argument: something is here; it acts under partial perception; it is not alone; action changes future-space; time closes paths; description is not permission; gates keep the record honest; protection needs reasons; correction is not repair; dirty conflicts remain dirty; responsibility follows visibility only through role, evidence, and control; power must answer; civilization must remain correctable and support shared viable becoming; artificial systems must not hide behind opacity or optimized false goods.

Only now should the operator machinery appear.

```
use_protocol_position :=  
  after_human_spine  
  not_inside_human_spine
```

This order protects the book from becoming a form before it becomes an argument.

A case record is useful only if the reader understands why the fields matter. Otherwise the template becomes another bureaucratic ritual: complete the boxes, produce the finding, claim discipline, move on.

Mechanical Ethics should not become that.

The use protocols are not the soul of the book. They are the handle. They help a reader, reviewer, institution, or system apply the argument without silently changing it.

The shortest adequate record is usually the best record.

```
protocol_minimization_rule :=
  use_smallest_record
  that_preserves_route_visibility
+ uncertainty_status
+ affected_scope
+ timing
+ viable_becoming_change
+ residue
+ responsibility_remaining
```

A low-stakes case should not become a ceremonial file. A high-stakes case should not be compressed until the path effect disappears. The point is not to create paperwork. The point is to keep a live route from question to finding without hiding the parts that matter.

The protocol should expand only when the situation demands it.

```
anti_bureaucracy_rule :=
  record_enough_to_keep_answerability_alive
- perform_framework_for_its_own_sake
```

17.2 14.2 How to Build a Case Record

A case record begins with the decision or question on the table.

Without that, the finding floats.

```
case_record_begins_with :=
  decision_or_question_on_the_table
```

The live question should be concrete enough to route. Not “what is justice here?” but “should this decision proceed?” “should this channel be treated as adequate correction?” “should this system pause enforcement?” “what repair duties remain after late correction?” “does this case require escalation because standing or authority remains unresolved?” “does the claimed benefit hide remaining burden?”

A minimum viable case record contains:

```
minimum_viable_case_record :=
  decision_or_question_on_the_table
+ transition_description
+ at_least_one_affected_scope
+ future_space_comparison_for_that_scope
+ viable_becoming_record
+ at_least_one_clock_using_TRACE_bands
+ first_gates_results
+ router_finding
+ prescription_audit
```

Everything else should be filled when known. Missing fields must be marked missing or unknown, not silently invented.

The discipline is simple: do not pretend to know what the record does not show.

17.3 14.3 TRACE on the Left, ME on the Right

For full operator use, the case record should keep description and prescription separated.

```
left_side := TRACE_or_descriptive_case_map
right_side := ME_prescriptive_response
```

Left side asks:

```
descriptive_questions :=
  what_changed
  + who_or_what_was_affected
  + how_future_space_changed
  + what_clocks_apply
  + who_estimated
  + what_channels_exist
  + what_burden_shifted
  + what_claimed_good_appears
  + what_residue_remains
```

Right side asks:

```
prescriptive_questions :=
  which_gates_fail
  + which_scopes_need_protection
  + which_duties_follow
  + whether_conflict_is_dirty
  + whether_claimed_good_survives_audit
  + who_must_answer
  + whether_power_is_answerable_enough
  + what_router_finding_is_safe
```

The left side must not quietly become permission. The right side must not silently rewrite the left side.

If the descriptive record says unknown, Mechanical Ethics must either work with the unknown, seek more mapping, hold, escalate, or explain why reversible action remains acceptable despite uncertainty.

A clean case record preserves the firewall.

17.4 14.4 Viable Becoming Record

The value-spine patch requires one new discipline: case records must not track harm alone.

They must also track what opens, what closes, and what remains viable for affected scopes.

```
viable_becoming_record :=
  affected_scope
  + continuation_change
  + orientation_change
  + capacity_change
  + relation_change
  + repairability_change
  + meaningful_future_space_change
  + confidence
```

The record can be short. It can use plain language. It does not require false measurement. But it must ask the question.

A transition may reduce one kind of harm while damaging trust. It may expand access while shifting burden. It may preserve safety while narrowing relation. It may speed correction while reducing contestability. It may create opportunity for one scope while closing opportunity for another.

The viable becoming record prevents the case from being flattened into one word: harm, benefit, efficiency, safety, progress, compliance, or risk.

```
viable_becoming_record_rule :=
  no_single_label_can_replace
  scoped_change_record
```

This field should remain uncertain where the record is uncertain.

17.5 14.5 Claimed Good and Positive Audit

Mechanical Ethics now needs a positive audit as well as a harm audit.

A claimed good may be real. It may justify action. It may preserve life, repair a route, expand access, reduce burden, rebuild trust, support learning, coordinate shared work, or prevent worse loss.

But a claimed good can also hide burden.

```
claimed_good_audit :=
  what_good_is_claimed
+ for_whom
+ what_viable_becoming_expands
+ what_viable_becoming_contracts
+ who_carries_cost_or_risk
+ what_correction_stays_live
+ what_residue_remains
```

This is the operational form of benefit is not absolution.

```
benefit_is_not_absolution_operational :=
  claimed_good
  requires
    affected_scope_record
  + burden_record
  + correction_route
  + residue_check
```

A positive audit should not become optimism theatre. It should not list benefits as a shield. It should test whether the claimed benefit actually reaches affected scopes, whether the route remains contestable, and whether remaining burden is visible.

A good case record can therefore say: this transition appears to create a real good, under stated uncertainty, but it still imposes residue and duties.

17.6 14.6 Router Labels as Findings, Not Authorizations

Router labels are often where misuse begins.

They look official. They compress complexity. They can be copied into decision memos. They can become institutional shields.

So every router label must be read narrowly.

```
ME_router :=
  NO_STRUCTURAL_OBJECTION_FOUND
| PROCEED_ONLY_WITH_NAMED_CONSTRAINTS
| PAUSE
| HALT
| ESCALATE
| REPAIR_REQUIRED
| COMPENSATION_REQUIRED
| INDEPENDENT_REVIEW_REQUIRED
| HOLD
| UNRESOLVED_VALUE_CRUX
| REGRESSION_RISK
```

Every router output means:

```
router_output :=
  finding
+ recommendation_or_constraint
+ remaining_actor_responsibility
```

It does not mean:

```
router_output != authorization
```

NO_STRUCTURAL_OBJECTION_FOUND is especially narrow. It means no structural objection was found in the available record for the named decision or question. It does not mean the action is right, lawful, legitimate, safe, final, complete, or immune from later challenge.

A failed unanswered gate forbids that finding.

```
if any first_gate == fail and unanswered:
    router_finding != NO_STRUCTURAL_OBJECTION_FOUND
```

A failed claimed-good audit can also block clean routing.

```
if claimed_good_used_to_clear_action
and claimed_good_audit_missing_material_burden:
    router_finding != NO_STRUCTURAL_OBJECTION_FOUND
```

This consistency rule prevents benefit laundering.

17.7 14.7 Prescription Audit

Every prescription should be audited.

Not every case needs a long essay. But every case needs enough route to see why the output follows.

A prescription audit asks:

```
prescription_audit :=
    decision_or_question_on_the_table
    + facts_read
    + primitives_used_without_redefinition
    + value_priority_added
    + claimed_good_audit
    + scopes_protected_and_why
    + uncertainties_remaining
    + contaminated_estimates
    + duties_and_structural_pressure
    + dirty_or_unrepaired_loss
    + residue_persisting
    + responsibility_remaining
    + laundering_risk
    + fail_state_considered
```

The most important field may be value_priority_added.

Description alone does not decide. Mechanical Ethics adds value priorities when it says preserve correction, protect vulnerable scopes, avoid irreversible action under uncertainty, keep contestability live, record residue, choose the least-unreparable path in dirty conflict, or preserve viable becoming where a scope depends on an actor-controlled route.

Those priorities should be visible.

If Mechanical Ethics hides its own value additions, it becomes the same kind of authority it criticizes.

17.8 14.8 Residue Ledger

A residue ledger records what remains.

It should be short when residue is small. It should be serious when residue is serious. It should not become ceremonial.

```
residue_ledger :=
    loss_not_restored
    + time_not_returned
    + trust_damage
    + opportunity_expired
    + record_scar
    + recurrence_risk
    + acknowledgement_needed
    + viable_becoming_not_restored
```

The ledger should answer:

```
residue_questions :=
  what_remains
+ who_carries_it
+ can_it_be_reduced
+ who_still_owes_attention
+ what_prevents_recurrence
+ what_future_remains_closed
```

Residue is where official closure often lies.

A case may be administratively closed and ethically open. A payment may be restored and time still lost. A decision may be corrected and trust still damaged. A dirty choice may be justified and the sacrificed scope still gone.

The residue ledger prevents the file from pretending the world has rewound.

17.9 14.9 Challenger Function

Every serious use of Mechanical Ethics needs a challenger function.

The challenger asks whether the framework is laundering what it claims to expose.

```
challenger_function :=
  search_for_hidden_loss
+ search_for_erased_scope
+ search_for_fake_certainty
+ search_for_hidden_value_priority
+ search_for_permission_drift
+ search_for_benefit_laundering
+ search_for_residue_erasure
```

The challenger does not need to be hostile in personality. It needs to be structurally independent enough to see what the first actor misses or benefits from missing.

A challenger may be a person, panel, reviewer, affected-scope representative, external auditor, court, ombuds route, community witness, technical auditor, or deliberately adversarial review process.

The function matters more than the label.

A captured challenger is not a challenger. A decorative challenger is not a challenger. A reviewer who cannot see the record is not a challenger. A reviewer punished for disagreement is not a challenger.

Mechanical Ethics should be built with teeth against itself.

17.10 14.10 Worked Case Type 1 — Human Administrative Route

A human administrative route case asks whether a public or institutional process has changed a person's future-space through decision, delay, record, classification, or burden.

Typical fields:

```
administrative_case :=
  decision
+ affected_person_or_group
+ record_or_rule
+ delay_or_burden
+ correction_route
+ viable_becoming_record
+ claimed_good_audit
+ residue
```

Likely questions:

Does correction arrive before hardening?

Does the affected person carry the system's uncertainty?

Is the appeal route reachable?

Was contradiction preserved?

Who set the clock?

What viable becoming changed?

What remains after correction?

This case type is where Mechanical Ethics should be most practical. It must not require the affected person to master the whole framework. The burden should remain on the system and the reviewer to make the path visible.

17.11 14.11 Worked Case Type 2 — Platform or Algorithmic System

A platform or algorithmic case asks whether technical mediation has changed future-space through classification, ranking, removal, recommendation, scoring, automation, or opacity.

```
platform_algorithmic_case :=  
  system_output  
  + affected_scope  
  + proxy_or_objective  
  + opacity_level  
  + appeal_or_correction_route  
  + burden_shift  
  + viable_becoming_record  
  + claimed_good_audit  
  + monitoring_and_owner
```

Likely questions:

What proxy is being optimized?

Who carries false positives or false negatives?

Can the affected scope understand and contest the decision?

Does human review have meaningful control or only theatre?

Who owns failure when the system acts at scale?

Does the system support viable becoming, damage it, or both?

Does correction arrive before account loss, income loss, reputational loss, or access loss hardens?

The central risk is treating technical complexity as moral distance.

Mechanical Ethics should not allow that.

17.12 14.12 Worked Case Type 3 — Dirty Conflict

A dirty conflict case asks what to do when every available action causes material loss.

```
dirty_conflict_case :=  
  loss_by_path  
  + claimed_good_by_path  
  + affected_scopes  
  + viable_becoming_by_scope  
  + irreversibility  
  + repairability  
  + least_reachable_scope  
  + value_tilt  
  + residue
```

Likely questions:

Is action required, or is the sense of necessity manufactured?

Who is harmed by each path?

What good is claimed by each path?

Which viable-becoming claims are incompatible?

Which losses are least repairable?

Which scopes cannot contest?

What value priority is being added?

What minimum floors apply, if any are known?

What residue remains after the chosen path?

What repair duties follow?

This case type must never return a clean answer where loss remains dirty.

17.13 14.13 Worked Case Type 4 — Positive Route With Burden

The value spine needs a fourth case type.

Some transitions are genuinely beneficial but still ethically incomplete.

```
positive_route_with_burden_case :=  
  claimed_good  
  + affected_scopes  
  + viable_becoming_expanded  
  + viable_becoming_contracted  
  + burden_carrier  
  + correction_route  
  + residue
```

Likely questions:

What becomes newly reachable?

For whom?

Who pays the cost?

Is the benefit real or only a metric?

Can affected scopes contest the burden?

Does the positive route preserve repairability?

What residue remains despite the benefit?

This case type keeps the framework from becoming harm-only while also preventing benefit laundering.

17.14 14.14 Low-Fit Case Returning COMPRESSION_ONLY

Not every case needs Mechanical Ethics.

Some cases are ordinary coordination, low-stakes preference, reversible scheduling, routine choice, or simple information update. Mechanical Ethics may still describe the case, but it should not inflate itself.

```
low_fit_case :=  
  reversible  
  + low_material_loss  
  + affected_scopes_not_erased  
  + correction_live  
  + no_serious_gate_failure  
  + no_material_claimed_good_audit_needed
```

The proper result may be COMPRESSION_ONLY at the descriptive level and a modest ME finding such as no structural objection found under named ordinary constraints.

This is not failure.

It is good discipline.

A framework that turns every meeting move into a moral crisis is unusable. A framework that cannot distinguish low-stakes coordination from hard conflict will train people to ignore it.

Mechanical Ethics should be able to stay small when the case is small.

17.15 14.15 How This Part Should Be Used

Part 14 is not a new bureaucracy.

It is a way to keep the book's argument from drifting when applied.

Use the full operator schema for serious cases. Use the minimum viable record for early worked examples. Use the smoke test to check that the system does not overinflate modest cases. Use the challenger function to catch laundering.

Use less than the full schema when the case is small. Use more only when the missing detail would materially change the finding. The case record should grow because the situation demands it, not because the framework wants to display itself.

```
use_protocol_summary :=  
  name_question  
  + separate_description_from_prescription  
  + fill_minimum_record  
  + run_gates  
  + record_viable_becoming_change  
  + audit_claimed_good  
  + derive_duties  
  + check_dirty_conflict  
  + attach_responsibility  
  + audit_prescription  
  + record_residue
```

The next part closes the first draft by naming limits, failures, and open questions.

A book that cannot name its limits should not be trusted with hard cases.

18 Part 15 — Limits, Failures, and Open Questions

Status: first draft. Incomplete. Not canon.

18.1 15.1 Why the Book Must End With Limits

A book about constraint must constrain itself.

Mechanical Ethics is most dangerous if it sounds complete too early. It has built a route from minimal origin to prescriptions, but that route is not final moral law. It is an auditable structure for action under uncertainty.

```
book_status :=  
  provisional  
  + auditable  
  + revisable  
  + not_final_moral_law
```

The book has made strong claims.

Action changes future-space. Time changes repair. Correction is not repair. Residue remains. Description is not permission. Dirty conflict must not be cleaned. Responsibility requires more than visibility. Power must answer. Civilization should remain correctable and support shared viable becoming. Artificial systems must not hide loss, burden, or false goods behind opacity.

Those claims may still fail in use.

They may overreach. They may under-protect. They may create paralysis. They may be captured by institutions. They may be turned into forms. They may hide value judgments behind procedural language. They may misuse viable becoming as a slogan. They may work well on administrative cases and badly on intimate, cultural, spiritual, ecological, or geopolitical cases.

So the ending must not celebrate arrival.

It must preserve revision.

18.2 15.2 What ME Cannot Decide

Mechanical Ethics cannot decide everything.

It cannot settle all moral standing. It cannot produce a complete theory of legitimacy. It cannot solve punishment and desert. It cannot fully define dignity. It cannot always choose between incommensurable losses. It cannot make dirty conflict clean. It cannot guarantee correct action from incomplete evidence. It cannot validate TRACE. It cannot remove actor responsibility.

It also cannot prove a final theory of the good.

```
ME_cannot_decide :=  
  final_moral_status_all_scopes  
  + complete_legitimacy_theory  
  + full_punishment_desert_theory  
  + all_acceptable_sacrifice  
  + all_dirty_conflicts  
  + all_incommensurable_values  
  + final_theory_of_the_good
```

This does not make the book useless.

A map can be useful without being the territory. A warning can be valid without solving the whole case. A brake can prevent disaster without steering the whole vehicle. A residue ledger can expose loss without knowing the final theory of justice.

Mechanical Ethics is strongest when it identifies what must not be hidden before harder decisions proceed.

It is weakest when asked to become a total doctrine.

18.3 15.3 Final Theory of the Good Remains Open

The value spine added a positive structure to the book: viable becoming, care, trust, learning, repair, cooperation, play, meaning, love, shared worlds, and good routes made ordinary.

That was necessary. A framework that only names harm cannot say enough about why repair matters, why trust is worth rebuilding, why capacity and relation matter, or why civilization should preserve more than bare survival.

But the positive structure is still a candidate.

```
viable_becoming_status :=  
  candidate_structural_account  
  + not_final_theory_of_good  
  + must_remain_scoped  
  + must_remain_revisable
```

Viable becoming is useful because it is mechanically inspectable. It asks whether a transition changes continuation, orientation, capacity, relation, repairability, and meaningful future-space for named scopes under uncertainty.

It does not prove that these are the only goods.

It does not solve conflicts between them.

It does not rank every scope.

It does not decide when one scope's becoming may permissibly narrow another's.

It does not replace love, dignity, beauty, worship, grief, loyalty, excellence, freedom, justice, joy, or forms of meaning that do not compress neatly into the current grammar.

So the value spine must remain disciplined:

```
value_spine_guard :=
  use_viable_becoming_as_probe
  not_as_total_doctrine
  + name_scope
  + name_time
  + name_cost
  + name_uncertainty
  + name_residue
```

The book may say that viable becoming is a strong structural candidate for tracking good and bad transitions.

It must not say that it has finished goodness.

18.4 15.4 UNKNOWN Is Not Failure

UNKNOWN is not failure.

It is a valid state.

```
UNKNOWN :=
  material_field_not_settled
  + no_safe_silent_inference
```

A framework that cannot say unknown will lie. It will invent facts, inflate confidence, force premature routing, or hide uncertainty behind authority.

But unknown is not automatic paralysis either.

If action is reversible, low-stakes, and correction remains live, an unknown may be tolerable. If delay would cause greater loss, action under uncertainty may be needed. If the unknown concerns an irreversible action against a vulnerable scope, hold or escalation may be required.

```
unknown_response_depends_on :=
  reversibility
  + stakes
  + timing
  + affected_scope
  + correction_capacity
  + dirty_conflict_status
```

Unknown should change the route. It should not disappear.

18.5 15.5 When ME Makes Things Worse

Mechanical Ethics can make things worse.

That must be said plainly.

It can make things worse by slowing action where delay is the harm. It can make things worse by creating paperwork instead of repair. It can make things worse by giving institutions better language for self-protection. It can make things worse by turning every small coordination issue

into ethical theatre. It can make things worse by hiding value judgments behind structured fields. It can make things worse by allowing powerful actors to decide that they have satisfied the protocol.

```
ME_regression_risks :=
  paralysis
+ bureaucracy_theatre
+ authority_launders
+ overclassification
+ hidden_value_priority
+ false_completion
```

The value-spine layer adds its own failure modes.

```
positive_value_language_failure_modes :=
  viable_becoming_used_as_slogan
+ claimed_good_used_as_clearance
+ benefit_launders
+ flourishing_paternalism
+ capacity_language_used_to_coerce
+ trust_language_used_to_silence_warning
+ care_language_used_to_capture_dependency
+ meaning_language_used_to_erase_loss
```

This is a serious risk. Positive language is easier to weaponize than negative constraint. It sounds generous while hiding control.

An institution can say it is expanding capacity while imposing dependency. A platform can say it is building community while trapping attention. A state can say it is protecting future generations while silencing present affected scopes. A system can say it is preserving trust while punishing the witness who made repair possible.

So the positive value layer must remain tied to records, not slogans.

```
positive_claim_safety_rule :=
  claimed_good_must_pass
  affected_scope_record
+ burden_record
+ contestability_check
+ residue_check
```

This is why REGRESSION_RISK must exist as a router output.

Sometimes applying the framework may degrade the situation. The correct move may be to stop using it, use a narrower tool, seek affected-scope testimony, defer to domain law, escalate to independent review, or act under emergency constraints with residue recorded.

No framework should assume its own use is always beneficial.

Mechanical Ethics must include the possibility that Mechanical Ethics is the wrong tool for a given moment.

18.6 15.6 Value Lock-In

A system can freeze today's values into tomorrow's machinery.

That is value lock-in.

```
value_lock_in :=
  present_value_priority
  embedded_in_durable_system
  reducing_future_correction_or_contest
```

Mechanical Ethics itself could contribute to value lock-in if its open wounds are treated as solved or if its current gates become permanent doctrine without challenge.

This risk is serious in artificial systems, law, public administration, education, markets, and infrastructure. Once a value priority becomes code, procedure, contract, model, metric, or institutional habit, changing it may become expensive or impossible.

Value lock-in is not always bad. Some commitments are meant to resist easy reversal. Stable floors may be necessary for trust, safety, and ordinary life.

But commitment and correction are in tension.

```
commitment_correction_tension :=
  stable_floor_needed
  + future_revision_needed
```

Mechanical Ethics has not fully solved that tension.

It should preserve explicit review points, open-wound records, affected-scope challenge, and refusal to pretend that provisional structures are final law.

Viable becoming also creates lock-in risk. If the current candidate account is embedded too deeply, future goods that do not fit its shape may become harder to see.

```
viable_becoming_lock_in_risk :=
  candidate_good_language
  becomes
  recognition_filter
```

The candidate must remain a probe, not a cage.

18.7 15.7 Authority Creep

Authority creep happens when a framework built to constrain power becomes a source of power.

```
authority_creep :=
  constraint_language
  becomes
  institutional_authority_shield
```

This is a major risk for Mechanical Ethics.

An institution could say: we ran ME, therefore the action is ethical.

A platform could say: the residue ledger is complete, therefore no further duty remains.

A state could say: the dirty-conflict protocol selected this path, therefore the sacrifice is clean.

A company could say: no structural objection found, therefore proceed.

A system could say: the viable-becoming audit found net expansion, therefore the burdened scope is answered.

All of those are misuse.

The book has repeated the guard because repetition is necessary:

```
ME_output != permission_to_act
```

Authority creep should be checked by independent challenge, prescription audit, claimed-good audit, residue visibility, and the constant reminder that actors remain responsible after the finding.

If a Mechanical Ethics record makes power less answerable, the record has failed.

18.8 15.8 Paralysis Risk

Mechanical Ethics can also freeze action.

If every uncertainty triggers hold, if every affected scope creates veto, if every residue prevents closure, if every dirty conflict becomes unbearable, then the framework becomes paralysis.

Paralysis is not harmless.

Inaction changes the world. Delay can harden loss. Opportunity can expire. A system that refuses to act until clean may allow dirtier harm to continue.

```
paralysis_risk :=
  overconstraint
  + time_passing
  + avoidable_loss_from_delay
```

The answer is not to abandon constraint.

The answer is to make the time dimension explicit. If delay causes loss, record that. If action must proceed under uncertainty, name the uncertainty, use constraints, preserve correction, and record residue. If no clean answer exists, treat the case as dirty rather than pretending that waiting is neutral.

Mechanical Ethics should resist both reckless action and self-protective delay.

A good hold is protective.

A bad hold is hidden action by delay.

18.9 15.9 Open Problems

The open problems should remain visible.

```
open_problems :=
  final_floor_foundation_for_scope_protection
  + final_theory_of_the_good
  + acceptable_sacrifice
  + legitimacy_and_coercion
  + punishment_and_desert
  + dignity_beyond_recognition
  + artificial_moral_standing
  + future_mind_recognition
  + corrigibility_vs_binding_commitment
  + value_lock_in
  + validation_of_usefulness
```

Some of these are philosophical. Some are institutional. Some are technical. Some are political. Some are probably not solvable in one book.

The point of naming them is not to weaken the structure. It is to prevent false closure.

A reader should know where the book is firm and where it is provisional.

Firm: do not hide loss. Do not erase affected scopes. Do not treat uncertainty as certainty. Do not pretend repair where residue remains. Do not treat description as permission. Do not clean dirty conflict by language. Do not build responsibility from hindsight alone. Do not use benefit as absolution.

Provisional: exact floors, final standing across all scopes, full legitimacy theory, punishment and desert, artificial standing, final alignment doctrine, complete conflict resolution, final account of good life or flourishing.

That distinction is part of the honesty contract.

18.10 15.10 How This Book Should Be Tested

Mechanical Ethics should be tested by use, not admiration.

A good test asks whether the book changes the quality of a case record, catches hidden loss, prevents permission drift, improves repair, clarifies claimed-good burdens, or clarifies when hold or escalation is needed.

```
testing_questions :=
  did_ME_find_hidden_loss
  + did_ME_prevent_scope_erasure
  + did_ME_preserve_uncertainty
  + did_ME_detect_late_correction
  + did_ME_record_residue
  + did_ME_avoid_hindsight_blame
```



```
+ did_ME_keep_dirty_conflict_dirty
+ did_ME_audit_claimed_good
+ did_ME_track_viable_becoming_without_sloganizing
+ did_ME_avoid_overrouting_low_fit_cases
```

Tests should include low-fit cases where the right answer is modest. If the framework inflates every routine case, it is unusable.

Tests should include hostile cases where the actor wants permission. If the framework becomes a shield, it fails.

Tests should include dirty conflicts where every path causes loss. If the framework cleans the loss, it fails.

Tests should include administrative cases where correction arrives late. If the framework treats correction as repair, it fails.

Tests should include artificial-system cases where opacity hides responsibility. If the framework lets automation dissolve answerability, it fails.

Tests should include positive cases where a real benefit exists but burden remains. If the framework uses the benefit to erase residue, it fails.

Validation remains unclaimed.

18.11 15.11 Ending Without Triumph

This book began thin.

Something is here. It perceives partially. It can change something. It does not know everything. Time passes. It is not alone.

From that beginning, the book has built pressure: action interferes with other continuations; future-space can narrow or expand; viable becoming can be supported or blocked; correction can arrive too late; opportunity can expire; description is not permission; gates keep the record honest; protection needs reasons; repair must touch the path; dirty conflict remains dirty; responsibility requires visibility plus role and control; power must answer; civilization must remain correctable; artificial systems must not hide loss behind opacity; protocols must record claimed goods as well as harms.

That is not everything.

It is enough to begin.

```
ending_claim :=
  these_are_best_current_prescriptive_answers
  under_this_structure
  with_these_limits
  and_open_wounds_visible
```

Mechanical Ethics should not end by saying it has solved ethics.

It should end by saying it has made a promise:

Do not hide the loss.

Do not erase the scope.

Do not fake certainty.

Do not call correction repair when the path remains harmed.

Do not make dirty loss clean by naming a procedure.

Do not let claimed good erase remaining burden.

Do not let positive language become a permission machine.

Do not let power outrun answerability.

Do not let the actor disappear into the machine.

And when the answer is not known, say so before the unknown becomes someone else's permanent burden.